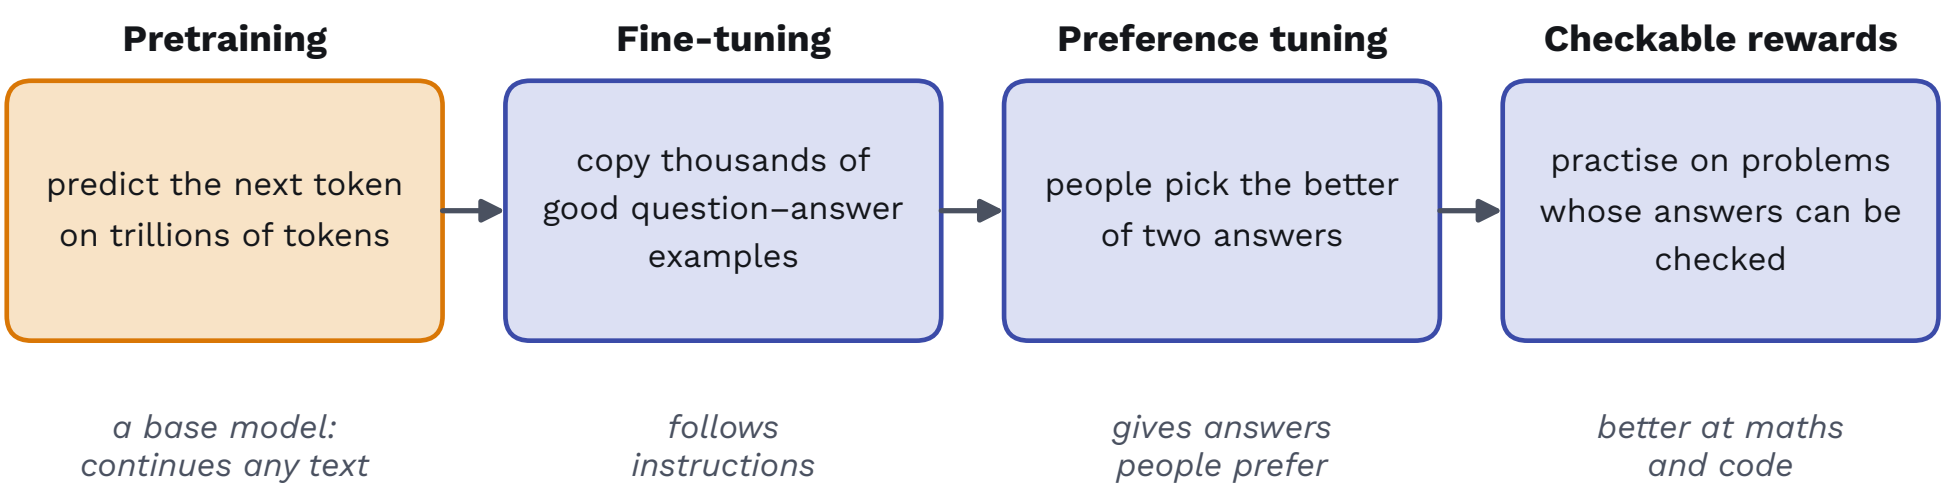
Pretraining
predict the next token on trillions of tokens
a base model: continues any text
Fine-tuning
copy thousands of good question–answer examples
follows instructions
Preference tuning
people pick the better of two answers
gives answers people prefer
Checkable rewards
practise on problems whose answers can be checked
better at maths and code