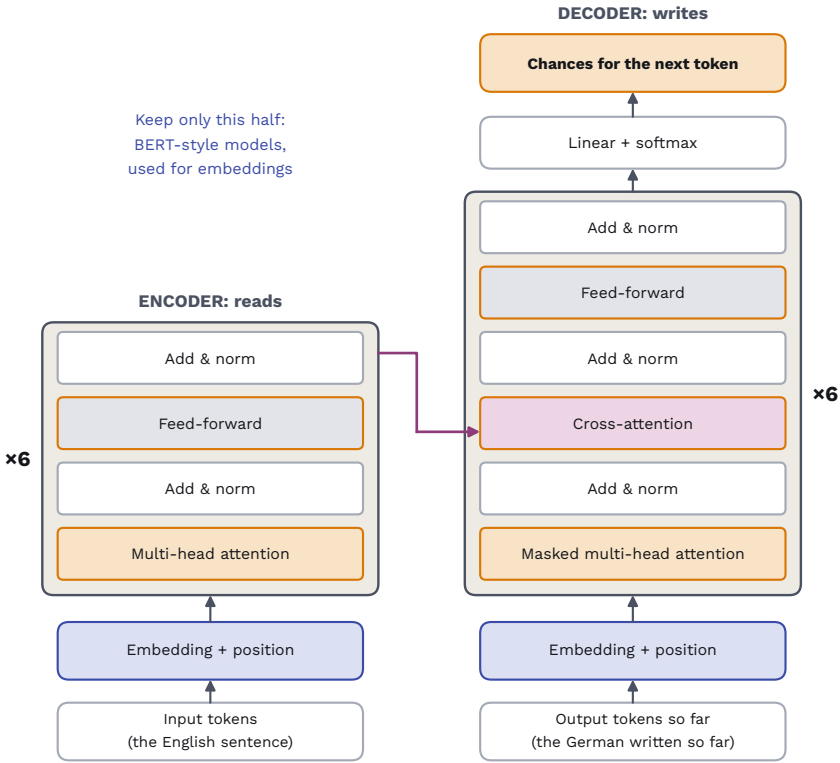
Keep only this half (no cross-attention): GPT-style chat models
DECODER: writes
Chances for the next token
Linear + softmax
Add & norm
Feed-forward
Add & norm
Cross-attention
Add & norm
Masked multi-head attention
×6
Embedding + position
Output tokens so far
(the German written so far)
Keep only this half:
BERT-style models,
used for embeddings
ENCODER: reads
Add & norm
Feed-forward
Add & norm
Multi-head attention
×6
Embedding + position
Input tokens
(the English sentence)