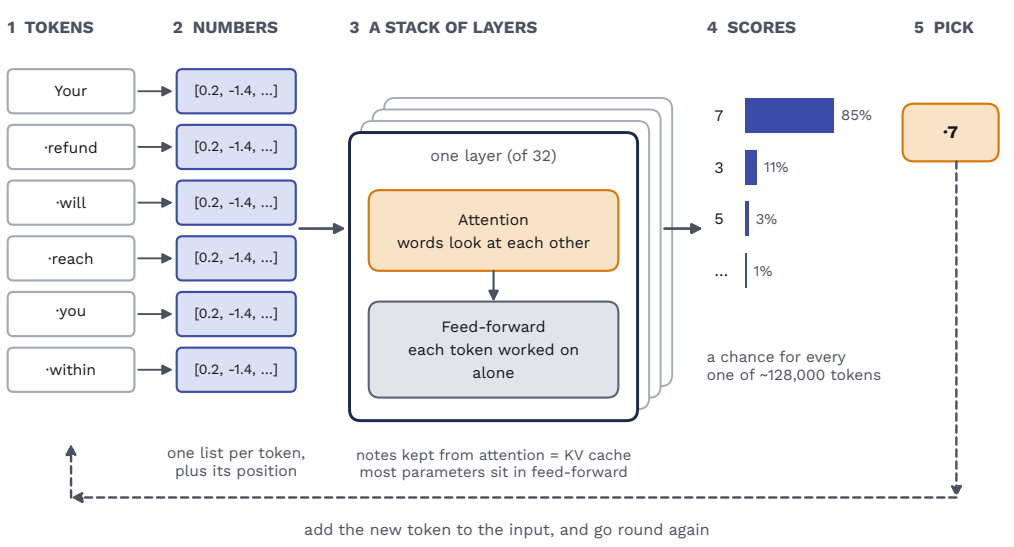
1 TOKENS
2 NUMBERS
3 A STACK OF LAYERS
4 SCORES
5 PICK
Your
·refund
·will
·reach
·you
·within
[0.2, -1.4, ...]
[0.2, -1.4, ...]
[0.2, -1.4, ...]
[0.2, -1.4, ...]
[0.2, -1.4, ...]
[0.2, -1.4, ...]
one layer (of 32)
Attention
words look at each other
Feed-forward
each token worked on
alone
7
3
5
...
85%
11%
3%
1%
·7
a chance for every
one of ~128,000 tokens
one list per token,
plus its position
notes kept from attention = KV cache
most parameters sit in feed-forward
add the new token to the input, and go round again