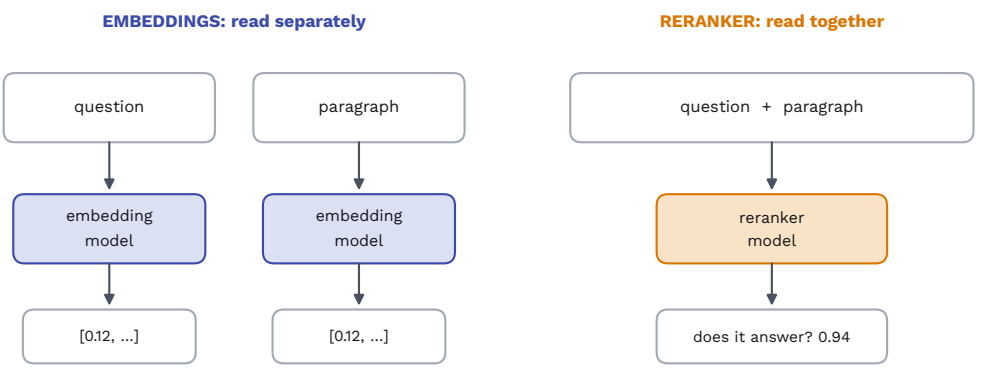
EMBEDDINGS: read separately
question
paragraph
embedding model
embedding model
[0.12, ...]
[0.12, ...]
compare the two lists: fast, paragraphs embedded once in advance
RERANKER: read together
question + paragraph
reranker model
does it answer? 0.94
sees "unopened" against "opened": accurate, but one model run per paragraph