Table of contents

Preface vii

I The model and what goes into it xi

1 What an AI engineer builds 1

2 How LLMs work, for engineers 8

3 Prompts and context 21

II Memory, tools and knowledge 33

4 Memory 34

5 Tools and structured output 44

6 Search inside out 55

7 RAG, end to end 64

8 Advanced RAG 72

III Measuring and deciding 81

9 Evals 82

10 Decide, don’t generate 93

IV Agents 101

11 Agents: the loop 102

12 Guardrails and safety 111

13 MCP, skills and agent protocols 123

14 Long tasks and many agents 133

V Production 143

15 Freshness, scale, caching and cost 144

16 Running it in production 152

17 Designing an AI system, start to finish 160

Appendices 167

A The code, and the versions it was checked with 167

B Glossary 170

References 175

About the author 181