

Most of what an AI agent does isn't writing. **It's deciding.**

Is this alert real? Which team gets this ticket? Is this action safe? Most agents hand every one of those small decisions to a large language model, then dig the answer out of a paragraph. It works, slowly and expensively, and the model sounds just as sure when it's wrong.

This book shows you a better way to build the decision layer. You'll get probabilities you can trust, lines drawn from what each mistake costs, and a clear rule for when to hand a case to a person. Every step is built with runnable code, on a synthetic but realistic security team, using Jev, a new System One model, and a free mock that needs no API key.

[[PRAISE]]

Inside, you'll learn to:

- Check whether a model's probabilities mean what they say, and fix them when they don't
- Turn a probability into act, review or escalate, with lines drawn from real costs and real capacity
- Run a fair six-way bake-off: rules, logistic regression, a text classifier, two LLM methods and Jev
- Build a hybrid agent in which Jev decides and the LLM reads and writes
- Build your own small System One model, with typed heads and calibration built in

Real threats seen by a person same six analysts



From the book's synthetic
case study (Chapter 18)

For students, engineers building agents, and tech leads deciding what to build.

About the author

Sridhar Mukkandi [[AUTHOR BIO]]
sridhar.authorhub@gmail.com



Labs, figures and code

Scan for the companion repository

Artificial Intelligence / Machine Learning
[[PRICE INR]] · [[PRICE USD]]

github.com/Mukkandi-Sridhar/JEVBook


BARCODE AREA

2 × 1.2 in, keep clear

Decide, Don't Generate

Sridhar Mukkandi

Decide, Don't **Gen**er**ate**



Jev, System One Models, and the Decision Layer of Agentic AI

Runnable labs · 130+ figures · no API key needed

Sridhar Mukkandi