

Most of what an AI agent does isn't writing. **It's deciding.**

Is this alert real? Which team gets this ticket? Is this action safe? Most agents give every one of those small decisions to a large language model. Then code has to find the answer inside a paragraph. It works, but it's slow and expensive, and the model sounds just as sure when it's wrong.

This book shows you a better way to build the decision layer. You'll get calibrated probabilities that mean what they say, thresholds set by what each mistake costs, and a clear rule for when to give a case to a person. Every step is built with runnable code, around a realistic (synthetic) security team. You'll work with Jev, a new System One model, through a free mock that needs no API key.

Inside, you'll learn to:

- Check whether a model's probabilities mean what they say, and fix them when they don't
- Turn a probability into act, review or escalate, with thresholds set by real costs and real capacity
- Compare six ways to make the same decision, from hand-written rules to LLMs to Jev
- Build a hybrid agent in which Jev decides and the LLM reads and writes
- Build your own small System One model, with typed heads and calibration built in

For students, engineers building agents, and tech leads deciding what to build.

About the author

Sridhar Mukkandi is an applied AI engineer who builds agents and the decision systems behind them. He writes for engineers who want AI systems they can trust.
sridhar.authorhub@gmail.com

An independent guide. Not affiliated with TypeSafe AI.




Labs, figures and code
Scan for the companion repository

Artificial Intelligence / Machine Learning

github.com/Mukkandi-Sridhar/decide-dont-generate

Real threats seen by a person same six analysts




From the book's synthetic case study (Chapter 18)

Decide, Don't Generate

Sridhar Mukkandi

Decide, Don't Generate



Jev, System One Models, and the Decision Layer of Agentic AI

Runnable labs · 130+ figures · no API key needed

Sridhar Mukkandi