

Most of what an AI agent does isn't writing. It's deciding.

Is this alert real? Which team gets this ticket? Is this action safe? Most agents hand every one of those small decisions to a large language model, then dig the answer out of a paragraph. It works, slowly and expensively, and the model sounds just as sure when it's wrong.

This book shows you a better way to build the decision layer. You'll get probabilities you can trust, lines drawn from what each mistake costs, and a clear rule for when to hand a case to a person. Every step is built with runnable code, on a synthetic but realistic security team, using Jev, a new System One model, and a free mock that needs no API key.

[[PRAISE]]

Inside, you'll learn to:

- Check whether a model's probabilities mean what they say, and fix them when they don't
- Turn a probability into act, review or escalate, with lines drawn from real costs and real capacity
- Run a fair six-way bake-off: rules, logistic regression, a text classifier, two LLM methods and Jev
- Build a hybrid agent in which Jev decides and the LLM reads and writes
- Build your own small System One model, with typed heads and calibration built in

Real threats seen by a person

same six analysts

before 54%

after 89%

From the book's synthetic case study (Chapter 18)

For students, engineers building agents, and tech leads deciding what to build.

About the author

Sridhar Mukkandi [[AUTHOR BIO]]

sridhar.authorhub@gmail.com



Labs, figures and code

Scan for the companion repository

Artificial Intelligence / Machine Learning

[[PRICE INR]] · [[PRICE USD]]

github.com/Mukkandi-Sridhar/JEVBook

BARCODE AREA

2 × 1.2 in, keep clear

Decide, Don't Generate

Sridhar Mukkandi

Decide,

Don't

likely

the

ated

Gen

er

ate

act

review

escalate

Jev, System One Models, and the Decision Layer of Agentic AI

Runnable labs · 130+ figures · no API key needed

Sridhar Mukkandi