

An LLM writes its answer, one piece at a time



...then your code has to read the sentence

Jev returns a typed answer with probabilities, in one pass

