

Output so far:

```
{"verdict": "}
```

The model's next-token scores, and what the schema allows:

mal  allowed

ben  allowed

Sure  masked out

I  masked out

unk  masked out

...  masked out

Schema: verdict must be
"malicious" or "benign"

Disallowed tokens get zero chance.
The output can't break the format.
It can still be wrong.