

Structured outputs and JSON mode

1 Asking for JSON mostly works

Prompt only: 95.3% came back valid and the right shape. The rest were wrapped in chat, broken, or had extra fields.

2 Constrained decoding fixes the format

Tokens that would break the schema get zero chance. The format becomes guaranteed. The answer doesn't.

3 Validate anyway

Parse into a strict type, reject unknown fields, retry or fail safe. Treat model output as untrusted input.

4 A stated confidence is text too

The mock wrote only 6 different values, 98% of them 0.9 or more. Ranking power: AUC 0.66.

5 A token probability isn't an answer probability

It's split across spellings and starts. In the mock it ranked better (AUC 0.84) but was overconfident (ECE 0.053).

6 Measure, whatever you use

A model built to return probabilities scored ECE 0.030 on the same text. Research on real LLMs is mixed. Check yours.