

Three kinds of memory an agent can use

Working memory

the context window: what the model
can see right now

Reference memory

documents fetched on demand:
policies, runbooks, tickets (RAG)

Case memory

what happened last time: past alerts,
verdicts and analyst notes

Every write and every read is a small decision: keep this? fetch that? is it relevant enough to trust?