

RAG and memory

1 Retrieve, then generate

Find the passages that match the question and put them in the prompt. The model answers from your documents.

2 Retrieval is the weak link

If the right chunk isn't retrieved, nothing downstream can fix it. Our simple retriever put the answer first 81% of the time, in the top 3 94%.

3 Chunk size matters

Too small loses context; too big dilutes the match. Test sizes on your own questions.

4 Know when you don't know

A similarity threshold let the system abstain on 92% of questions it couldn't answer while answering 72% of those it could.

5 Context is a budget

Instructions, documents, memory and history all compete for the same window. Every token costs time and money.

6 Memory is full of decisions

What to store, what to fetch, whether it's relevant, whether to trust it: small typed decisions, at high volume.