

Where agents break

1 Inputs can lie

One planted sentence moved 50% of real threats into auto-close (from 13%). Decision models read text too.

2 Take facts from trusted systems

With the IT-tool fact taken from a trusted record instead of the text: 14%.

3 Retrieval fails quietly

66% of the agent's policy lookups returned a weak match. Check relevance before using it.

4 Demos are easy cases

Wrong on 0.5% of easy alerts, 43% of hard ones. Test on the hard ones.

5 Budgets stop runaways

Cap steps, time and spend. Make “stop” a decision with a threshold, not a hope.

6 A decision layer catches most of it

Allow-lists, typed decisions with calibrated probabilities, cost-based thresholds and human doors.