

VERIFIED

read in the public SDK

- POST /v1/systemone
- state + model + questions
- three types: choice, score, noul
- a probability for every option; default model "jev-latest"
- output tokens currently free (SDK docs)

VENDOR-REPORTED

claimed; not independently checked here

- "parallel sampler": all outputs in one pass
- trained with RLCD: reinforcement learning for calibrated decisions
- 70–500 ms latency
- \$0.042 per million input tokens
- two orders of magnitude faster than LLMs

UNKNOWN

not published

- model size and layers
- training data
- how RLCD's reward is defined
- calibration on your data
- failure modes on adversarial input