VERIFIED	VENDOR-REPORTED	UNKNOWN
read in the public SDK	*claimed; not independently checked here*	*not published*
• POST /v1/systemone	• “parallel sampler”: all outputs in one pass	• model size and layers
• state + model + questions	• trained with RLCD: reinforcement learning for calibrated decisions	• training data
• three types: choice, score, noul	• 70–500 ms latency	• how RLCD’s reward is defined
• a probability for every option; default model “jev-latest”	• $0.042 per million input tokens	• calibration on your data
• output tokens currently free (SDK docs)	• two orders of magnitude faster than LLMs	• failure modes on adversarial input