

Inside Jev: what we know and what we don't

1 Sort every claim

Verified, vendor-reported, unknown. Most of what's written about new models mixes the three.

2 The interface is public

State plus typed questions in; a probability for every option out. You can read it in the SDK.

3 The speed claim has a mechanism

All answers in one pass instead of token by token.
Vendor-reported: 70–500 ms per request.

4 The price is the story

At the listed price, a million 500-token decisions cost about 21(*vendor – reported*), *against* 660 for our illustrative LLM.

5 Check claims against each other

The Doom demo's numbers imply about 4,630 tokens per decision. Plausible, and checkable.

6 Calibration can't be taken on trust

“Calibrated” is a claim about data. Yours isn't theirs. Chapter 18 tests it.