

Testing Jev’s calibration yourself

1 Both claims can be true

A model can be calibrated where it was checked, and miscalibrated on data that differs. Calibration belongs to model and data together.

2 Same model, two companies

ECE 0.013 at Kestrel (7.6% attacks), 0.036 at Harbor (3.5%). Harbor’s risks come out about twice too high.

3 Fix the base rate first

Adjusting the odds for Harbor’s base rate cut ECE to 0.010; Platt scaling on 300 labels to 0.007.

4 Small samples lie

With 100 labels, a model whose true ECE is 0.013 typically measures about 0.043. Report a range, and get more labels.

5 Check choices too

For a choice: does the top label’s confidence match how often it’s right? Then check each label.

6 Scores until proven otherwise

Treat new model outputs as scores. Once tested and fixed on your data, treat them as probabilities.