

1 Hand-written rules

reads fields

returns yes or no

needs an expert's afternoon

2 Logistic regression

reads fields

returns a probability

needs thousands of labels

3 Trained text classifier

reads raw text

returns a probability

needs thousands of labels

4 LLM, structured output

reads raw text

returns JSON: verdict +
stated confidence

needs a prompt

5 LLM-as-judge

reads raw text

returns a 1–10 rating

needs a prompt and a rubric

6 Jev (mock)

reads fields as JSON

returns a probability

needs typed questions