1 Hand-written rules

reads	fields
returns	yes or no
needs	an expert's afternoon

2 Logistic regression

reads	fields
returns	a probability
needs	thousands of labels

3 Trained text classifier

reads	raw text
returns	a probability
needs	thousands of labels

4 LLM, structured output

reads	raw text
returns	JSON: verdict + stated confidence
needs	a prompt

5 LLM-as-judge

reads	raw text
returns	a 1–10 rating
needs	a prompt and a rubric

6 Jev (illustrative)

reads	fields as JSON
returns	a probability
needs	typed questions