

	ranking (AUC)	caught at 240/day	calibration (ECE)	cost per million	time per decision	labels needed	same answer twice?	explains itself?
Rules	0.66	50%	0.009	< \$1	1 μs	none	yes	fully
Logistic regression	0.90	87%	0.008	< \$1	10 μs	~3,000	yes	weights
Text classifier	0.75	67%	0.011	\$1	100 μs	15,000+	yes	word weights
LLM, JSON	0.71	60%	0.025	\$740	1.4 s	none	not always	writes reasons*
LLM-as-judge	0.75	69%	0.102	\$552	583 ms	none	not always	writes reasons*
Jev (mock)	0.89	86%	0.012	\$21+	70–500 ms†	none; 300 to calibrate	yes (mock)	no

Bold: best in column. † vendor-reported. * reasons written after the fact may not be how the verdict was reached (Chapter 6). Classic costs: compute only (assumption). LLM: illustrative.