


# The bake-off: six ways to make a decision

## 1 Six contestants, one fair test

Each was trained or tuned on three weeks of Kestrel alerts (14,973) and scored on the fourth (5,027), on ranking, calibration, cost, speed, stability and labels needed.



## 2 Where Jev loses

With enough labels and clean fields, logistic regression matches or beats it: AUC 0.896 against 0.892, at almost no cost per decision. Rules win on explainability. Nothing is cheaper than a rule.

## 3 Where it wins

No labels to start, typed answers that always parse, and probabilities that rank and are close to calibrated. It reads raw text better than a bag-of-words classifier.

## 4 The LLM's real problem

Stated confidences snap to a few values, so they can't rank. 3.2% of JSON answers broke. At temperature 0.7, 5.2% of verdicts changed on a second ask.

## 5 Read the design, not the winner

The mock LLM was built as a noisier reader than mock Jev. The accuracy gap between them is a design choice, not a finding. Run this contest on your own data.

## 6 The scorecard has eight columns

Pick the method that wins on the columns your decision cares about. Often that's two methods together.