


Act, review, or escalate

1 One line is not enough

A single cut at 0.5 auto-closes 37 real threats a day. Three zones let machines handle the easy ends and people handle the middle.



2 Act, review, escalate

Act: the machine closes it alone. Review: an analyst checks. Escalate: someone is paged now.



Kestrel: act below 0.0313, escalate at 0.34 and above

3 Costs set the lines

Write down what each action costs as a function of P. The cheapest line at each P is your zone.

$P \times \text{miss}$

review + misses

$(1-P) \times \text{page}$

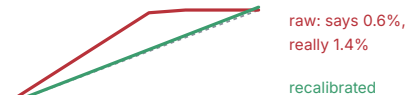
4 Capacity moves the lines

Six analysts can review 240 alerts a day, not 638. Queue size is a business decision.



5 Only honest P can be thresholded

The raw mock's “almost surely fine” alerts were 2.5× riskier than it claimed. Recalibrate on recent history before you trust any line.



6 Log every decision

Jev gives no reasons, so audit the policy and the evidence: what it saw, what it said, what you did.

{alert, P, calibrated P,
policy v3, action, time}

Watch it every week

Zone rates, queue size and calibration drift. Audit a random slice of ACT, because otherwise you never see the mistakes the machine closed on its own. Fail safe to REVIEW, never to ACT.

zone rates

queue vs capacity

calibration (ECE)

audit 3% of ACT

fail safe = REVIEW