

Build your own System One model

1 Three heads, one pass

A shared encoder reads the fields once; a sigmoid, a softmax and an ordinal head answer the three types.

2 Log loss everywhere

Each head is trained with a proper scoring rule, so honest probabilities are what training rewards.

3 And still overconfident

On unseen weeks the network was a little too sure: temperatures of about 1.2 for each head fixed it.

4 Stop before it memorises

Train too long without a penalty and held-out loss climbs. A temperature can't fully repair that.

5 Respectable, not magic

TinyJev: AUC 0.874, ECE 0.009. Logistic regression and mock Jev still rank a little better here.

6 Plug it in

A transport answers the SDK's requests with your model. Every lab in the book runs against it.