Decide, Don’t Generate

Decide, Don’t Generate

Jev, System One Models, and the Decision Layer of Agentic AI

Sridhar Mukkandi

First edition

Decide, Don’t Generate: Jev, System One Models, and the Decision Layer of Agentic AI

Copyright © 2026 Sridhar Mukkandi. All rights reserved.

No part of this book may be reproduced, stored or transmitted in any form without written permission from the author, except for brief quotations in reviews and articles. The code in this book is released under the MIT licence at github.com/Mukkandi-Sridhar/decide-dont-generate.

Independently published. First edition, 2026.

Jev is a product of TypeSafe AI. This book is independent: the author is not affiliated with TypeSafe AI, and the book is not endorsed or sponsored by it. All trademarks belong to their respective owners.

About the numbers in this book. The author did not have access to the Jev API while writing. Every Jev response, chart and table in this book was produced by `jev-mock-synthetic`, a deterministic stand-in that copies Jev’s public interface but not its behaviour. Such material is labelled “Synthetic — not measured on real Jev”. Claims made by the vendor are labelled “vendor-reported” and cited. The security data is synthetic too: Kestrel Logistics, Harbor Pharma and everyone in them are invented.

The information in this book is provided without warranty. The author assumes no responsibility for errors or omissions, or for damages resulting from its use. Nothing in this book is legal, security-compliance or medical advice.

Cover design and typesetting by the author. Text set in Source Serif 4, Inter and JetBrains Mono, all under the SIL Open Font License.

Contact: sridhar.authorhub@gmail.com

Table of contents

Preface vii

How to read this book ix

Prologue: two nights x

1 What “learning” means 1

2 Probability, and how a machine learns it 8

3 Calibration: when 0.8 really means 80% 15

4 From probabilities to actions 22

5 How modern AI models work 28

6 RAG, agents, and where they break 34

7 Fast and slow: System One models and Jev 40

8 Asking the right question: choice, score, noul 46

9 Testing a model’s calibration yourself 51

10 The bake-off: six ways to make a decision 57

11 Act, review, or escalate 63

12 Decision patterns, and a hybrid agent 70

13 One team, start to finish, and beyond 77

14 From model to service 84

15 What changes now 91

Appendices 98

Glossary 98

Cheat sheets 102

Bibliography 105

Index 108

About the Author 110

Preface

This is a book about making decisions with AI.

Every day, software makes millions of small decisions. Is this message a scam? Which team should handle this complaint? Is this payment safe? Should this AI agent take that action? More and more of these decisions are now given to large language models like ChatGPT. LLMs are excellent at reading and writing. But a decision is not a paragraph. It is a choice between a few options, made without being fully sure, where every wrong choice has a cost.

To make such decisions well, you need three things:

- a **probability you can trust**, which means what it says;
- a **threshold based on the cost of mistakes**, not on habit;
- a clear rule for **when to hand the case to a person**.

Most people building with AI today do not have all three. This book explains each of them from the beginning, in simple language, without any code.

In 2026, a company called TypeSafe AI released Jev, which it calls the first **System One model**: a model that never writes text, and only answers clear questions with probabilities, in one fast step (TypeSafe AI 2026). This book shows how to judge such a model for yourself. But the main idea is bigger than any one product: decisions deserve their own fast, testable layer in an AI system, and that layer needs its own skills.

How to trust the numbers in this book

The examples are assumed. The book follows one imagined company, Kestrel Logistics, and its security team. All numbers about this company are from an assumed example, chosen to explain the ideas clearly. Charts made this way are marked **synthetic**.

Jev's numbers are illustrations. The numbers shown for Jev are examples to explain how such a model can be tested, not measurements of the real product.

Company claims are labelled. When the book reports what TypeSafe says about Jev, such as its speed or price, it is marked *vendor-reported*, and the source is given.

Every claim goes into one of three groups: what can be checked, what the company says, and what nobody outside the company knows yet. Chapter 7 explains this habit. It is a good habit for anything you read about new AI models, including this book.

How to read this book

If you are new to machine learning, read the book in order. Chapters 1 to 4 are the foundation: what a probability means, how to check it, and how to turn it into an action. Chapters 5 and 6 give a short, simple tour of neural networks, LLMs and AI agents.

If you already work with LLMs, still read Chapters 3 and 4, because everything after them depends on these two chapters. Then continue from Chapter 7.

If you lead a team or a product, read Chapters 1 to 4, then Chapters 7, 10, 11, 13 and 15. After that, you will be able to ask your team the four questions that matter most: What does a mistake cost? Where is the threshold? Can we trust the probabilities? Who checks the warnings?

The boxes

A **Key idea** box, with a thick bar on its left, holds one sentence worth remembering. At the end of each chapter, a **Keep these** box lists all the key ideas of that chapter.

A **Where this breaks** box shows the limits of what the chapter just explained. Please read these boxes carefully; they tell you when a method can be trusted and when it cannot.

A **Set the threshold** box gives you a small decision to make, with real costs. The answer is given in brackets.

A **Going deeper** box has optional maths. You can skip it; nothing later depends on it.

Prologue: two nights

Picture a shipping company's security team. Its computers raise a warning, called an **alert**, whenever something looks suspicious: a sign-in from a strange country, an email with an odd link, a laptop talking to an unknown server. Someone has to look at each alert and decide: real attack, or nothing?

Tonight one person is on duty. Checking an alert properly takes about twelve minutes, so they can check about five an hour.

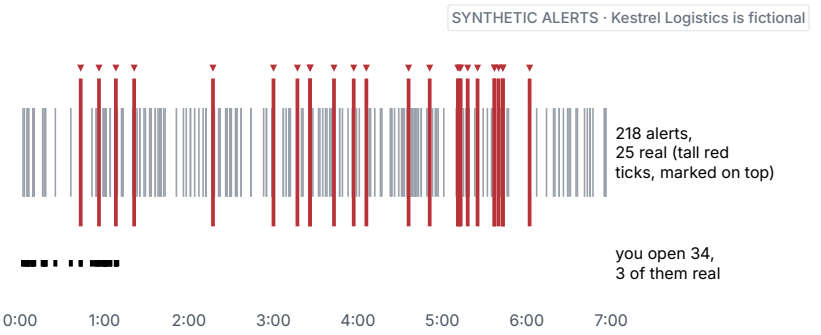


Figure 1. One night, midnight to 7 a.m. Each line is an alert; the real attacks are red, taller and marked on top. Working in the order they arrive, the person on duty checks the ones on the left. The rest wait for the morning.

Figure 1 shows the whole night. 218 alerts arrived between midnight and seven. 25 of them were real attacks. Working through them in the order they arrived, the person on duty checked 34, and only 3 of those were real.

They made no mistakes. Every alert they checked, they judged correctly. The problem was the order. The queue decided what they looked at, and a queue only knows what arrived first.

Now picture a second night.

A model reads every alert the moment it arrives. It writes nothing. It answers one question with a number: *how likely is this to be real?* And that number can be trusted: of the alerts it calls "one in ten", about one in ten really are attacks.

Alerts it's almost sure are harmless get closed, with a random few kept aside to check later. Alerts it's fairly sure are dangerous wake someone up at once. Everything in between goes to the person on duty, most likely first. The person still makes the decisions that matter, on the right alerts.

That second night is what this book builds. It needs little new technology and a lot of clear thinking: what a probability really is and how to tell if one can be trusted; what each mistake costs, and where that puts the line between “act” and “ignore”; and when a machine should decide and when it should hand the case to a person. The same thinking runs a bank’s fraud check, a hospital’s waiting list and every AI agent choosing its next step. The security team is just the example, because its mistakes are easy to count.

Next: before a machine can decide anything, it has to learn. What does that actually mean?

CHAPTER 1

What “learning” means

Anyone who starts learning AI today hears many new words at once: machine learning, deep learning, LLM, RAG, agents. It feels like seven different subjects. But all these words belong to one story. Each new idea came because the old idea had a problem it could not solve.

In this chapter we go through that story from the beginning, in simple steps. At the end, you will see where this book fits.

BY THE END OF THIS CHAPTER

- Understand the difference between a computer that follows rules and a computer that learns.
- Explain what a *model* is, in one sentence.
- See the full picture: AI, machine learning, deep learning, LLMs, RAG, agents and Jev.

Rules work, until they don’t

Many systems in our daily life follow simple rules. A water tank switch turns the motor off when the tank is full. Your bank blocks a UPI payment if it is above your daily limit. A college portal shows “absent” if attendance is below 75%. Nobody “trained” these systems. A programmer wrote the rule, and the computer follows it.

Rules are good. They are fast and cheap, and when something goes wrong you can read the rule and understand why. If a simple rule solves your problem, use the rule.

The problem starts when the situation keeps changing. Think about fake messages that try to steal passwords. This is called **phishing**. The IT team of a company writes a rule: *if an email has the word “password”, mark it as suspicious.* On Monday, the rule blocks the IT team’s own “reset your password” emails, so they add an exception. On Wednesday, the attacker writes “password” with a zero, so they add another rule. On Friday, the attacker stops using the word completely and just sends a link to a fake login page. The list of rules keeps growing, and it is always one step behind

the attacker. You may have seen the same thing with fake “KYC update” SMS messages: the wording changes every week.

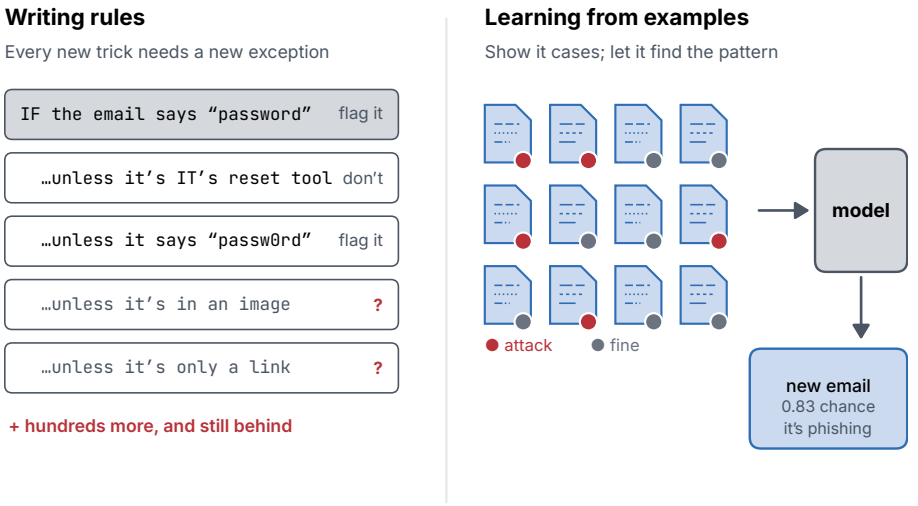


Figure 1.1. Two ways to catch phishing. On the left, every new trick forces a new exception, and the list never catches up. On the right, we stop writing rules and show the computer labelled examples instead.

So we change the approach. Instead of writing the rules ourselves, we collect thousands of old emails where we already know the answer: this one was phishing, this one was normal. Then we ask the computer to find the pattern by itself (Figure 1.1).

● KEY IDEA

Machine learning means the computer finds the rules by itself, by studying examples where we already know the correct answer.

What is a model?

After the computer studies the examples, we get something we can use on new cases. You give it a new email, and it gives back an answer. This thing is called a **model**.

You can imagine a model as a machine with many small knobs. During learning, the computer keeps turning the knobs until the machine gives the right answer for most of the examples. After that, we stop turning the knobs and use the machine on new emails, where nobody knows the answer yet.

● KEY IDEA

A model is a function whose settings come from examples, not from a programmer.

You already use many models. When your bank sends an SMS like *"Did you just make a payment of ₹40,000 in another city? Reply NO if not you,"* no person wrote a rule about you. A model learned your normal spending pattern, and this payment did not match it.

A model can learn the wrong thing

To make the ideas concrete, this book uses one example again and again. Let us imagine a logistics company called Kestrel Logistics. Like any company, its computers raise warnings when something looks suspicious, for example a login from a strange country. These warnings are called **alerts**. Suppose that in four weeks the security team receives 20,000 alerts, about 714 per day, and only 7.6% of them are real attacks; the rest are false alarms. Since this is an example, we can assume we know the correct answer for every alert, which lets us check every idea properly.

Each alert comes with a *threat score* between 0 and 1. A higher score means outside security sources think it is more dangerous. Let us try a simple rule: *if the score is above some number t, mark the alert as an attack.* A cut-off number like t is called a **threshold**.

Which value of t should we choose? We let the computer try every value from 0 to 1, count the mistakes for each value, and keep the value with the fewest mistakes. This is learning in its smallest form.

The computer chooses 0.90. With this threshold it makes 1,487 mistakes. But look at this: a rule that *never* marks anything as an attack makes 1,515 mistakes, almost the same number. Out of 1,515 real attacks, the "best" rule catches only 73 (Figure 1.2).

Why did this happen? There are about twelve harmless alerts for every real attack. So the easiest way to reduce the number of mistakes is to almost never raise an alarm. The computer did exactly what we asked. Our question was wrong.

● KEY IDEA

A model learns only what we ask it to learn. If we ask the wrong question, it learns the wrong thing.

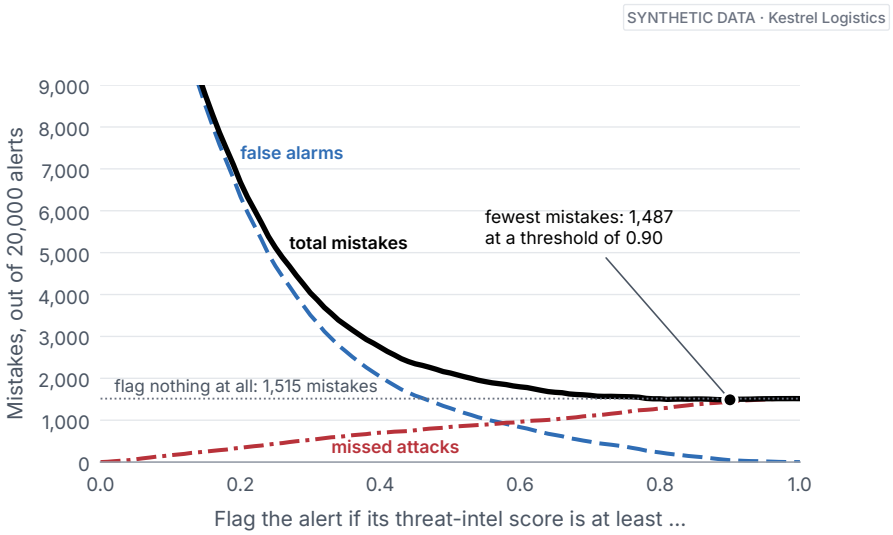


Figure 1.2. Every possible threshold and the mistakes it makes. Low thresholds bury the team in false alarms; high ones miss attacks. The total is lowest at the far right, just below the dotted line for “flag nothing at all”.

We made two mistakes in the question. First, we counted every mistake as equal. But missing a real attack is much more costly than one false alarm. Chapter 4 fixes this. Second, we forced the answer to be only *yes* or *no*.

A better answer: how likely?

Let us ask a different question: *among alerts with a similar score, how many were real attacks?*

Figure 1.3 shows the answer. There is no clear line where alerts suddenly become dangerous. Even in the lowest-score group, 2% were attacks. That looks small, but the group has 7,270 alerts, so more than a hundred attacks are hiding there. And in the highest-score group, only 62% were attacks.

So for an alert with a score of 0.55, the honest answer is not “attack” or “safe”. The honest answer is: *about one in five alerts like this is an attack.*

This kind of answer is much more useful. It tells the team how worried to be, which alerts to check first, and, once we know the cost of each mistake, where to draw the line.

● KEY IDEA

A good model gives an answer and also says how sure it is. And that “how sure” must be correct.

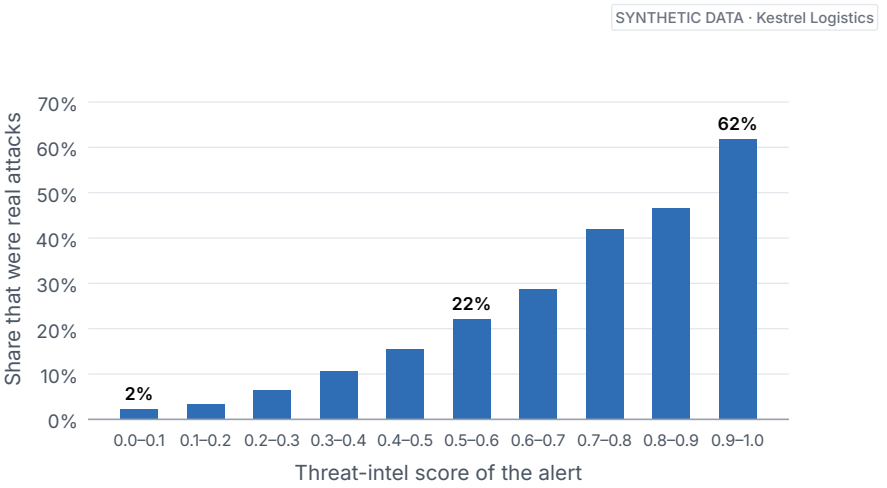


Figure 1.3. Group the alerts by threat score and count how many were real attacks. The share climbs steadily from 2% to 62%. Even at the top, plenty are harmless; even at the bottom, a few are attacks.

When a model's "how sure" is correct, we say the model is **calibrated**. You will see this word many times in this book, and Chapter 3 explains it fully.

The bigger story, and where this book fits

The threat score was one simple number. But a photo has lakhs of pixels, and an email is a long sentence. Nobody can write rules for them by hand. So researchers built models with many layers that learn their own clues. This is **deep learning**. Then they trained very large deep-learning models on huge amounts of text, with one task: guess the next word. These models learned language itself. They are called **large language models**, or LLMs, like ChatGPT. An LLM knows a lot about the world but nothing about your company, so we give it your documents to read before it answers. This is called **RAG**. Finally, we let the model use tools and take actions step by step. This is an **agent**. Chapters 5 and 6 explain all of these.

When you watch an agent work, you notice something interesting. Most of its steps are not writing. They are small decisions with only a few possible answers: *Is this alert real? Which team should handle this ticket? Is this action safe?* Today, people usually send each such question to an LLM. The LLM writes a full paragraph, and then more code has to find the actual answer inside that paragraph. It works, but it is slow and costly for such a small job.

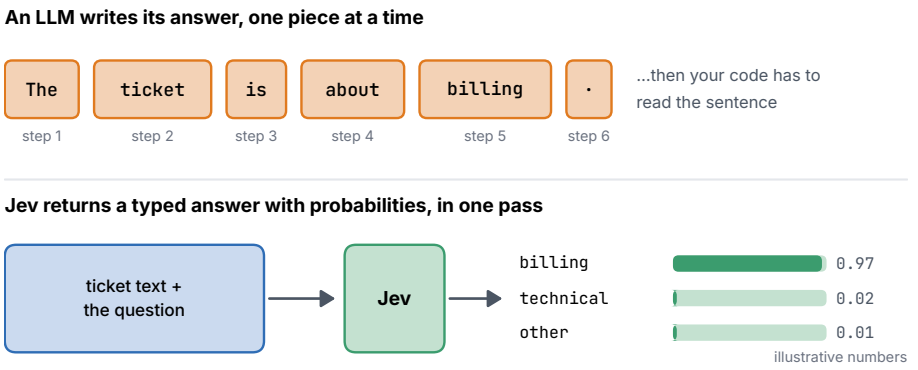


Figure 1.4. The same question, two ways. An LLM writes its answer piece by piece, and your code has to read the sentence. A decision model returns the answer as data, with a probability for every option, in one pass.

In 2026, a company called TypeSafe AI released a model named **Jev** for exactly these small decisions. Jev does not write any text. You give it the situation and a clear question, and it gives back the answer with a probability for each option (Figure 1.4). TypeSafe calls it a *System One* model. The name comes from the psychologist Daniel Kahneman, who called our fast, automatic thinking “System 1” (Kahneman 2011). In this book, the numbers shown for Jev are examples to explain the ideas, not measurements of the real product. They are marked *synthetic*.

Figure 1.5 shows the whole story on one page. The next three chapters stay at the base of this map, because everything above it depends on one idea: an answer should come with a probability, and that probability should be correct.

KEEP THESE

- **Machine learning** means the computer finds the rules by itself, by studying examples where we already know the correct answer.
- A **model** is a function whose settings come from examples, not from a programmer.
- A model learns only what we ask it to learn. If we ask the wrong question, it learns the wrong thing.
- A good model gives an answer and also says how sure it is. And that “how sure” must be correct.

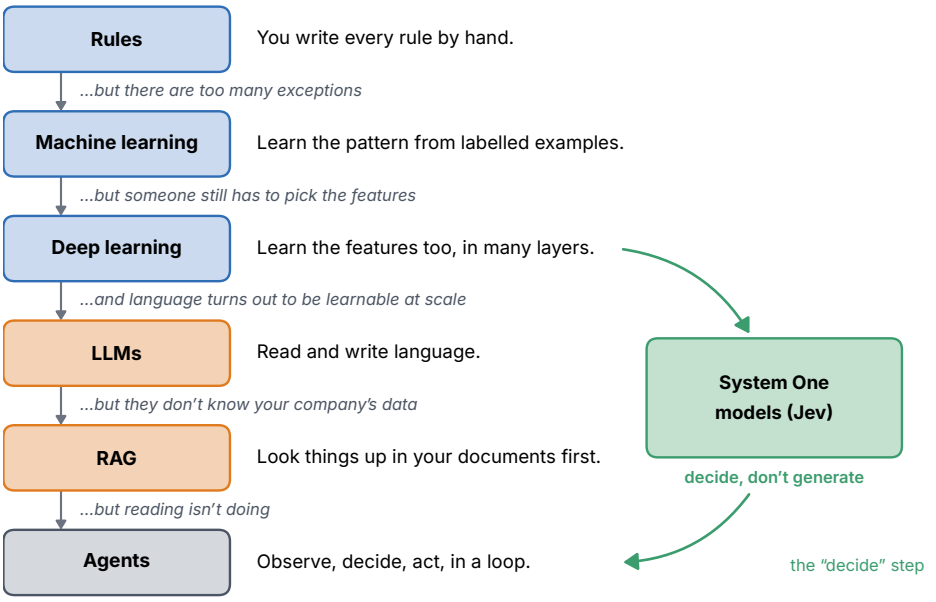


Figure 1.5. The story so far. Each step exists because the one before it hit a problem it couldn't solve. System One models like Jev grow out of deep learning and fit into the "decide" step of an agent.

EXERCISES

1. Name one system you use that follows fixed rules, and one that has learned from examples. How did you know?

Next chapter: what does "70% chance" really mean? There is a common trap here that even doctors fall into.

CHAPTER 2

Probability, and how a machine learns it

Many people think probability is only about coins and dice, something from school maths that has nothing to do with building AI. This is a costly mistake. In this chapter we learn to read a probability correctly, see a common trap that even doctors fall into, and see how a machine learns probabilities from examples.

BY THE END OF THIS CHAPTER

- Understand what "70% chance" really means, in a way you can check.
- Avoid the base-rate trap, and update a probability when new evidence comes.
- See how a machine learns a probability, and two ways it can fool you.

What "70% chance" means

Suppose the weather app says there is a 70% chance of rain tomorrow, and it does not rain. Was the app wrong?

One day cannot tell us. We need to look at many days. Take all the days when the app said "70%". If it rained on about 70 out of every 100 such days, the app is doing its job. Tomorrow was simply one of the other 30 days.

● KEY IDEA

A probability tells us how often something happens among many similar cases.

This is useful because we can check it. When a model says "this alert is 70% likely to be an attack", it is making the same kind of promise as the weather app, and we can check it in the same way.

The trap that fools doctors

Try this question before reading further. A test catches 99% of real attacks. It also gives a false alarm on 5% of normal activity. Real attacks are rare: about 1 in every 1,000 events. The test now raises an alarm. What is the chance that it is a real attack?

Most people answer about 95%. The correct answer is about 2%. In a famous study from 1978, doctors at a top American medical college were asked the same question about a medical test, and the most common answer was also 95% (Casscells et al. 1978).

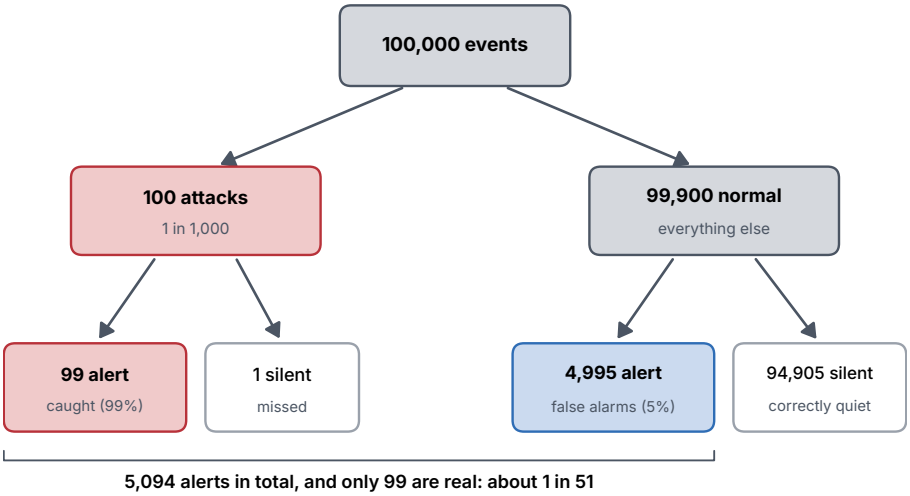


Figure 2.1. The same puzzle, counted out. Of 100,000 events, 100 are attacks and the detector catches 99 of them. But 5% of the 99,900 normal events also set it off: 4,995 false alarms. When the alarm sounds, you're holding one of 5,094 alerts, and only 99 of them are real.

The easy way to see this is to count, as in Figure 2.1. Take 1,00,000 events. Only 100 are attacks, and the test catches 99 of them. But 5% of the 99,900 normal events also raise an alarm, which is 4,995 false alarms. So out of 5,094 alarms, only 99 are real. That is about 2%.

The chance before we see any evidence, here 1 in 1,000, is called the **base rate**. Forgetting the base rate is one of the most common mistakes with probability. This is why a fraud model that is "95% accurate" can still send a bank lakhs of false alarms. Whenever a probability surprises you, convert it into counts like this (Gigerenzer and Hoffrage 1995).

● KEY IDEA

When something is rare, it stays fairly unlikely even after an alarm.

There is a related mistake. “How often do attacks have a bad score?” and “How often is an alert with a bad score an attack?” are two different questions. In our example, only 27% of attacks come with a bad threat score, but an alert with a bad score is an attack 46% of the time.

Updating with new evidence

How should a new clue change our belief? It is easiest with **odds**. A probability of 20% means odds of 20 to 80, or 0.25. Each clue has a strength called the **likelihood ratio**: how many times more common the clue is among attacks than among harmless alerts. In our example, a bad threat score is about 10.2 times more common among attacks.

● KEY IDEA

New odds = old odds × strength of the clue. This is Bayes’ rule.

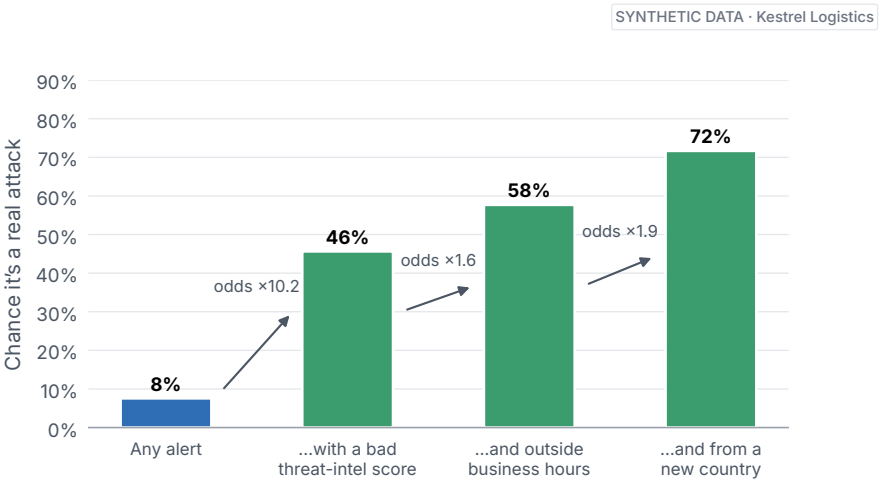


Figure 2.2. Starting from the base rate and multiplying the odds by each clue’s likelihood ratio. Three ordinary clues take an alert from “probably nothing” to “probably an attack”.

An ordinary alert starts at 8%. With a bad threat score, it goes up to 46%. If it also happened outside office hours, it becomes 58%. If the login was also from a new country, it becomes 72% (Figure 2.2). One caution: this assumes each clue gives new information. In real life, clues often overlap, so in practice we let a model learn how much each clue is worth.

Why accuracy can mislead

Suppose our rule is: *mark an alert as an attack if its threat score is 0.5 or more.* Three common measures:

- **Precision**: when the rule raises an alarm, how often is it a real attack? Here, 34%.
- **Recall**: out of all real attacks, how many does the rule catch? Here, 45%.
- **Accuracy**: out of all alerts, how many did the rule judge correctly? Here, 89%.

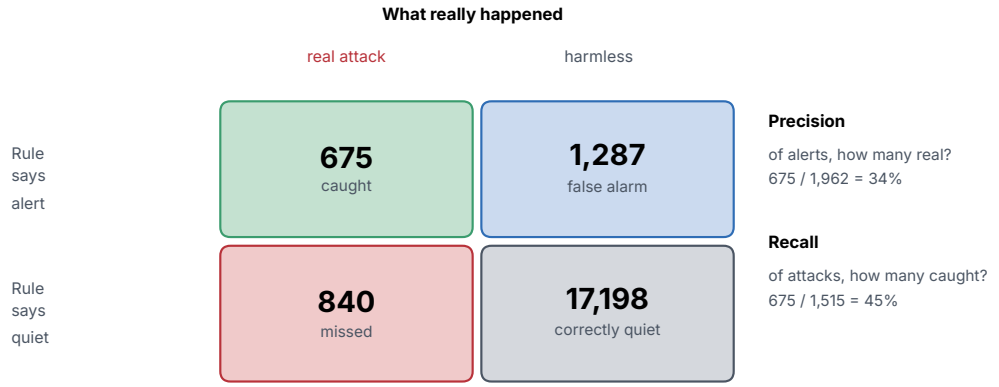


Figure 2.3. Every one of the 20,000 alerts, sorted by what the rule said and what really happened. Precision and recall each ignore one of the two columns, which makes them fair judges of rare events.

The accuracy of 89% sounds good. But a rule that *never* raises any alarm has an accuracy of 92%. When one outcome is rare, accuracy mostly measures how common the other outcome is. So for rare events, look at precision and recall.

How a machine learns a probability

Bayes' rule multiplies odds. Adding is easier than multiplying, so we take the logarithm of the odds, and then each clue simply adds a number. This gives a very simple model: start with a base number, add a **weight** for each clue, and convert the total into a probability using an S-shaped curve. This model is called **logistic regression**. Banks and hospitals have used it for decades for credit scores and risk scores.

To find good weights, we need a way to score a probability. If the model says 90% and the event happens, it should lose very few points. If it says

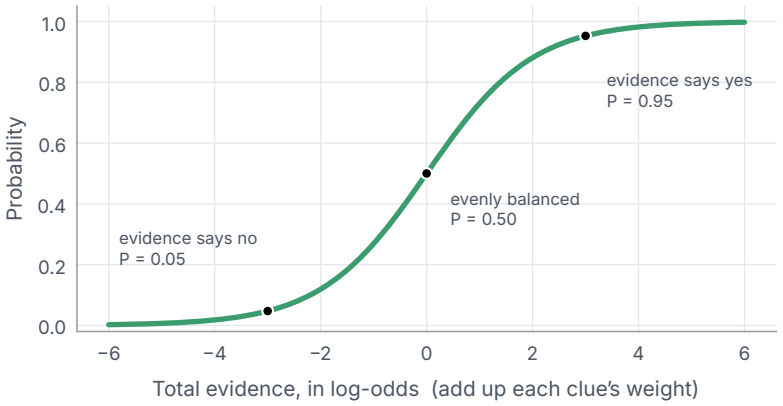


Figure 2.4. The S-curve (the logistic, or sigmoid, function) turns total evidence into a probability. Every step to the right multiplies the odds by the same amount, which is why the curve flattens out: once you're very sure, more evidence changes the probability only a little.

1% and the event happens, it should lose many points, because it was confident and wrong. This score is called **log loss**.

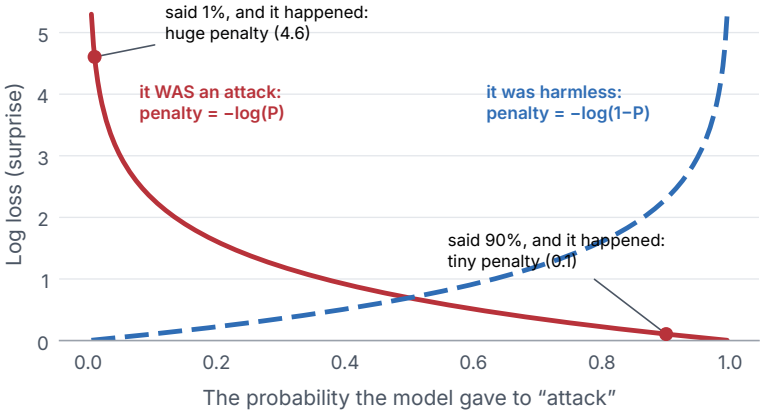


Figure 2.5. Log loss for a single alert. If it really was an attack (red, solid curve), the penalty is small when the model gave "attack" a high probability. It grows very fast as that probability gets close to zero. Harmless alerts (blue, dashed curve) work the same way from the other side.

Log loss has an important property. The only way to get the best score, on average, is to give honest probabilities (Figure 2.5). If a model exaggerates and says 99% when the truth is 80%, it loses more over many cases. So models trained with log loss tend to give honest probabilities.

Training means finding the weights with the lowest log loss. Imagine standing on a hill in thick fog. You cannot see the bottom, but you can feel which way the ground slopes. So you take a small step downhill, and repeat. This method is called **gradient descent**.

● KEY IDEA

Machine learning is one simple loop: predict, measure the mistake, adjust a little, and repeat.

Even large language models are trained with this same loop, only with a much bigger model and much more data.

Two ways a model fools you

Overfitting means memorising instead of learning. A very flexible model can remember every small detail of its training data, including random noise. It then looks excellent on that data but performs badly on new cases. The only way to catch this is to test the model on data it has never seen.

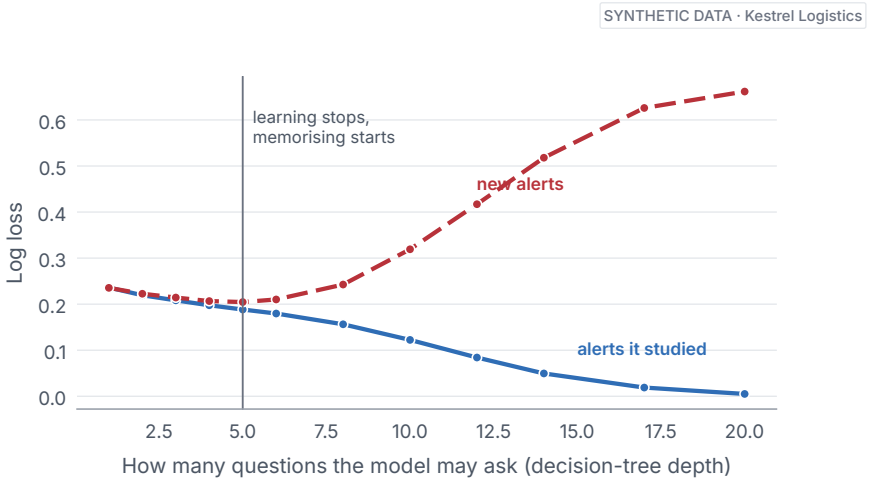


Figure 2.6. A decision tree allowed only a few questions learns general patterns. Allowed dozens, it memorises individual training alerts: its loss on them approaches zero while its loss on new alerts climbs.

Leakage means a clue secretly contains the answer. For example, if we train with the analysts' notes, results look amazing, because notes like "confirmed phishing" are written *after* someone has already decided. When the model has to decide, that note does not exist yet. So for every clue, ask: *will I really have this information at the moment of the decision?*

KEEP THESE

- A **probability** tells us how often something happens among many similar cases.
- When something is rare, it stays fairly unlikely even after an alarm.
- New odds = old odds × strength of the clue. This is Bayes' rule.
- Machine learning is one simple loop: predict, measure the mistake, adjust a little, and repeat.

EXERCISES

1. A spam filter catches 98% of spam and wrongly flags 1% of normal emails. If only 1% of all emails are spam, what fraction of flagged emails are actually normal? Use counts, as in Figure 2.1.

Next chapter: if a model says 80%, does it really mean 80%? We learn how to check, and how to fix it.

CHAPTER 3

Calibration: when 0.8 really means 80%

If a model says 80%, does it really mean 80%? This sounds like a silly question, but it can cost a lot of money. Nothing inside a model forces its 80% to match reality. It matches only if the model was built carefully and then checked. Checking this is the most important habit in this book.

BY THE END OF THIS CHAPTER

- Understand what *calibrated* means, and why it is different from accurate.
- Check calibration using a reliability diagram and one simple number.
- Fix a model whose probabilities are wrong, and check it group by group.

What weather forecasters do well

In the 1970s, researchers in the United States compared weather forecasts with what actually happened. When forecasters said "30% chance of rain", it rained on about 30% of those days. When they said "70%", it rained on about 70% (Murphy and Winkler 1977).

Most people are not like this. If you ask someone for a range they are "90% sure" about, the true answer falls inside that range much less than 90% of the time. People are usually too confident.

Good forecasters are not always right; nobody can be. They are good because their numbers mean what they say (Figure 3.1). This is called **calibration**.

● KEY IDEA

A model is calibrated when its 80% comes true about 80% of the time.

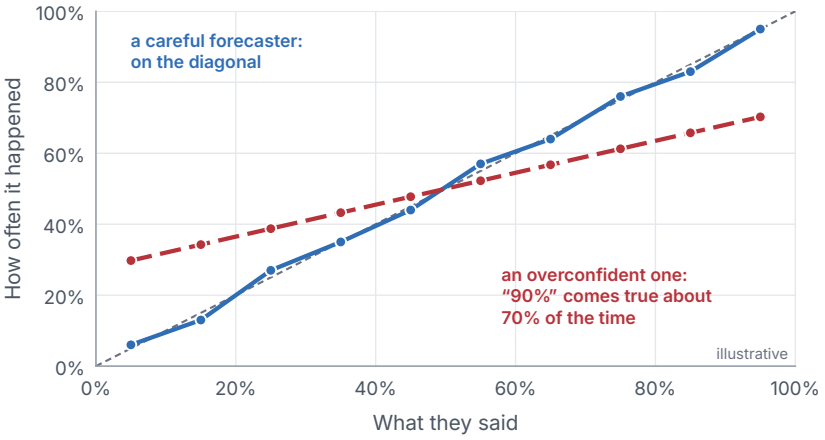


Figure 3.1. Two forecasters. For each thing they said ("30%", "70%" and so on), we count how often it actually happened. The careful forecaster sits on the diagonal: their 70% happens about 70% of the time. The overconfident one says 90% about things that happen only about 70% of the time.

Ranking and calibration are two different things

A model's numbers can be good at two separate jobs.

Ranking: does the model give higher numbers to attacks than to harmless alerts? If we sort alerts by the model's number, do real attacks come to the top? This is measured by **AUC**. It is the chance that a random attack gets a higher number than a random harmless alert. An AUC of 1 is perfect; 0.5 is no better than tossing a coin.

Calibration: when the model says 0.3, are about three out of ten such alerts real attacks?

In Figure 3.2, both models have the same AUC of 0.888. But model B says 40% for alerts that are attacks only a few times in a hundred. If you only sort alerts, B is as good as A. But the moment you use its numbers to make a decision, B will mislead you. Most leaderboards report only ranking. For decisions, calibration usually matters more.

● KEY IDEA

Ranking well and being calibrated are different skills. Good decisions need both.

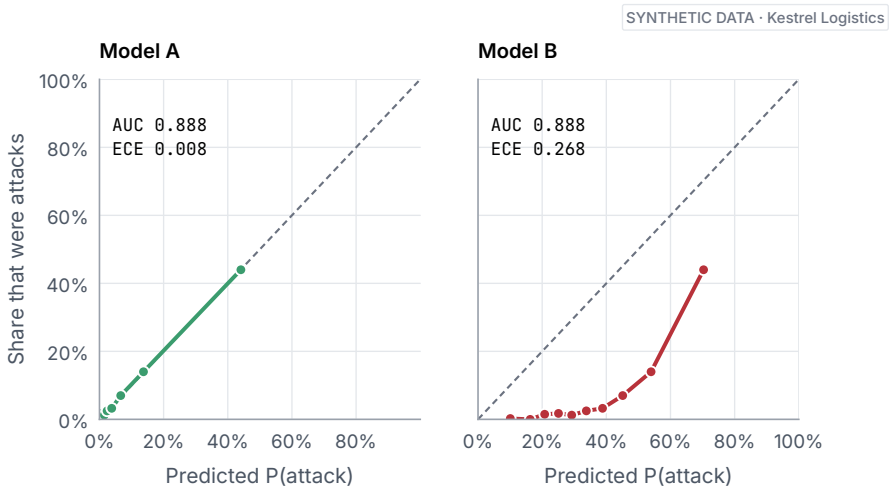


Figure 3.2. Two models, identical at ranking (same AUC) but not at calibration. Model A is the logistic regression from Chapter 2. Model B is the same model with its probabilities distorted: it still sorts the alerts in the same order, but its numbers no longer mean anything.

How to check calibration

Do what the weather researchers did. Take predictions where you already know the outcome. Put them in groups by what the model said: 0–10%, 10–20%, and so on. In each group, compare the average prediction with how often the event really happened. When we draw this, it is called a **reliability diagram.** If the points lie on the diagonal line, the numbers can be trusted.

Always look at how many cases are in each group, shown at the bottom of Figure 3.3. A point based on only twelve alerts is weak evidence.

For a single number, we use the **expected calibration error**, or **ECE.** It is the average gap between "what the model said" and "what actually happened", across all groups (Naeini et al. 2015). In our example, the model's ECE is 0.008, which means it is off by less than one percentage point on average. Model B's ECE is 0.27. But remember that ECE is an average, so it can hide one bad area. Always look at the diagram too.

Why calibration goes wrong, and how to fix it

Rebalancing the training data. This is a very common cause. Attacks are rare, only 8% in our example. So someone trains the model on equal numbers of attacks and normal alerts, "so the model sees more attacks". The ranking stays fine, but the model now believes attacks are common. Its

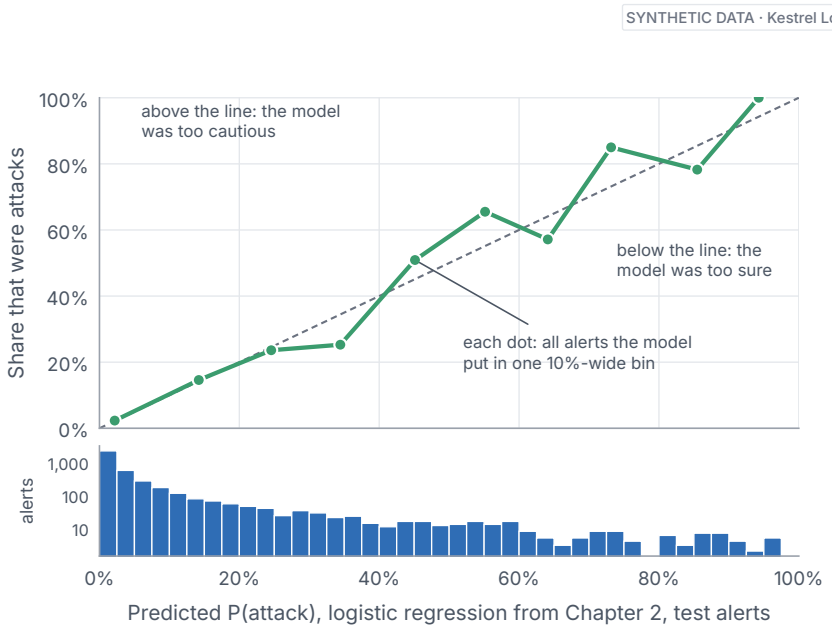


Figure 3.3. The reliability diagram of our Chapter 2 model on alerts it never saw. Points on the dashed diagonal mean the model's numbers can be trusted as they are. Underneath, a histogram (log scale) shows how many alerts fall in each part of the range: most are near zero, and only a few are above 50%.

average prediction becomes 29%. If we use these numbers for decisions, the team gets far too many alerts.

Overfitting pushes probabilities towards 0 and 1. **Very large neural networks** also tend to be overconfident (Guo et al. 2017). And **changes in the world**, like a sudden increase in attacks, can make last month's calibrated model wrong today.

The good news is that a model which ranks well but is badly calibrated is usually easy to fix. We do not need to train it again. We keep aside some data the model has never seen, called a **calibration set**, and learn a small correction on top of the model:

- **Platt scaling** learns two numbers that stretch and shift the model's output (Platt 1999). This is usually the first choice.
- **Temperature scaling** learns one number that softens overconfident answers. It is the standard fix for neural networks.
- **Isotonic regression** learns any rising staircase shape (Zadrozny and Elkan 2002). It is flexible but needs more data.

For the rebalanced model, Platt scaling brings the ECE down from 0.216 to 0.010. Temperature scaling does not help at all; the ECE stays at 0.215. The

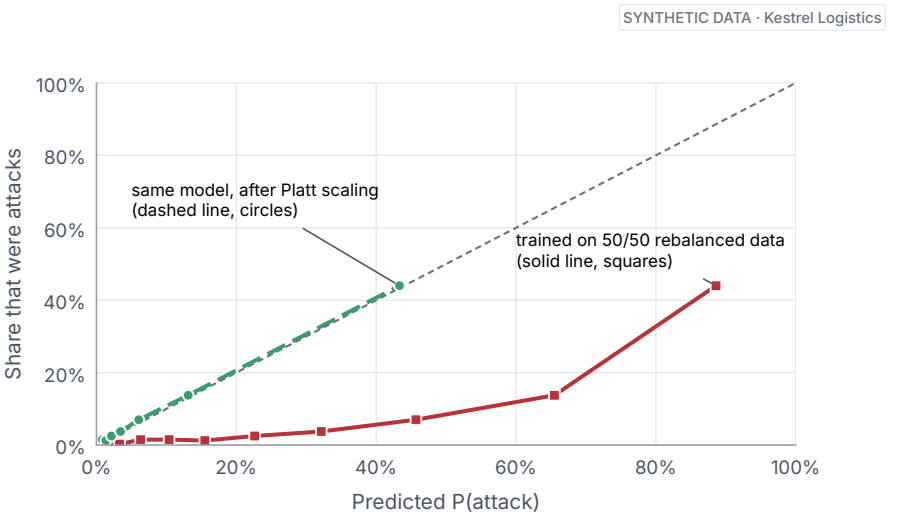


Figure 3.4. A model trained on 50/50 rebalanced data (red, solid line with squares) thinks attacks are about ten times more common than they are, and its probabilities sit far below the diagonal. After Platt scaling fitted on a separate calibration set (green, dashed line with circles), the same model is well calibrated again.

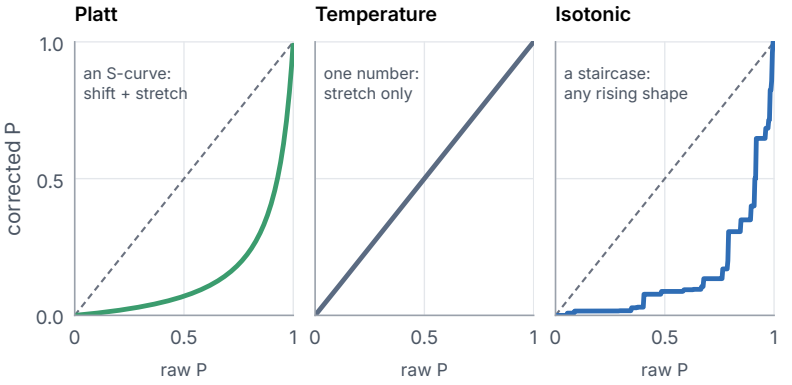


Figure 3.5. The three corrections, each fitted to the rebalanced model on the calibration set. Platt scaling learns an S-curve that can shift and stretch. Temperature scaling learns one number that only stretches around the middle, and here it barely changes anything. Isotonic regression learns any rising staircase.

reason is useful to know. Rebalancing did not make the model overconfident. It *shifted* all its predictions upwards. Temperature can only stretch or squeeze around the middle; it cannot shift everything down.

● KEY IDEA

Temperature scaling fixes overconfidence. It cannot fix a wrong base rate.

Correct on average, wrong for one group

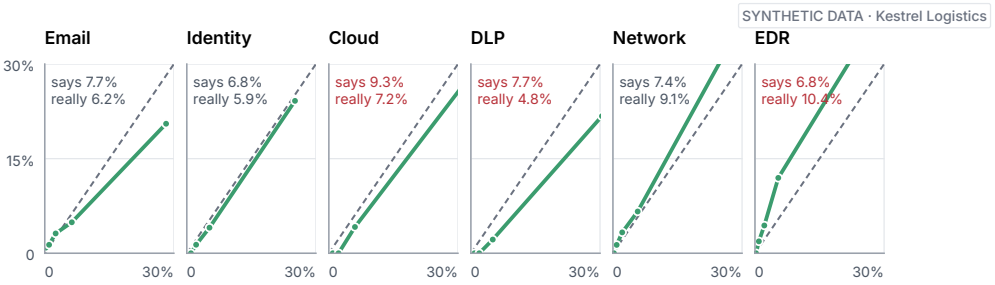


Figure 3.6. Calibration checked separately for each alert source. Across all sources together, this model is well calibrated, with an ECE of about half a percentage point. But it overestimates the risk of DLP and Cloud alerts and underestimates EDR alerts, which really are more dangerous here. Averaging hides both mistakes.

A model can be well calibrated overall but wrong for one group. Suppose a model cannot tell which system raised each alert. Its overall ECE is an excellent 0.005. But for alerts from laptops and servers (EDR in Figure 3.6), it says 6.8% when the truth is 10.4%. For data-leak alerts (DLP), it says 7.7% when the truth is 4.8%. The errors cancel out in the average, and real attacks on laptops get closed without anyone looking.

● KEY IDEA

Check calibration separately for every group you treat differently, not only on average.

These groups can be sources, customer types, languages and regions. When a model affects people's lives, they must also include groups of people, such as age or gender.

⚠ WHERE THIS BREAKS

Calibration is true only for data similar to the data you checked it on, and it slowly fades as things change. Also, calibration says nothing about one single prediction: a perfectly calibrated 70% is still wrong three times out of ten. Keep fresh data for checking, and check again regularly.

● KEEP THESE

- A model is **calibrated** when its 80% comes true about 80% of the time.
- Ranking well and being calibrated are different skills. Good decisions need both.
- Temperature scaling fixes overconfidence. It cannot fix a wrong base rate.
- Check calibration separately for every group you treat differently, not only on average.

EXERCISES

1. Explain to a friend, without technical words, the difference between a model that ranks well and a model that is calibrated.

Next chapter: now that we have probabilities we can trust, where should we draw the line between “act” and “ignore”?

CHAPTER 4

From probabilities to actions

Now we have probabilities we can trust. The next question is: *where should we draw the line?* Most people choose 0.5, because it is the default in almost every software library and it feels neutral. But 0.5 is not neutral. It hides a strong assumption about your business, and usually that assumption is wrong.

BY THE END OF THIS CHAPTER

- Turn the cost of two kinds of mistakes into one threshold.
- Understand why 0.5 is only a special case.
- Treat "ask a person" as a real choice, and meet the three doors: act, review, escalate.

The umbrella question

During the monsoon, the forecast says there is a 30% chance of rain. Should you carry an umbrella?

You already know the answer depends on two costs. Carrying an umbrella you did not need is a small trouble; call it 1. Getting fully wet on the way to an important meeting is much worse; call it 4. If you carry the umbrella, you always pay 1. If you do not carry it, you pay 4, but only when it rains. So on average you pay 4 × the chance of rain.

The two costs are equal when the chance of rain is 20% (Figure 4.1). Above 20%, carry the umbrella. So at 30%, you carry it, even though it will probably not rain.

One simple formula

Any yes-or-no decision can go wrong in two ways:

- a **false alarm**: you acted, but there was no need;
- a **miss**: you did not act, but you should have.

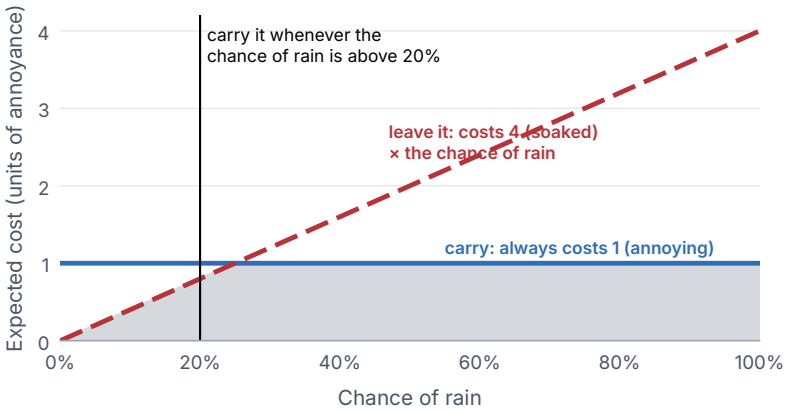


Figure 4.1. The umbrella decision. Carrying always costs 1. Leaving it costs 4 times the chance of rain. The lines cross at 20%: above that, carry it.

● KEY IDEA

Act when the probability is more than: cost of a false alarm ÷ (cost of a false alarm + cost of a miss).

For the umbrella: 1 ÷ (1 + 4) = 0.2. Now see what happens when both costs are equal: the formula gives 1 ÷ 2 = 0.5. So 0.5 is correct only when a false alarm and a miss cost exactly the same. That is almost never true in real problems.

● KEY IDEA

The 0.5 threshold is correct only when both kinds of mistakes cost the same.

Different costs give different thresholds (Figure 4.2). For example, if we let a security alert close automatically, a missed attack may cost about 630 times more than a quick check by a person. Then the threshold falls to about 0.0016.

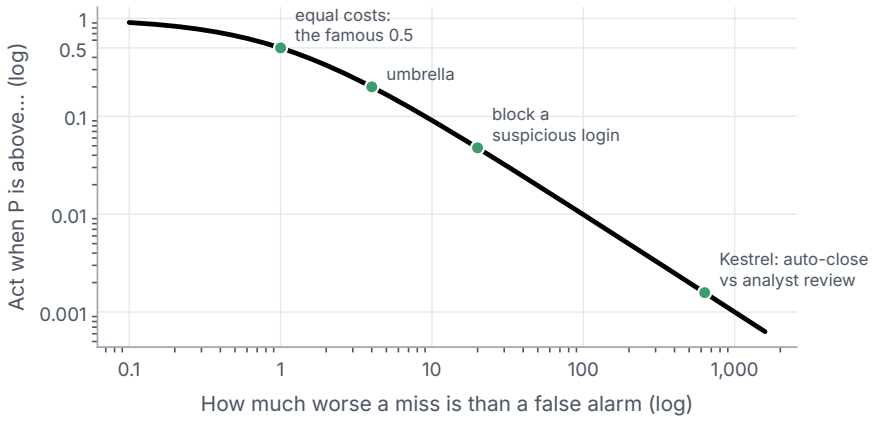


Figure 4.2. The right threshold for any ratio between the cost of a miss and the cost of a false alarm (both axes on a log scale). Equal costs give 0.5. As misses get more expensive, the threshold drops fast.

GOING DEEPER optional

Where the formula comes from. If you act, you pay the false-alarm cost only when the case was harmless, so on average $(1 - P) \times C_{fa}$. If you do not act, you pay the miss cost only when it was real, so $P \times C_{miss}$. Act when the first is smaller, and rearrange: $P > C_{fa} / (C_{fa} + C_{miss})$ (Elkan 2001).

Does it work in practice?

In our example company, the team must decide for each alert: block it or allow it? Suppose wrongly blocking a harmless activity costs about $200 of lost work, and letting a real attack through costs about $4,000. The formula says: block when the probability is above 0.048. If we try all possible thresholds one by one, the best one is almost exactly the same. Using 0.5 instead costs 2.3 times as much.

This is why calibration comes first. The rebalanced model from Chapter 3 gives numbers that are too high. If we use the formula with its numbers, the cost is 62% higher (Figure 4.3). Its best threshold is somewhere else, and we could only find it by trial and error, again and again whenever the costs change.

KEY IDEA

With calibrated probabilities, the costs directly tell us the threshold. Without calibration, we are only guessing.

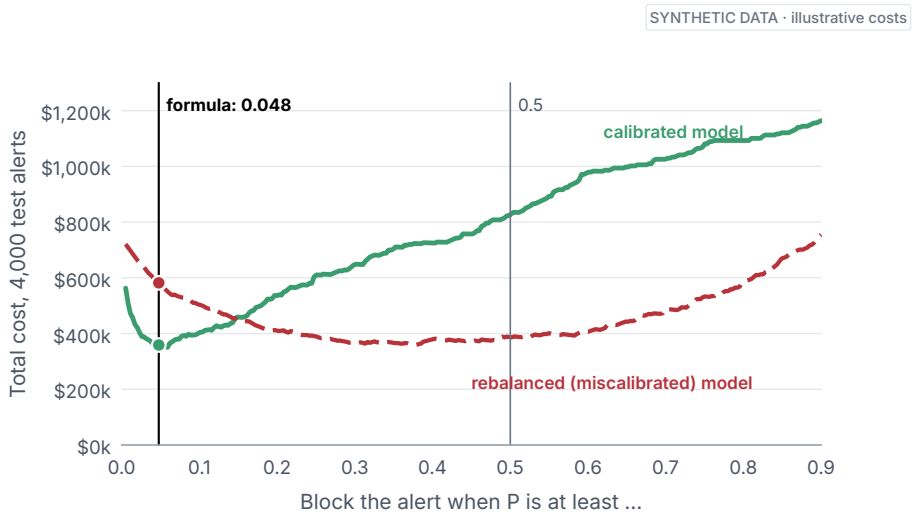


Figure 4.3. Total cost on the test alerts at every threshold. For the calibrated model (green, solid line), the formula's threshold is near the lowest point of the curve. For the same model trained on rebalanced data (red, dashed line), the same threshold lands far from its lowest point.

"I am not sure, ask a person" is also a choice

Usually there is a third option, and it is often the best one: do not decide, and send the case to a person. In 1970, the engineer C. K. Chow showed that a machine which is allowed to skip its least certain cases makes far fewer mistakes (Chow 1970).

If the model decides every alert, 6.1% of its decisions are wrong. If it decides only the 80% of alerts it is most sure about, the error falls to 2.2%. If it decides only the most certain half, the error is just 0.9% (Figure 4.4). This chart is very useful for managers. It answers their real question: *how much of this work can we safely give to the machine, and with how many mistakes?*

● KEY IDEA

If a model decides only the cases it is most sure about, its mistakes drop sharply.

Three doors

When we combine cost-based thresholds with "ask a person", we get the design used in the rest of this book.



Figure 4.4. Coverage against risk. Reading from right to left: if the machine decides every alert, you take its full error rate. Let it decide only the alerts it's surest about and send the rest to people, and the error rate among its decisions falls steeply. A better model (green, solid line) gives you more coverage at the same risk.

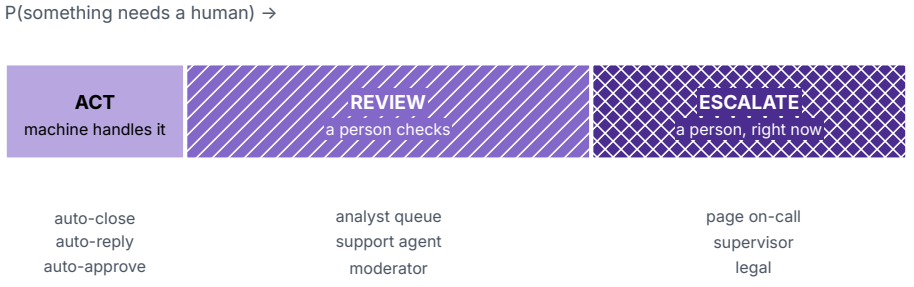


Figure 4.5. The three doors. At one end the model is confident nothing's needed, so the machine handles it. At the other end it's confident something's badly wrong, so a person is called in immediately. In between, where it's unsure, a person looks when they can.

- **Act**: the machine handles the case alone.
- **Review**: a person checks the case during normal work.
- **Escalate**: a person is called immediately.

The lines between these doors also come from costs. Chapter 11 shows how to set them, and what to do when there are not enough people for the review door.

⚠ WHERE THIS BREAKS

The costs we use are estimates, but the formula treats them as exact. Two habits help. First, try each cost at half and at double its value. If the threshold hardly moves, do not waste time arguing about the exact number. Second, remember that costs change from case to case: a false alarm on the CEO's laptop costs more than one on a test computer. And some costs are not money at all, such as trust, fairness and safety. Discuss them openly.

SET THE THRESHOLD

The sales team says a wrongly blocked customer email costs $800, not $200. Where should the blocking threshold go now?
(800 ÷ 4,800 ≈ 0.17. Fewer alerts will be blocked, and a few more attacks will get through. The model did not change; the business priorities changed, and the threshold should follow them.)

KEEP THESE

- Act when the probability is more than: cost of a false alarm ÷ (cost of a false alarm + cost of a miss).
- The 0.5 threshold is correct only when both kinds of mistakes cost the same.
- With calibrated probabilities, the costs directly tell us the threshold. Without calibration, we are only guessing.
- If a model decides only the cases it is most sure about, its mistakes drop sharply.

EXERCISES

1. A hospital test checks for a serious but treatable disease. Treating a healthy person by mistake costs ₹50,000. Missing a sick person costs ₹50,00,000. At what probability should the doctor start treatment?

Next chapter: how modern AI models work inside, from one artificial neuron to ChatGPT, and why this matters for decisions.

CHAPTER 5

How modern AI models work

Words like “deep learning” and “large language model” sound very complex. But they are built from four simple ideas. In this chapter we go through these four ideas, and then see why asking an LLM like ChatGPT to make a decision is not as simple as it looks.

BY THE END OF THIS CHAPTER

- Understand layers, embeddings and attention in simple words.
- See why big models often become overconfident.
- Understand why getting a clean decision out of an LLM is difficult.

Idea one: layers can draw any shape

One artificial **neuron** is the same as the logistic regression from Chapter 2: it multiplies each clue by a weight, adds them up, and turns the total into a probability. But one neuron can only separate cases with a straight line.

Some problems need more. For example, activity can be suspicious when it is unusual in *any* direction: far too much or far too little. A straight line cannot separate the normal middle from the unusual edges.

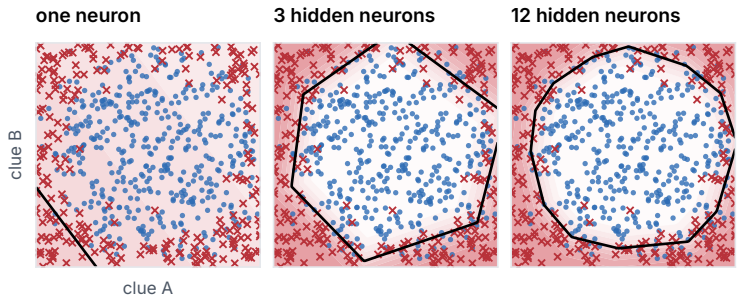


Figure 5.1. A toy problem. Blue dots are normal activity; red crosses surround them on every side. One neuron (left) can only draw a straight line through it. With a hidden layer of three neurons (middle) the boundary becomes a triangle; with twelve (right), close to a circle.

The solution is to pass the clues through a layer of neurons first. Each neuron draws its own line, and the next neuron combines these lines. Three lines make a triangle; twelve lines make almost a circle (Figure 5.1). A network with several such layers is called a **deep** network. Training works the same way as before: predict, measure the mistake, and adjust. The mistake is passed backwards through the layers to adjust each weight; this is called **backpropagation** (Rumelhart et al. 1986).

The big advantage of deep networks is that they can learn their own clues from raw data like images or text (LeCun et al. 2015). But when we already have clean, useful numbers, a deep network adds little. In our example, a small network does about as well as simple logistic regression (0.180 against 0.175).

● KEY IDEA

Deep networks are worth it for raw, messy data. For clean numbers, start with a simple model.

Idea two: words as numbers

A computer can only calculate with numbers, so words must become numbers. The trick comes from the study of language: *we understand a word by the words around it.* "Phishing" and "fraud" appear in similar sentences, near words like "email", "link" and "password". If we turn each word into a short list of numbers based on its neighbours, words used in similar ways get similar lists. This list is called an **embedding**. Similar documents also get similar embeddings, and this is how AI search finds related text.

But embeddings capture what a text is *about*, not the exact numbers inside it. To an embedding, an alert with a threat score of 0.05 and one with 0.95 look almost the same. In our example, predicting from similar past alerts gives honest probabilities but weak ranking (AUC 0.68, against 0.89 for logistic regression).

Idea three: words that look at each other

"Invoice, not phishing" and "phishing, not invoice" have the same words but opposite meanings. **Attention** solves this. Each word looks at the other words in the sentence and decides which of them change its meaning. In "invoice, not phishing", the word "phishing" pays most attention to "not".

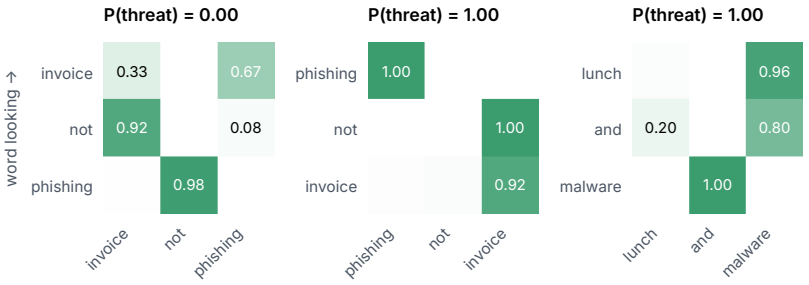


Figure 5.2. What each word attends to after training. In "invoice not phishing", the word "phishing" puts nearly all its attention on "not". Nobody told the model about negation; it found that looking at the word before was the cheapest way to get the answers right.

In Figure 5.2, nobody told the model about the word "not"; it learned this by itself. A model that only counts words reaches 84% accuracy on this task, while attention reaches 100%.

A **transformer** is many attention layers stacked together (Vaswani et al. 2017). There are two types. A **reader** sees the whole input at once, which suits understanding and deciding. A **writer** sees only the words before the current one, because it is producing text one word at a time. LLMs like ChatGPT are writers.

Idea four: very large models, and a side effect

If we train a huge writer to guess the next word on a huge amount of text, it learns grammar, facts and some reasoning (Kaplan et al. 2020). Then it is tuned with good example answers, and with feedback on which answers people prefer (Ouyang et al. 2022). But people usually prefer answers that sound confident. OpenAI reported that GPT-4's confidence was well calibrated before this tuning and less calibrated after it (OpenAI 2023).

There is a deeper reason too. If we train a network that is much bigger than needed for too long, its accuracy hardly changes, but its calibration error grows from 0.019 to 0.072 (Figure 5.3). Once the model gets most training cases right, the only way to improve its score further is to become more and more certain. It becomes certain about everything, even noise. Temperature scaling from Chapter 3 brings the error back to 0.046.

● KEY IDEA

Large models trained for too long become too confident. Their probabilities must be corrected.

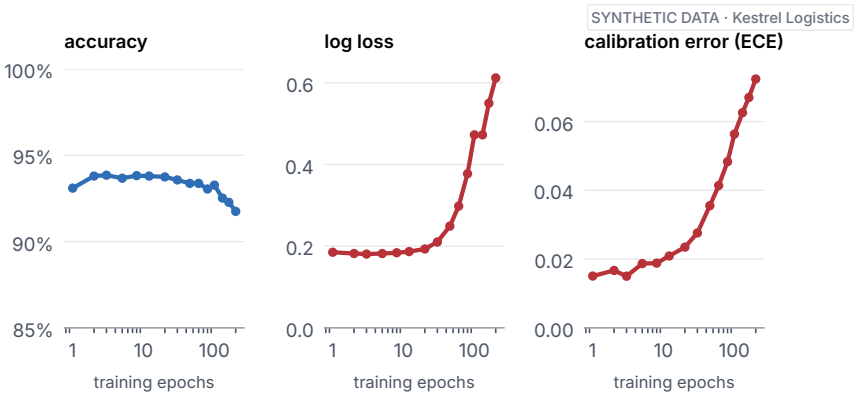


Figure 5.3. The same training run, on test alerts. Accuracy (left) barely moves. Log loss (middle) and calibration error (right) are flat for a while, then climb steeply. Watch only accuracy and you'd never notice.

How an LLM writes, and why this is a problem for decisions

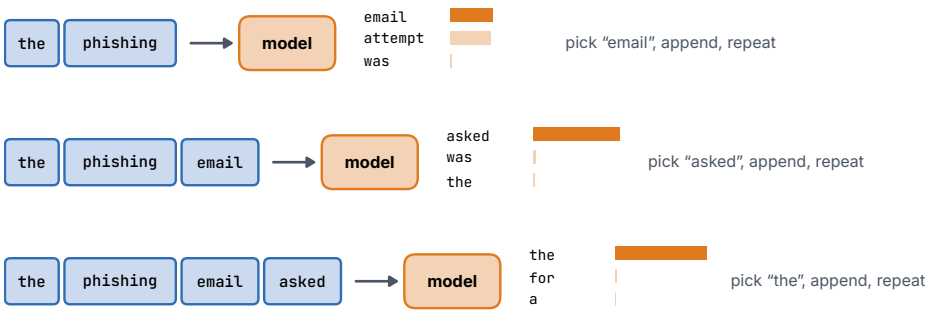


Figure 5.4. The generation loop. The model reads everything so far, produces a probability for every possible next token, one gets picked, it's appended, and the loop runs again.

An LLM writes one small piece of text at a time, called a **token**. At each step, it gives a probability to every possible next token, picks one, adds it, and repeats (Figure 5.4). This creates three problems for decisions.

Longer answers take more time and money. Each token waits for the previous one. In our example figures, a one-word answer takes about 0.5 seconds, but a long step-by-step answer takes about 25 seconds.

The same question can get different answers. The model usually picks the next token with some randomness. In our example, out of forty borderline alerts asked twenty times each, 39 got different answers at different times.

The answer is hidden inside text. If we ask for the answer in a fixed format (JSON), some answers still come back broken, with extra chatting or with invented fields. "JSON mode" fixes the *format*, not the *correctness*: a model can still write "benign" for a real attack in perfect format.

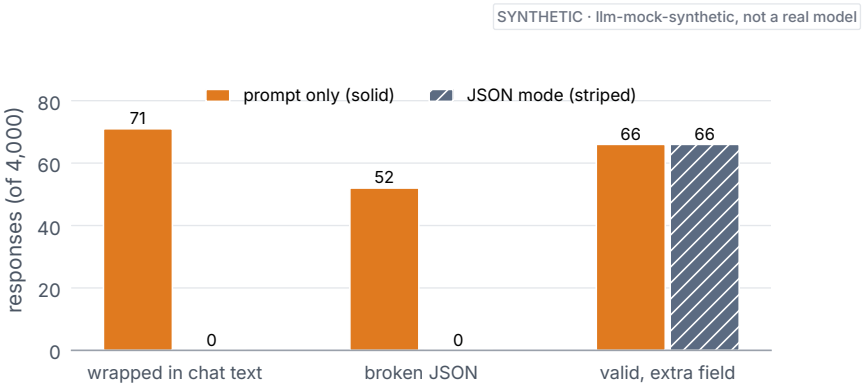


Figure 5.5. What came back when we asked for JSON, over 4,000 alerts. With only the prompt, some responses were wrapped in friendly chat, some were cut off mid-object, and some were valid JSON with a field nobody asked for. With JSON mode on, the first two problems vanish. The third doesn't.

● KEY IDEA

A correct format does not mean a correct answer.

Three kinds of "confidence"

Suppose an LLM replies: `"verdict": "malicious", "confidence": 0.95`. Can we trust this 0.95?

- **The written confidence** is just text that the model wrote. In our example, the LLM used only 6 different values, mostly 0.9 or more, so it hardly separates attacks from normal alerts.
- **The token probability** is the probability the model gave to the word "malicious". It ranks better but is overconfident, and it is the probability of a *word*, not of an answer.
- **A decision model's probability** is designed to be the probability of the answer itself (Figure 5.6). This is the gap that System One models try to fill.

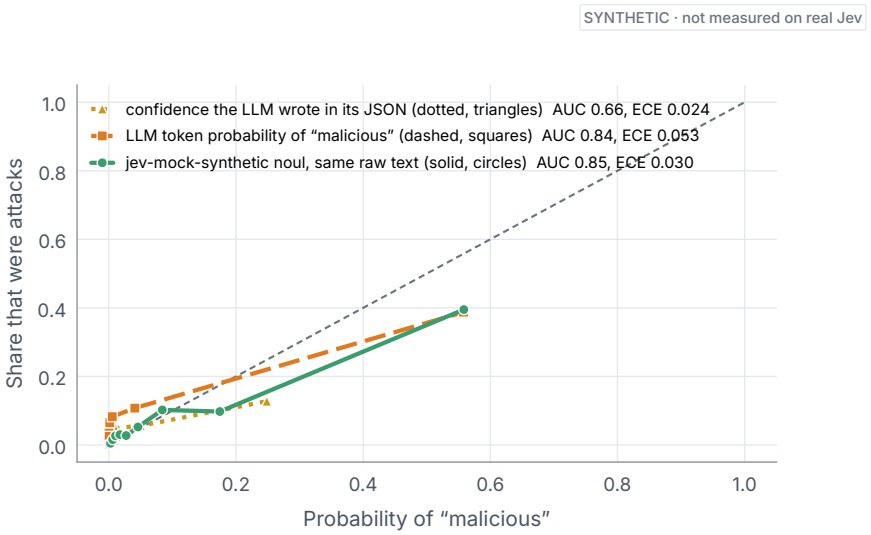


Figure 5.6. Reliability of three “confidences” on the same test alerts, all reading the same raw text. The stated confidence barely separates attacks from harmless alerts. The token probability ranks well but is overconfident. The decision model ranks best and also sits close to the diagonal. Synthetic: an illustration, not a measurement of real models.

KEEP THESE

- Deep networks are worth it for raw, messy data. For clean numbers, start with a simple model.
- Large models trained for too long become too confident. Their probabilities must be corrected.
- A correct format does not mean a correct answer.

EXERCISES

1. A company says its classifier is “97% accurate” and gives a confidence with every answer. You plan to act automatically when the confidence is above 0.9. What would you check first?

Next chapter: how AI systems read company documents and take actions, and why most of their weak points are small decisions.

CHAPTER 6

RAG, agents, and where they break

If you ask an LLM what your company's backup account is allowed to do, it will give a confident and well-written answer that is completely made up. It has never seen your company's documents. So we give it the right documents to read, and then we give it tools so it can take actions. This is where most of today's AI excitement is, and also where many hidden failures happen. As you read, notice one pattern: almost every weak point is a small decision.

BY THE END OF THIS CHAPTER

- Understand RAG, and why the search step is the weak link.
- See that most of an agent's work is making decisions.
- Learn how one planted sentence can mislead a model, and how to stop it.

Giving the model the right page

Retrieval-augmented generation, or **RAG**, works in three steps. First, search your documents for the passages most related to the question. Second, add those passages to the question. Third, ask the LLM to answer using them (Figure 6.1) (Lewis et al. 2020). In our example, with a set of company policies mixed with many similar-looking but useless documents, the search found the correct passage first 81% of the time.

The most important part of RAG is not the LLM; it is the search. If the correct passage is not found, the LLM simply builds the most likely-sounding answer from whatever it has. The answer will look fluent, will even mention a source, and will still be wrong.

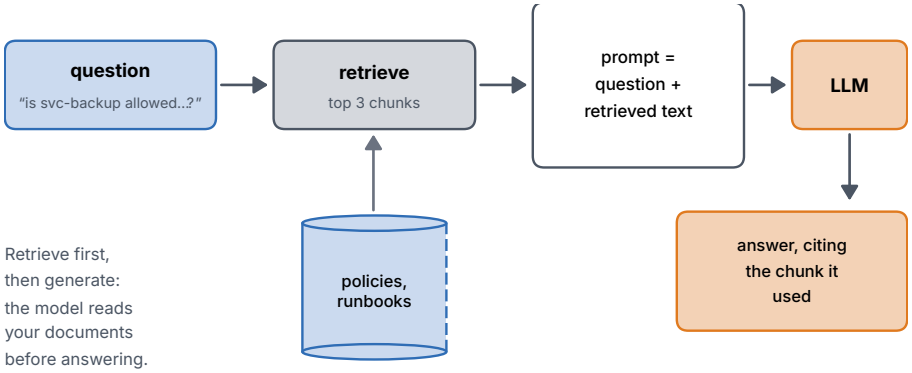


Figure 6.1. Retrieval-augmented generation. The question goes to a search step first, which pulls the most relevant chunks from your documents. The LLM sees the question plus those chunks, and answers from them, ideally citing which one it used.

KEY IDEA

A RAG system is only as good as its search step. Test the search first.

Some questions have no answer in the documents at all. A good system should then say "I don't know". A simple signal is how closely the *best* passage matches the question.

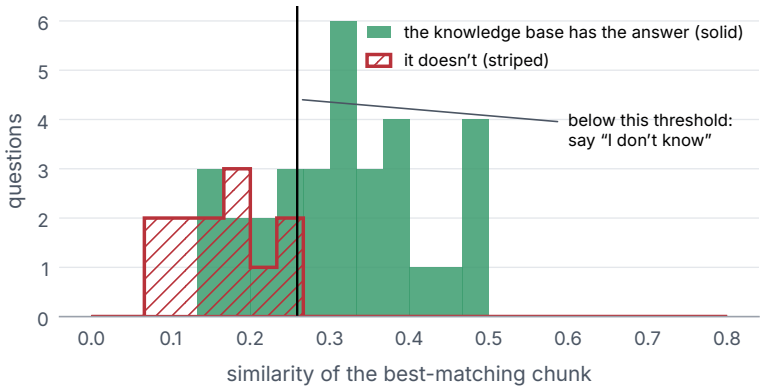


Figure 6.2. The similarity of the best-matching chunk, for questions the knowledge base can answer (green, solid bars) and questions it can't (red, striped outline). They overlap, but not much. Below a threshold, the system should decline to answer.

With a threshold on this match score, the system answered 72% of the questions that had answers, and only 8% of those that did not (Figure 6.2). So even a system built for writing depends on a small yes-or-no decision with a threshold.

A model in a loop: agents

An **agent** is a model working in a loop with tools: it looks at the situation, decides what to do, does it, and looks again. When people say an agent "uses a tool", the model does not press any button itself. It only writes which tool it wants and with what inputs (Yao et al. 2023). Your program checks this request and runs it. So each tool request is really a choice from a fixed list.

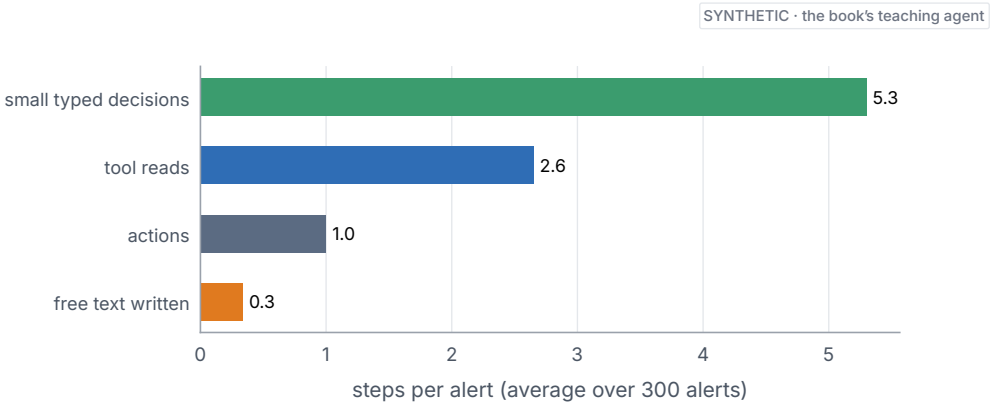


Figure 6.3. What the agent spent its steps on, averaged over 300 alerts. Small typed decisions dominate. Free text, the thing LLMs are famous for, is a small minority of the work.

Consider an example security agent. For each alert, it decides which evidence to collect, whether it has enough, and the final verdict, and then it acts. When we count its steps, decisions are 57% of all steps and writing is only 4% (Figure 6.3). Count the steps of any agent, and you will usually find the same pattern.

● KEY IDEA

Most of an agent's work is making small decisions, not writing.

Small mistakes multiply

If each step is correct 95% of the time, a task with ten steps succeeds only $0.95 \times 0.95 \times \ldots$ ten times, which is about 60%. With twenty steps, it is about 36%.

This is why agents look great in demos, which are short and easy, but fail in real work, which is long and messy (Figure 6.4). Since decisions are most of the steps, they carry most of the risk. If each decision comes with a trusted probability, an unsure step can be stopped and sent for checking.

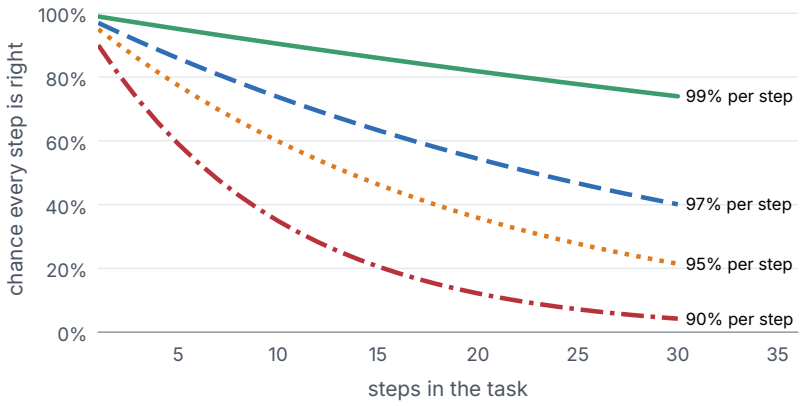


Figure 6.4. The chance that every step of a task goes right, for different per-step reliabilities. Small per-step error rates become large task failure rates as the number of steps grows.

When the input itself lies

Agents read text written by other people: emails, tickets, web pages. An attacker can hide instructions or false statements in that text. This is called **prompt injection** (Greshake et al. 2023). It is not only a problem for LLMs; any model that reads text written by an attacker can be misled.

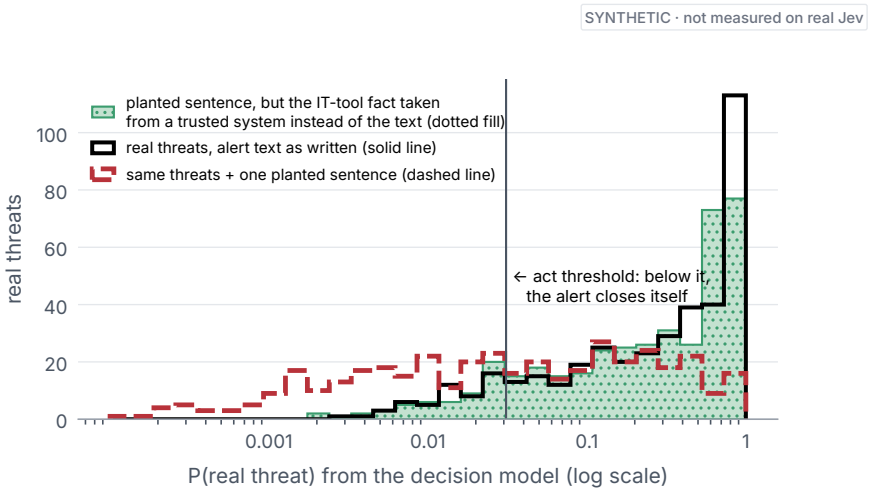


Figure 6.5. How likely the decision model thinks each real threat is, reading the alert text as written (black, solid line) and with one planted sentence (red, dashed line). The sentence pulls many real threats below the auto-close threshold. When the same fact comes from a trusted system instead (green, dotted fill), the planted sentence does nothing.

For example, some alerts contain the line “Matches approved IT tooling (change ticket on file)”, which strongly suggests the alert is harmless. Suppose an attacker adds this line inside a file name. In our example, the share of real attacks that would be closed without anyone checking jumps

from 13% to 50% (Figure 6.5). The fix is in the design, not in a smarter model: ask the company's ticket system directly whether such a ticket exists, and pass that answer separately, where the attacker cannot write. Then only 14% are closed.

● KEY IDEA

Any fact that makes a case look safer must come from a system the attacker cannot change.

The decision layer

Almost every protection in this chapter is the same thing: a check that becomes a decision, with a threshold based on the cost of mistakes. Put all of them together and you get a **decision layer** between what an agent wants to do and what actually happens.

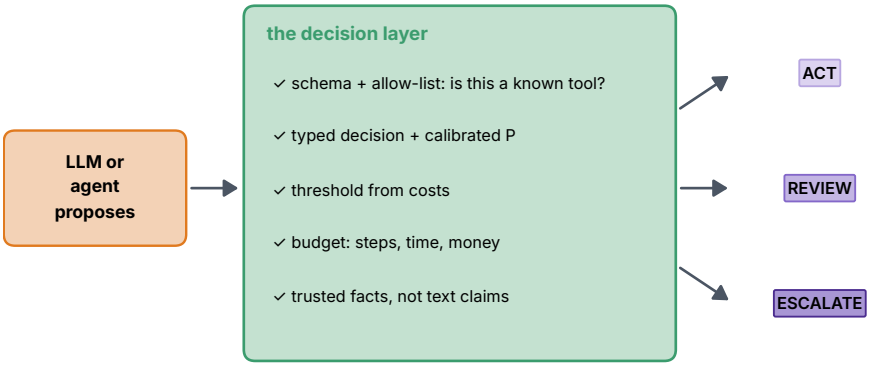


Figure 6.6. A decision layer between what an agent proposes and what actually happens: schema and allow-list checks, a typed decision with a calibrated probability, a threshold from costs, a budget and trusted facts, then one of three doors.

The decision layer checks each proposed action against what is allowed, asks clear questions, uses thresholds based on costs, takes risky facts only from trusted systems, and sends each case through one of the three doors: act, review or escalate (Figure 6.6). If anything fails, such as a timeout or an unreadable answer, the case goes to a person. None of this needs a smarter model. It needs decisions that are fast, cheap and calibrated.

SET THE THRESHOLD

An agent can close an alert, add a note, block a sender's website, or disconnect a laptop from the network. Which actions can it do automatically when confident, and which should always need a person?
(Closing an alert and adding a note are cheap and can be undone, so they can be automatic with a good threshold. Blocking a website can be undone but may hurt business, so use a stricter threshold and a list of partners it must never block. Disconnecting a laptop stops someone's work and may destroy evidence, so a person should decide.)

KEEP THESE

- A RAG system is only as good as its search step. Test the search first.
- Most of an agent's work is making small decisions, not writing.
- Any fact that makes a case look safer must come from a system the attacker cannot change.

EXERCISES

1. A company says its agent gets 90% of steps right, and a normal task has fifteen steps. What success rate would you expect for a full task?

Next chapter: the idea behind System One models, and what we really know about Jev.

CHAPTER 7

Fast and slow: System One models and Jev

When TypeSafe AI released Jev, it did not call it a better LLM. It called it the first of a new type of model, a *System One* model (TypeSafe AI 2026). In this chapter we first understand the idea behind this name. Then we separate what is actually known about Jev from what is only claimed.

BY THE END OF THIS CHAPTER

- Understand System 1 and System 2 thinking, in people and in software.
- Send cases between a fast system and a slow system based on doubt.
- Sort claims about a new model into three groups.

Two ways of thinking

What is 2 + 2? The answer came immediately. Now try 17 × 24. This feels different: you have to stop and work step by step. The psychologist Daniel Kahneman called these two ways of thinking **System 1**, which is fast and automatic, and **System 2**, which is slow and careful (Kahneman 2011). System 1 handles most of our daily life and is usually right, but it can also be confidently wrong. Try this: a bat and a ball cost ₹110 together, and the bat costs ₹100 more than the ball. The quick answer, ₹10 for the ball, is wrong. The ball costs ₹5 (Frederick 2005).

Software fits this idea well. Rules and simple classifiers work like System 1: fast and good with familiar patterns. LLMs, especially the ones that "reason" step by step, are closest to System 2: slow, general, and paid for by every word. In real work, as Figure 7.1 shows, most tasks are quick judgements that happen many times a day. Only a few need deep thinking. Yet today we often use slow System 2 tools for everything.

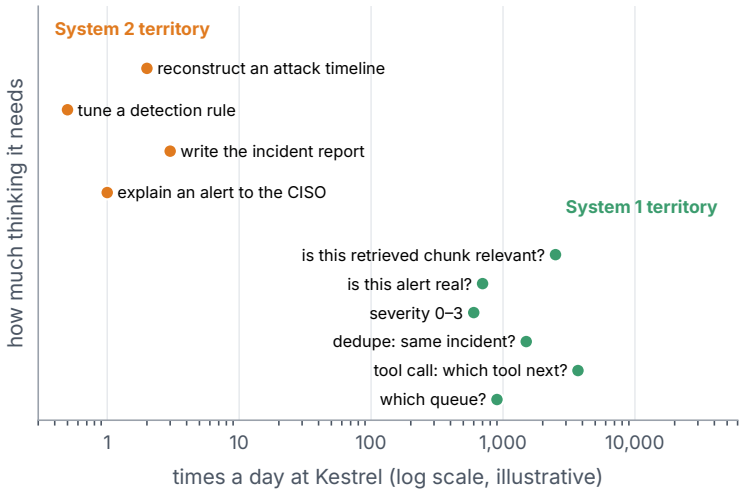


Figure 7.1. Tasks in and around a security team, placed by how often they happen and how much thinking each needs (illustrative). The frequent ones are grouped at the bottom right: quick judgements made hundreds or thousands of times a day. The few that need real reasoning happen rarely.

● KEY IDEA

Most real work needs fast System 1 judgement, but we often use slow System 2 tools for it.

Send the doubtful cases to the slow system

If we have a fast system and a slow system, how should we divide the work? A good expert handles most cases in seconds, and *notices* the cases that do not feel right. We can do the same: let the fast system see every case, and send only the cases it is unsure about to the slow system.

In our example, sending the 20% of alerts that the fast system is least sure about gives 79% of the total benefit that careful checking could give. Sending a random 20% gives only 3% (Figure 7.2). And we pay the slow cost only for that 20%.

● KEY IDEA

A fast system is useful in two ways: it is quick, and it knows which cases it should not decide.

This works only if the fast system is calibrated: its "50%" must really mean "could go either way". Kahneman and the psychologist Gary Klein agreed

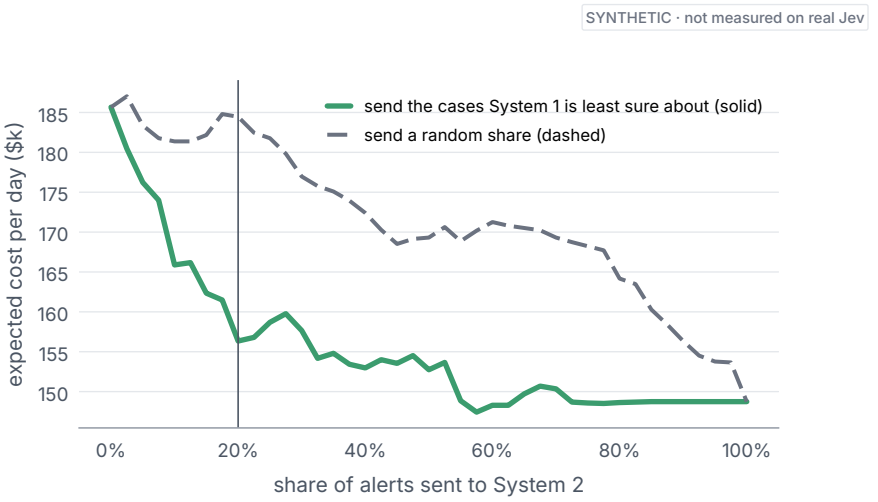


Figure 7.2. Expected cost per day as more alerts go to System 2. Sending the alerts System 1 is least sure about (green, solid line) gets most of the benefit early. Sending a random share (grey, dashed line) gets almost nothing until you send nearly everything.

that fast judgement can be trusted when the situation is regular, there is a lot of practice, and feedback comes quickly (Kahneman and Klein 2009).

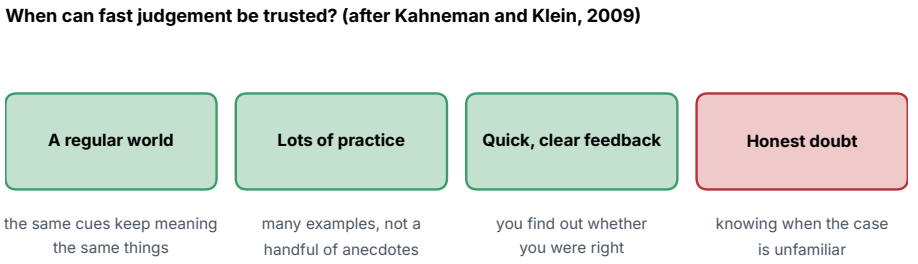


Figure 7.3. Conditions under which fast, intuitive judgement tends to be reliable, adapted from Kahneman and Klein. Alert triage mostly meets the first three. The fourth is what separates good experts, and good models, from confident ones.

For machines, we add a fourth condition, as in Figure 7.3: the system must honestly know when it is unsure.

What we know about Jev: three groups

When this book was written, Jev was new and in early access. The public information was its official software library, its documentation, news reports and a few early research papers. How much can we really know

about a model from outside? The best method is to sort every claim into three groups.

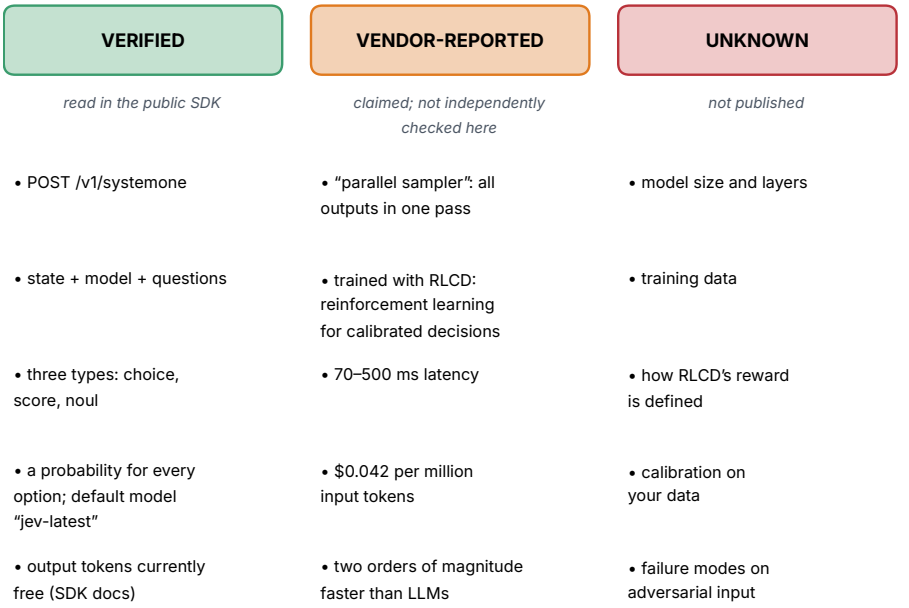


Figure 7.4. Every claim about Jev in this book, sorted into three piles. The first pile you can check yourself today. The second is what TypeSafe says. The third is what nobody outside the company can know yet.

● KEY IDEA

Sort every claim about a new model into three groups: checked, claimed by the company, and unknown.

Checked: how you use it. You send one request with three parts: the **state** (a description of the situation), the **model** name, and a set of clear **questions**. For each question, you get back an answer with a probability for every option. There are three types of questions: yes/no, pick-one, and rate-on-a-scale (Chapter 8).

Claimed by the company: two ideas. First, a "parallel sampler" that produces all answers in a single step instead of word by word, which TypeSafe says makes it fast. Second, a training method called **RLCD**, which is said to reward calibrated probabilities. The details are not public. But from Chapter 2 we know one way such a method could work: a scoring rule like log loss rewards honest probabilities.

request (what the SDK sends)

```
POST /v1/systemone

{

  "state": "I was charged twice…",

  "model": "jev-latest",

  "questions": {

    "topic":  {"type": "choice", "criteria": {…}},

    "urgent": {"type": "noul", "instructions": …},

    "tone":   {"type": "score", "criteria": […]}

  }

}
```

response

```
{

  "model": "jev-mock-synthetic",

  "answers": {

    "topic":  {"choice": "billing",

               "probabilities": {…}},

    "urgent": {"noul": 0.61},

    "tone":   {"score": 0.22,

               "probabilities": {…}}

  },

  "usage": {"input_tokens": 89, …}

}
```

Figure 7.5. What a request to Jev contains, and the shape of the answer that comes back. The values shown are an example.

Unknown: the model's size, design and training data; how well calibrated it is on *your* data; how it behaves when attackers write the input; and how its behaviour changes between versions. These are exactly the things a decision system depends on most.

Check if the company's numbers agree

TypeSafe reports a response time of about 70 to 500 milliseconds and a price of $0.042 per million input tokens, with output free. If each decision reads about 500 tokens, a million decisions cost about $21. With an example mid-range LLM price, the same million decisions would cost about $660. LLM prices vary a lot, so treat this only as a rough size. And for any speed claim, always ask: *faster than what, and for which task?* (regolo.ai 2026)

● KEY IDEA

When a company gives several numbers, check whether they agree with each other.

In its launch demo, Jev played the video game *Doom*, making about ten decisions per second at a reported cost of $7 per hour. That is 36,000 decisions per hour, or about $0.0002 per decision. At the listed price, that

means about 4,600 tokens of game description per decision, which is reasonable. The numbers agree. This does not prove they are true, but if they had disagreed, we would have learned something important.

● KEY IDEA

Unknown facts about a new model are not a reason to avoid it. They are a reason to test it on your own data.

⚠ WHERE THIS BREAKS

System 1 and System 2 are a way of describing thinking, not two real parts of the brain, and a System One model does not have "intuition"; it is a trained function. Sending doubtful cases to the slow system fails when the fast system is confidently wrong about a whole new type of case. Random checks and monitoring are needed to catch this. And company numbers come from the company's own tests; for a new model, behaviour can change quickly, so fix a version and test again.

● KEEP THESE

- Most real work needs fast System 1 judgement, but we often use slow System 2 tools for it.
- A fast system is useful in two ways: it is quick, and it knows which cases it should not decide.
- Sort every claim about a new model into three groups: checked, claimed by the company, and unknown.
- When a company gives several numbers, check whether they agree with each other.
- Unknown facts about a new model are not a reason to avoid it. They are a reason to test it on your own data.

EXERCISES

1. List five decisions in your own work that are quick System 1 judgements, and two that need slow System 2 thinking. What makes them different?

Next chapter: the three types of questions, and how to ask each one well.

CHAPTER 8

Asking the right question: choice, score, noul

A System One model asks you to put every question into one of three shapes: yes or no, pick one option, or rate on a scale. You may think real decisions are more complex than this. But look at your own daily work, and you will see that most complex decisions can be broken into a few simple questions of these three shapes: *Is this real? Which team? How urgent?* The difficulty is usually in how we ask the question, not in the shape of the answer.

BY THE END OF THIS CHAPTER

- Choose between yes/no, pick-one and scale questions.
- Avoid two common traps: the missing option and the misleading average.
- Make decisions using all the probabilities, not only the top answer.

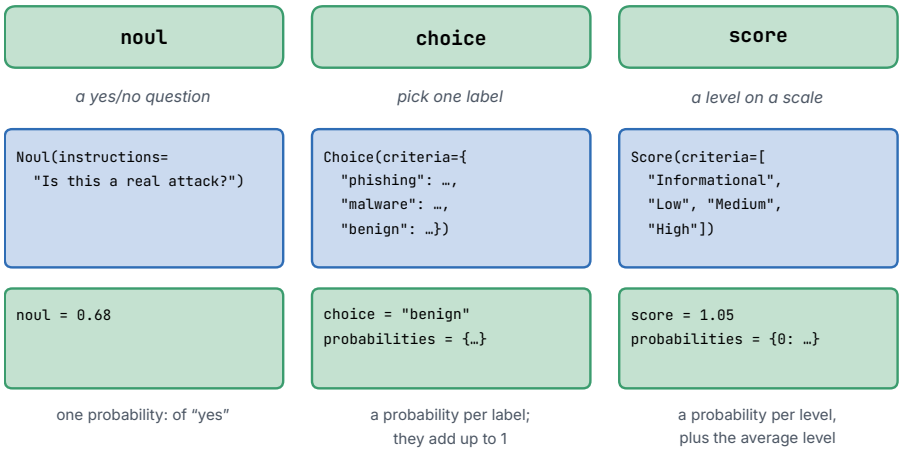


Figure 8.1. The three question types, what you send, and what comes back. A noul returns one probability. A choice returns a probability for every label. A score returns a probability for every level, plus the average level.

Yes or no: one fact at a time

TypeSafe calls a yes-or-no question a **noul**. It returns one number: the probability of "yes". This is the most useful type, and also the easiest to ask badly. Three habits help.

Ask about one fact at a time. "Is this a real attack on an important server?" combines two facts. If the answer is "no", you cannot know which part was false. Ask two separate questions, and combine the answers in your own program, where the logic is clear.

Say exactly what "yes" and "no" mean. Is an employee breaking a company rule "malicious"? Is a test attack paid for by the company "malicious"? Write it clearly: *yes* means "a real attacker or real harmful software is involved"; *no* means "normal activity, an approved test, or a false alarm".

Avoid double negatives. "Is this not unlikely to be safe?" confuses both people and models.

Pick one: always give a "none of these" option

A **choice** question returns a probability for every option, and all of them add up to 1. This has a hidden danger: the model *must* spread its belief across the options you gave it. If the correct answer is not in your list, the probability does not disappear. It goes to the wrong options.

In our example, a model was asked to label 400 harmless alerts twice. When "benign" (harmless) was one of the options, 96% were labelled benign. When "benign" was removed, every single alert got a threat label, most of them with high confidence.

The model was not exactly wrong: we asked *which type of attack*, and it answered (Figure 8.2). But anyone who reads "malware, 0.95" will take the wrong action.

> ● KEY IDEA
>
> **In a pick-one question, always include an option like "none of these", so every real case has a correct place to go.**

Also make sure only one option can be true at a time. If two options can both be true, ask separate yes/no questions instead.

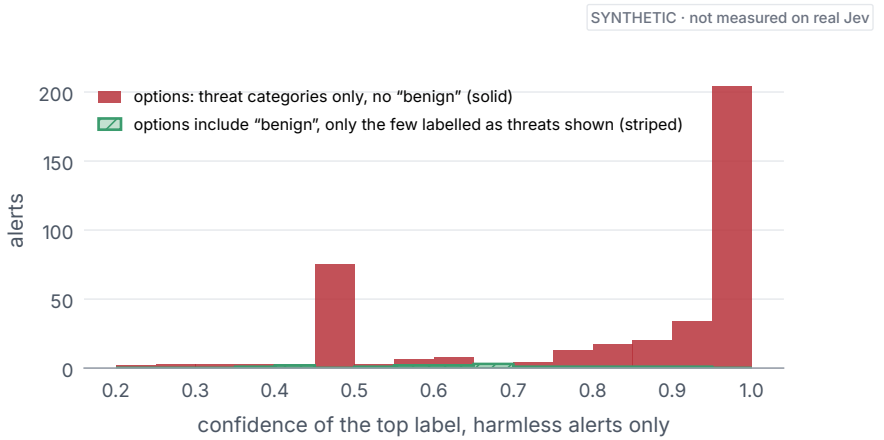


Figure 8.2. How confident the top label was for harmless alerts. With “benign” among the options (green, striped), almost all were labelled benign; only the few that weren’t are shown. With only threat categories to choose from (red, solid), every harmless alert got a threat label, most with high confidence.

Rate on a scale: do not trust the average

A **score** question returns a probability for each level of a scale, for example informational, low, medium and high, and also the *average* level. This average can mislead you badly.

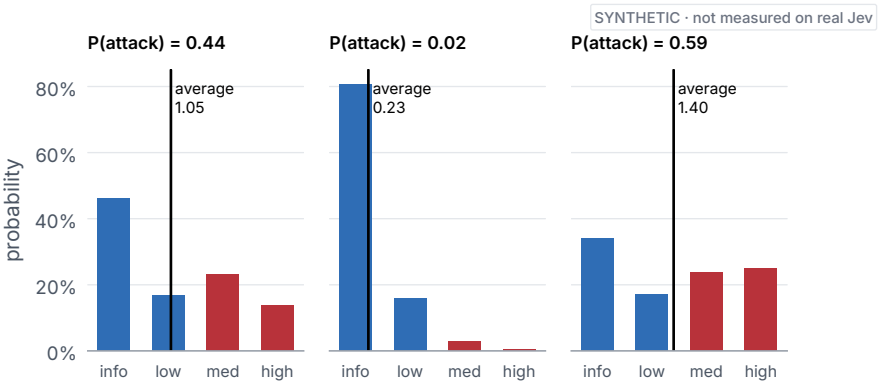


Figure 8.3. Severity probabilities for three alerts, with the average level marked. The first alert’s average sits at “low”. But its probability is split between “informational” and “medium or high”. Almost nothing is actually at “low”.

Look at the first alert in Figure 8.3. Its average level is 1.05, which reads as “low”. But almost no probability is actually at “low”. The model is really saying: *this is either nothing, or something serious*. There is a 37% chance it is medium or high. The average of two very different possibilities describes neither of them.

KEY IDEA

An average can hide the extremes. Set thresholds on the probability of the levels you care about.

A rule like "call someone if P(medium or high) is above 0.3" is clear and can be set using costs. A rule like "call someone if the average is above 1.5" is not.

Decide using all the probabilities

The most likely option is not always the one to act on. Suppose a pick-one question says: benign 0.40, stolen password 0.35, phishing 0.15, malware 0.10. "Benign" is the top answer, but together 60% says *some kind of attack*. Closing the alert would risk all of that.

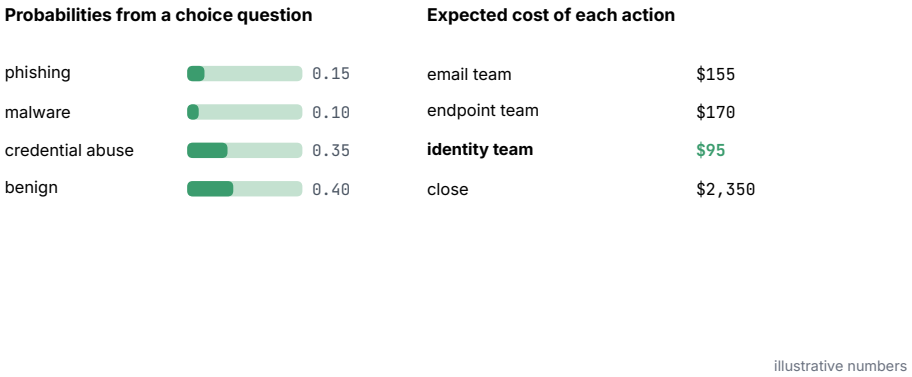


Figure 8.4. An illustrative choice between four labels, and the expected cost of four actions. "Benign" is the single most likely label, but closing the alert is the most expensive action, because the other 60% of the probability is spread across threats that are costly to miss.

The better way, as in Figure 8.4, is to calculate the expected cost of each action using all the probabilities, and choose the cheapest action. Here, that means sending the alert to the identity team. This is the formula from Chapter 4, extended to more options. It is also why we need a probability for every option. A text answer that only says "benign" throws away exactly the information we need.

When one situation needs several answers, ask all the questions together in one request, and combine the answers in your own program, where every rule is visible.

⚠ **WHERE THIS BREAKS**

These three shapes cannot express everything. A quantity ("how many megabytes?"), a name ("which username?") or several details from a long document still need something else, often an LLM (see Chapter 12). Answers to separate questions about the same case may not agree, so your program should handle that. And changing the wording of a question changes the probabilities, so test every wording change.

● **KEEP THESE**

- In a pick-one question, always include an option like "none of these", so every real case has a correct place to go.
- An average can hide the extremes. Set thresholds on the probability of the levels you care about.

EXERCISES

1. Design the questions for a decision you know, such as loan applications or customer complaints. Which are yes/no, which are pick-one, and which are scales? What "none of these" option does each pick-one question need?

Next chapter: the company says Jev's probabilities are calibrated, and a critic says they cannot be. Here is how to check this yourself.

CHAPTER 9

Testing a model's calibration yourself

A week after Jev was launched, two opinions were being shared. TypeSafe said Jev gives *calibrated* probabilities (TypeSafe AI 2026). The data scientist Alex Molas replied with a post titled "Jev can't be calibrated". His point was that calibration depends on *your* data, and Jev has never seen your data, so its outputs should be treated as scores, not probabilities (Molas 2026). Who is right? In a precise sense, both are. And you can settle the question for your own use in one afternoon.

BY THE END OF THIS CHAPTER

- Understand why calibration depends on your data.
- Fix a base-rate mismatch using a few hundred labelled cases.
- Know how many labelled cases a calibration check needs.

Why the critic has a point

From Chapter 2, a probability tells us how often something happens among similar cases. Now imagine two companies send the same alert text to the same model. The model reads the same input, so it gives both companies the same number. But suppose one company is attacked often and the other rarely. Then, among alerts like this one, the real share of attacks is different at the two companies. One number cannot be correct for both.

This is the base-rate trap from Chapter 2, now catching the model. A model can be perfectly calibrated for one mix of cases and wrong, always in the same direction, for a different mix.

● KEY IDEA

Calibration belongs to a model *together with the data it is used on*.

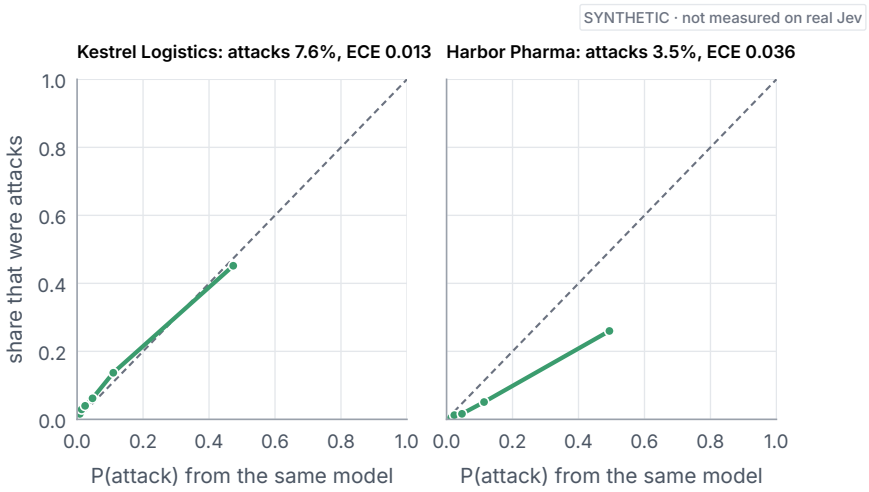


Figure 9.1. The same model's reliability at two companies. At Kestrel, whose mix of alerts matches what the model was tuned for, it sits close to the diagonal. At Harbor, where attacks are rarer, every probability is too high.

To see this, let us imagine a second company, Harbor Pharma. Its alerts come from the same kind of systems and look the same, but real attacks are rarer: 3.5% of alerts, compared with Kestrel's 7.6%. The same model has a calibration error of 0.013 at Kestrel, but 0.036 at Harbor, and the error is always in the same direction. Harbor's average predicted risk is 7.1%, while the real rate is 3.5% (Figure 9.1). The model did not change. Only the situation changed.

Fixing it

This kind of error is often easy to fix. If the model's probabilities were right for one base rate, and you know your own base rate, you can adjust the odds by the ratio of the two. This is Bayes' rule again (Saerens et al. 2002). With Harbor's base rate known, its ECE falls to 0.010. **Platt scaling**, fitted on just 300 of Harbor's labelled alerts, does even better: 0.007.

Platt scaling can also fix probabilities that are too spread out or too bunched together, not only probabilities that are too high (Figure 9.2).

How many labelled cases do you need?

A **labelled case** is a past case where we know the final outcome. Small samples can badly mislead a calibration check.

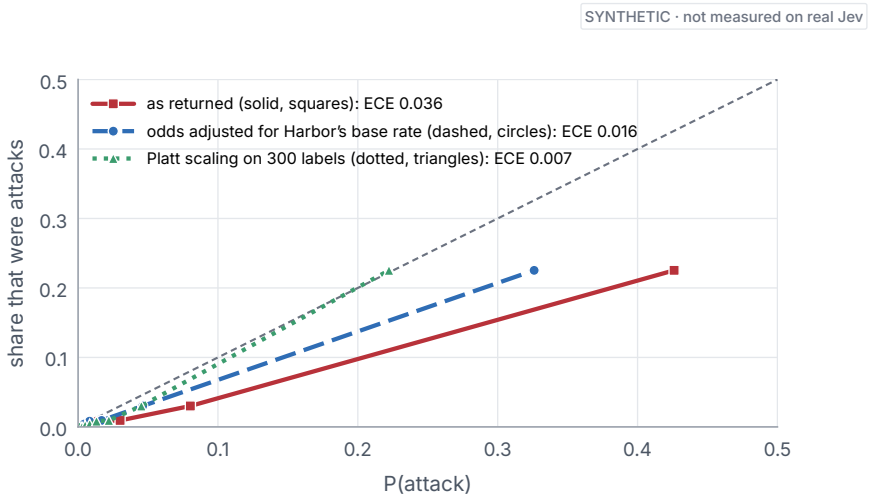


Figure 9.2. Harbor's alerts, on the cases not used for fitting: as returned (red, solid line), after adjusting the odds for Harbor's base rate (blue, dashed line), and after Platt scaling on 300 labelled alerts (green, dotted line). Both fixes bring the curve most of the way back to the diagonal.

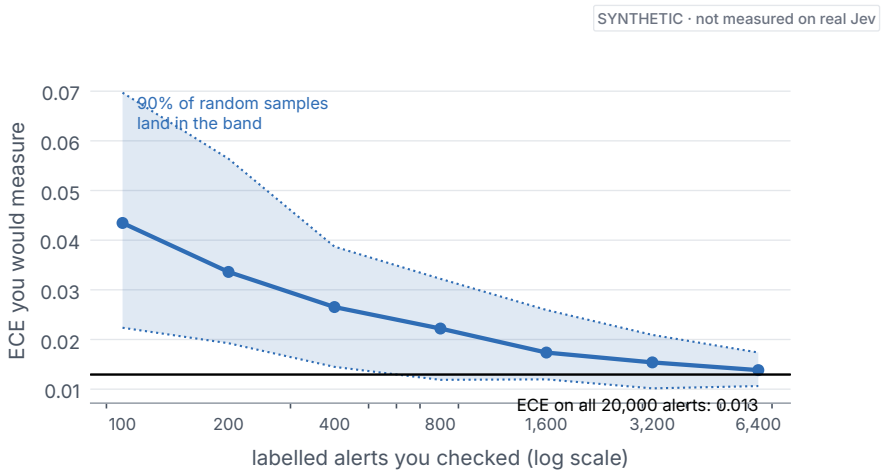


Figure 9.3. The ECE you'd measure from a random sample of labelled alerts, for different sample sizes, on a model whose ECE over all 20,000 alerts is shown by the labelled black horizontal line. With a hundred labels the measurement is both noisy and too high on average: a calibrated model looks miscalibrated. The error shrinks as samples grow.

In our example, the model's real ECE over all 20,000 alerts is 0.013. But if we measure it on a random 100 labelled alerts, the typical result is 0.043, more than three times too high. One time in twenty, it even shows 0.070 (Figure 9.3). With few cases in each group, the results jump around by chance, and ECE adds up these jumps without caring about their direction. So random noise always makes ECE look worse.

● KEY IDEA

With few labelled cases, ECE looks worse than it really is. Report a range, and do not judge a model on a hundred cases.

A simple rule: a few hundred labelled cases are enough to *fix* a big error that always goes in one direction. A few thousand are needed to *confirm* that a model is calibrated within one or two percentage points.

For pick-one questions, first check the top answer: when its confidence is about 0.8, is it correct about 80% of the time? In our example, the top category has an average confidence of 89% and is correct 90% of the time.

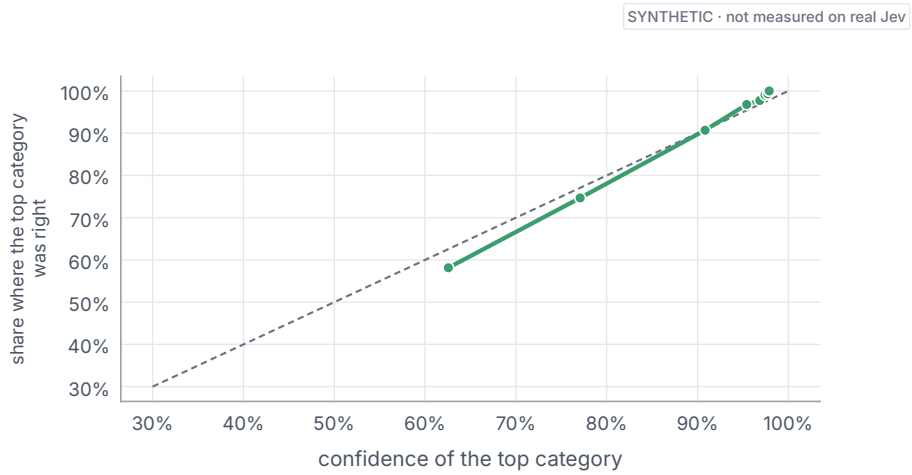


Figure 9.4. Top-label calibration for the category question on 2,500 alerts: the confidence of the model's top category against how often that category was right.

A routine you can repeat

1. **Collect labelled cases** from recent cases whose outcome is now known. Include a random sample of cases the machine closed by itself; otherwise, you will only have labels for cases that people looked at.

2. **Draw the reliability diagram**, with the number of cases in each group visible, and report ECE as a range. Look most carefully near your thresholds.
3. **Check each group separately**, for every group you treat differently.
4. **Fix the problem**: use a base-rate adjustment if that is the issue; otherwise, use Platt or temperature scaling on your labelled cases.
5. **Repeat** this regularly, and whenever the model version or your data changes.

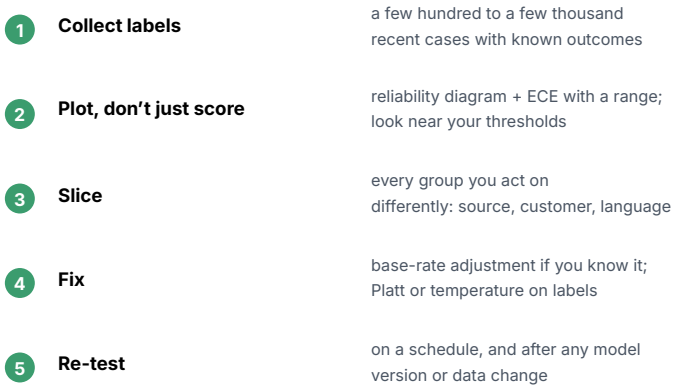


Figure 9.5. A testing routine for any model that returns probabilities, Jev included.

So, the company may be right that Jev is calibrated on the data it tested. The critic is right that this does not prove it is calibrated on yours. This routine turns the argument into a measurement.

● KEY IDEA

Treat a new model's outputs as scores until you have tested them on your own data.

⚠ WHERE THIS BREAKS

Labelled cases are not perfect either. People disagree, some attacks are never discovered, and some outcomes become clear only later: an alert closed today may turn out to be part of an attack found next month. So use cases old enough that their outcome is settled. And be careful with cases that have labels only because of the model's own decisions: if the model closed a case, nobody checked it.

KEEP THESE

- Calibration belongs to a model *together with the data it is used on*.
- With few labelled cases, ECE looks worse than it really is. Report a range, and do not judge a model on a hundred cases.
- Treat a new model's outputs as scores until you have tested them on your own data.

EXERCISES

1. You have 150 labelled cases and you measure an ECE of 0.05. A colleague says the model is badly calibrated. What would you reply, and what would you do next?

Next chapter: six different ways to make the same decision, compared fairly. Jev does not win every round.

CHAPTER 10

The bake-off: six ways to make a decision

In every team that makes many decisions, there is a long-running debate. The senior expert says the hand-written rules are enough, because anyone can read them. The data scientist says a trained model would do better, if someone labels enough data. The new engineer says an LLM understands cases better than both. And now someone has read about Jev. Each of them is right about something, because each is looking at a different quality. So let us put all the numbers on one page: six methods, one decision, one fair comparison. A comparison like this, where everyone makes the same thing, is sometimes called a "bake-off".

BY THE END OF THIS CHAPTER

- Compare six ways of making the same decision.
- Judge methods on more than accuracy.
- Choose a method using a scorecard.

The six methods

The decision is: *is this alert a real attack?* Every method gets three weeks of labelled history (14,973 alerts) to learn from, and is then tested on a fourth week it has never seen (5,027 alerts, of which 412 are real attacks). The six methods are shown in Figure 10.1:

- hand-written **rules**;
- **logistic regression** on the alert's clean data fields;
- a simple **text classifier** that counts words in the raw text;
- an **LLM** asked to answer in a fixed format;
- an **LLM as judge**, giving each alert a rating from 1 to 10;
- **Jev**, reading the clean data fields.

An important note: the numbers for the LLM and for Jev in this comparison are illustrative examples, not measurements of the real products, and

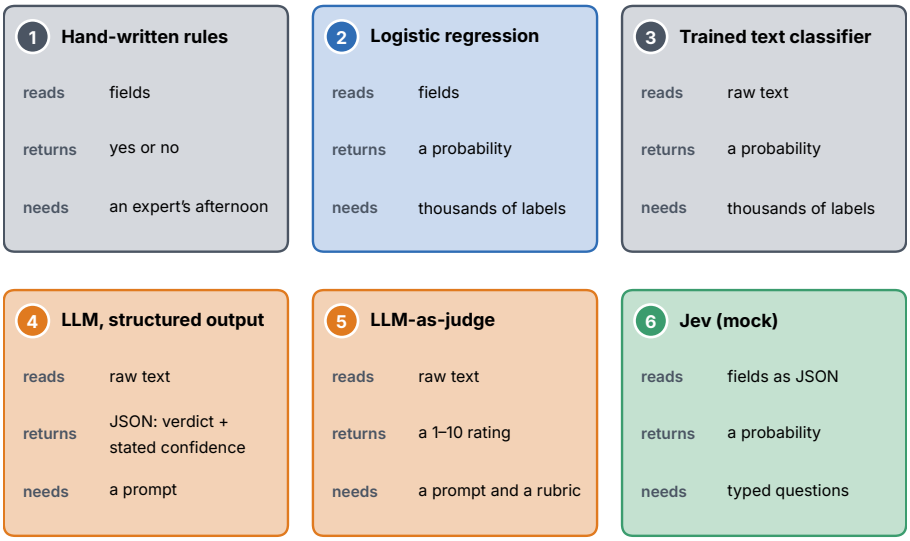


Figure 10.1. The six contestants: what each one reads, what it returns, and what it needs before it can start.

the example LLM was set up to read alerts less accurately. So if the LLM ranks below Jev, do not take that as proof about the real models. What you *can* rely on are the patterns that come from how each *type* of method works: how written confidences behave, how often text answers break, what happens when you ask twice, and what labels cost.

The results

Ranking. Logistic regression wins, by a small margin: AUC 0.896 against Jev's 0.892 (Figure 10.2). This is not luck. It was trained on fifteen thousand of this company's own labelled alerts, with clean data, for a decision that depends on that data in a fairly simple way. This is exactly where classic machine learning is strongest. Rules come last (0.66): a rule only says yes or no, so it cannot tell which of its many "yes" cases to check first.

Calibration. Most methods are close to the diagonal. The LLM methods are not. The LLM's *written* confidence used only 6 different values across five thousand alerts, mostly 0.9, 0.95 and 0.99. A number that a model writes is just text, and text prefers round numbers. The LLM-as-judge rating from 1 to 10, when read as a probability, has the worst calibration of all (0.102).

Labels needed. Logistic regression needs about 3,000 labelled cases to catch up with Jev, which needs none to start. A team with months of his-

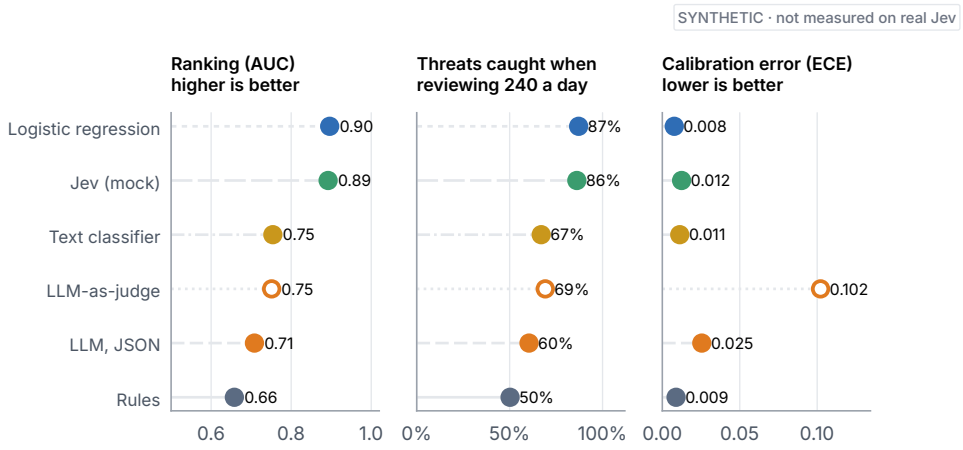


Figure 10.2. Three views of how well each method separates threats from harmless alerts, on the live week. Left: ranking quality (AUC). Middle: the share of all real threats found if the team reviews the top 240 alerts each day. Right: calibration error, where lower is better.

tory has these labels. A new team, or a completely new decision, does not.

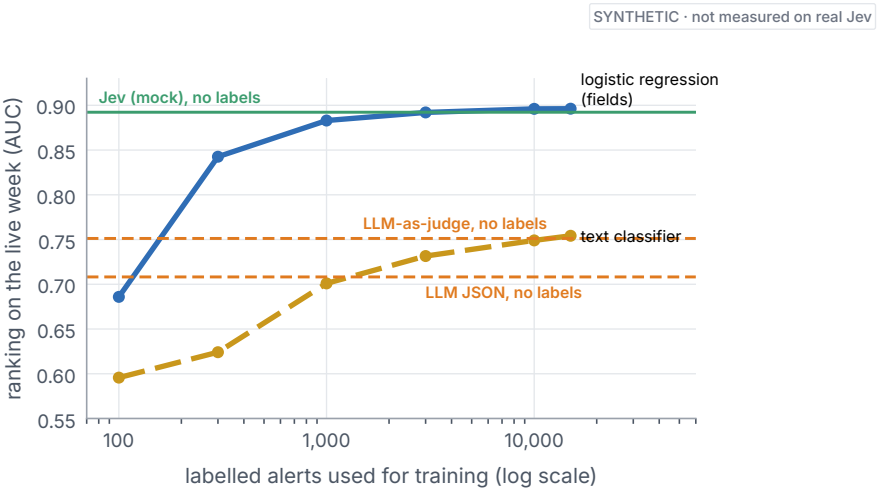


Figure 10.3. Ranking on the live week as the trained methods see more labelled alerts, averaged over random samples. Jev and the LLMs need no labels to start, so they're flat lines. Logistic regression catches up with Jev at about 3,000 labels.

● KEY IDEA

Labelled data is the hidden cost of trained models. A method that needs no labels can start today.

Speed and price. Nothing is faster or cheaper than a rule, and classic models are close behind: they take millionths of a second on computers you already have. Jev's figures, as reported by the company, are in between: about $21 per million decisions. The example LLM costs about $740 per million and takes about 1.4 seconds per answer.

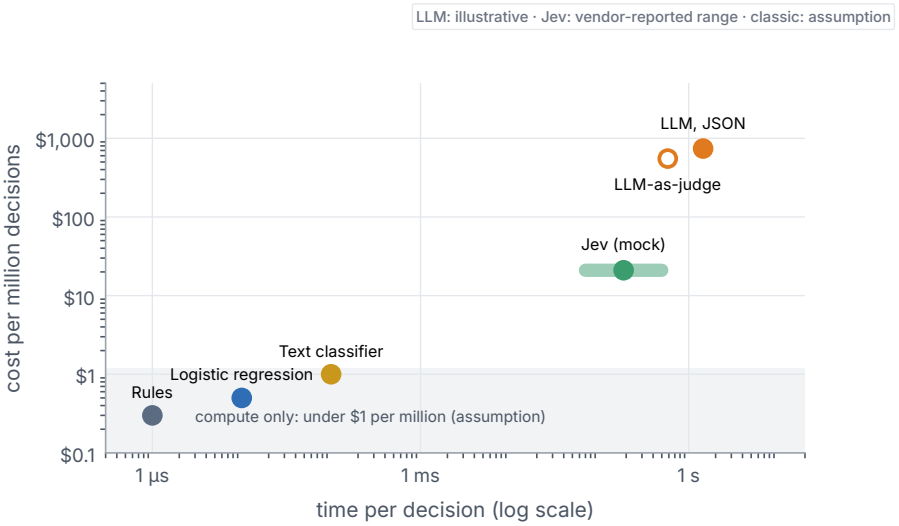


Figure 10.4. Time and cost per decision for each method, on log scales. Classic models run in millionths of a second, for fractions of a cent per million decisions (an assumption: computing cost only). The LLM figures are this book's illustrative assumptions; Jev's are vendor-reported.

● KEY IDEA

If a simple rule or logistic regression can make the decision well, use it.

Jev's price advantage is only against LLMs, not against everything.

Stability. When the same question was asked to the LLM five times, 5.2% of the answers changed at least once, and 3.2% of its answers could not be read by the program at all. Answers in a fixed typed shape cannot break like this.

The scorecard

SYNTHETIC · not measured on real Jev

	ranking (AUC)	caught at 240/day	calibration (ECE)	cost per million	time per decision	labels needed	same answer twice?	explains itself?
Rules	0.66	50%	0.009	< $1	1 µs	none	yes	fully
Logistic regression	**0.90**	**87%**	**0.008**	< $1	10 µs	~3,000	yes	weights
Text classifier	0.75	67%	0.011	$1	100 µs	15,000+	yes	word weights
LLM, JSON	0.71	60%	0.025	$740	1.4 s	none	not always	in words*
LLM-as-judge	0.75	69%	0.102	$552	583 ms	none	not always	in words*
Jev (mock)	0.89	86%	0.012	$21†	70–500 ms†	none‡	yes (mock)	no

Bold: best in column. † vendor-reported. ‡ none to start; about 300 to check and calibrate. Classic costs: compute only (assumption).
* reasons written after the fact, which may not be how the verdict was reached (Chapter 6). LLM figures: illustrative.

Figure 10.5. The whole contest on one page. Bold marks the best in each numeric column. Every accuracy number is synthetic; every cost and time is an assumption, an illustration, or vendor-reported, as marked.

Read Figure 10.5 column by column, and ask which columns matter for *your* decision.

- Thousands of labelled cases, clean data and a stable decision: logistic regression is hard to beat.
- Every decision must be explained to an auditor: rules are the clear winner.
- No labelled cases, raw text, or many different decisions: a System One model is the strongest choice, after a calibration check.
- The LLM is the weakest at deciding here, but it can do something no other method can: read anything, and *write*. Chapter 12 gives it exactly that job.

KEY IDEA

There is no single best method. There is only the best method for a particular decision.

⚠ WHERE THIS BREAKS

This is one decision in one example company, with illustrative numbers for the LLM and Jev. The test week has only about four hundred real attacks, so small differences in AUC may be due to chance. The real value of a bake-off is the *method*: run the same comparison on your own data, with the real models.

SET THE THRESHOLD

A new team will handle cloud storage alerts, a decision nobody has labelled yet. They will get about 300 labelled cases per week. Which method should they start with, and when should they reconsider?
(Start with a method that needs no labels. Use the first 300 labelled cases to check and fix its calibration. After a few weeks, train a logistic regression next to it. If both agree, keep the cheaper one. If they disagree, study the cases where they disagree.)

KEEP THESE

- Labelled data is the hidden cost of trained models. A method that needs no labels can start today.
- If a simple rule or logistic regression can make the decision well, use it.
- There is no single best method. There is only the best method for a particular decision.

EXERCISES

1. Make a scorecard for a decision in your own work. Which three columns matter most?

Next chapter: a probability is not an action. We set up three doors, act, review and escalate, with thresholds based on costs.

CHAPTER 11

Act, review, or escalate

An alert comes in at 2 a.m. The model says 0.37. We know this is a probability, and we know how to check it. But now the team asks the real question: *what should we do with it?* Close it? Keep it for the morning? Call someone immediately? This step, from a number to an action, is the least exciting part of an AI system, but it is where most of the money is saved or lost.

BY THE END OF THIS CHAPTER

- Set the act and escalate thresholds using costs.
- Adjust the policy to the number of people you actually have.
- Run the policy safely, with random checks, logs and daily monitoring.

Why 0.5 fails

Most tutorials say: above 0.5, treat it as an attack; below 0.5, treat it as safe. Let us test this on one realistic week in our example company. About 689 alerts per day get closed automatically. That sounds great, until we see that every day about 37 of these closed alerts were real attacks, and nobody looked at them.

The 0.5 rule makes two wrong assumptions. It assumes a missed attack and a false alarm cost the same. And it assumes there are only two possible actions. In fact, there are three.

Three doors

Every case can leave through one of three doors:

- **Act**: the system handles it alone, and no person reads it. Most cases should go this way, because most cases are harmless.
- **Review**: a person checks it during normal working hours. This takes about twelve minutes of a skilled person's time.

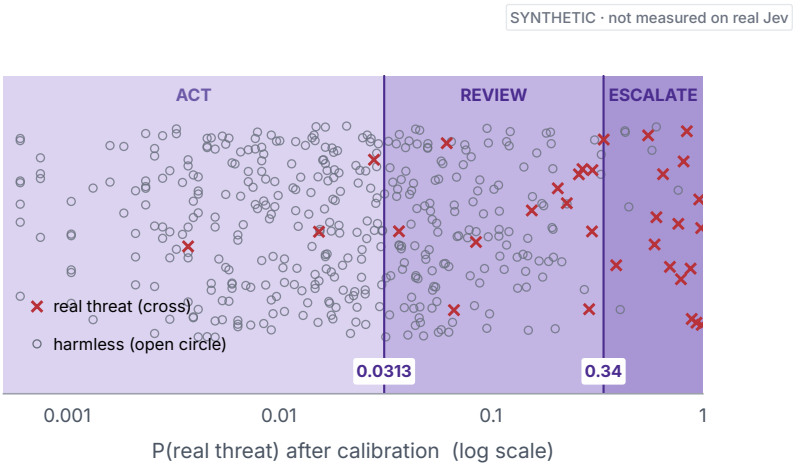


Figure 11.1. 420 alerts from the live week, placed by how likely the (calibrated) model thinks they are to be real threats. Note the log scale: most alerts are crowded at the far left, where the model is confident they're harmless. The red crosses, real threats, are mostly to the right, but not all of them.

- **Escalate:** someone is called immediately, even at 2 a.m., because it cannot wait.

The same idea works outside security. In customer support, the three doors are an automatic reply, a support agent and a supervisor. In checking social media posts, they are automatic approval, a human reviewer and the legal team.

> ● KEY IDEA
>
> **Act, review, escalate: the machine handles the clear cases at both ends, and people handle the unclear middle.**

In Figure 11.1, most harmless alerts are on the left, and most real attacks are on the right. But some attacks are scattered in the middle, and a few hide among the harmless ones. No two lines can put every alert in the right place. Our job is to make sure the mistakes we cannot avoid are the cheap ones.

Let costs decide the thresholds

For each door, we write down its average cost when the chance of a real attack is P. These are example estimates for a mid-sized company; the method matters more than the exact numbers.

- **Act.** If the alert was real, an attack got through. Expected loss about $10,000. Cost: $P \times \$10{,}000$.
- **Review.** Twelve minutes at $75 per hour is $15, and reviewers miss about one attack in twenty. Cost: $\$15 + P \times 5\% \times \$10{,}000$.
- **Escalate.** Calling someone for nothing costs about $400. Cost: $(1 - P) \times \$400$.

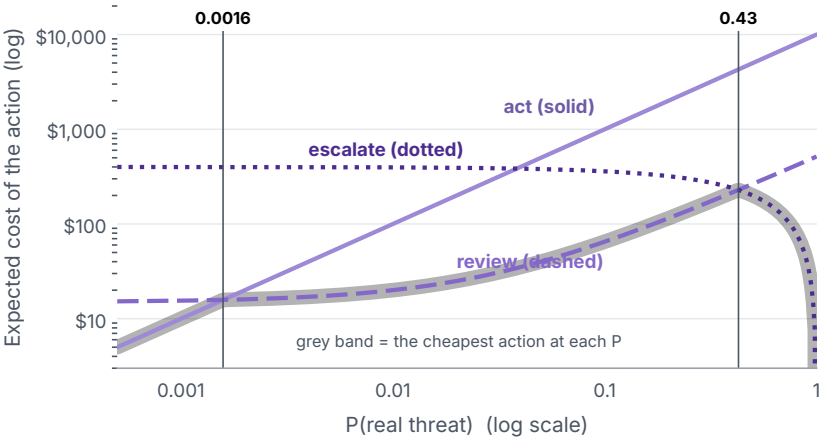


Figure 11.2. Expected cost of each action as the probability of a real threat rises (both axes on a log scale). Acting is cheapest when P is tiny, escalating when P is large, and review wins everywhere in between. The two crossing points are the thresholds.

For each value of P, we choose the cheapest door. The two points where the lines cross in Figure 11.2 are our thresholds, and we can find them with simple arithmetic: act when P is below about 0.0016, and escalate when P is above about 0.43.

Look at the first number again. The machine may close an alert by itself only when it is *99.8% sure* the alert is harmless. This is not over-caution; it comes directly from the costs. A real attack that nobody sees costs about 630 times more than one review.

Check the probabilities first. These lines assume that P is correct. In our example, for alerts the raw model called "below 2%", it said 0.6% on average, but the real rate was 1.4%, which is 2.5 times higher. Most models are overconfident at the extremes. Platt scaling on past data brings them close to the truth.

● KEY IDEA

Never set a threshold on a probability you have not checked against recent real outcomes.

The real limit: the number of people

If we use only the cost-based thresholds, about 638 alerts per day go to review. That is a full day's work for about 16 people. But the team has only 6 people, so it can review about 240 alerts per day. Anything beyond that is simply not looked at, and nobody chooses which ones are skipped.

● KEY IDEA

A policy that ignores how many people are available is only a wish.

So we ask a better question: *among the policies that fit the team we actually have, which one costs the least?* With such a policy, tested on a new week, the machine closes 432 alerts per day, people review 243, and 44 are escalated. The number of real attacks closed without anyone seeing them falls from 37 per day to about 6.1.

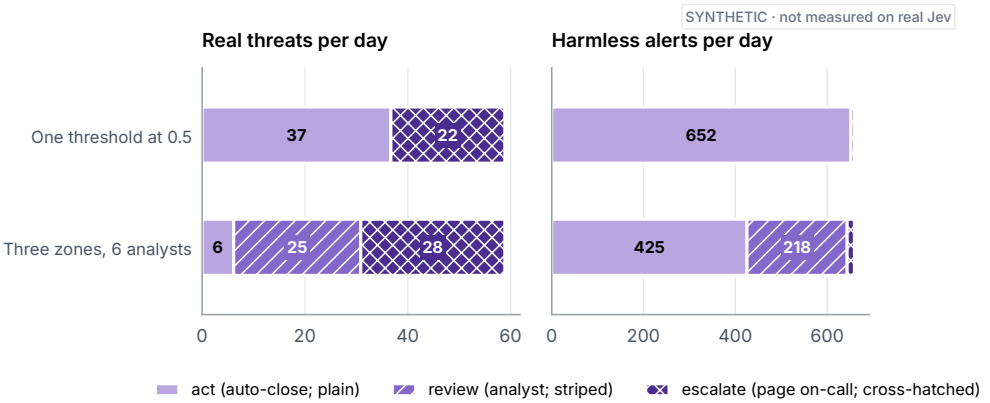


Figure 11.3. Where the live week's real threats (left) and harmless alerts (right) end up each day, under one threshold at 0.5 and under the three-zone policy that respects capacity. The light, plain segment on the left (act) is the number that matters: threats closed with no person involved.

About six per day is still not zero. Now think like the manager who approves the budget. Figure 11.4 shows the best policy for every team size. Going from six people to eight reduces unseen attacks from 6.1 to 2.7 per day, which avoids about $34,000 of expected loss every day. The model does not make this choice. But now the choice has a clear price, instead of being a guess. The lower curve also shows that giving the model cleaner input data helps as much as adding more people.

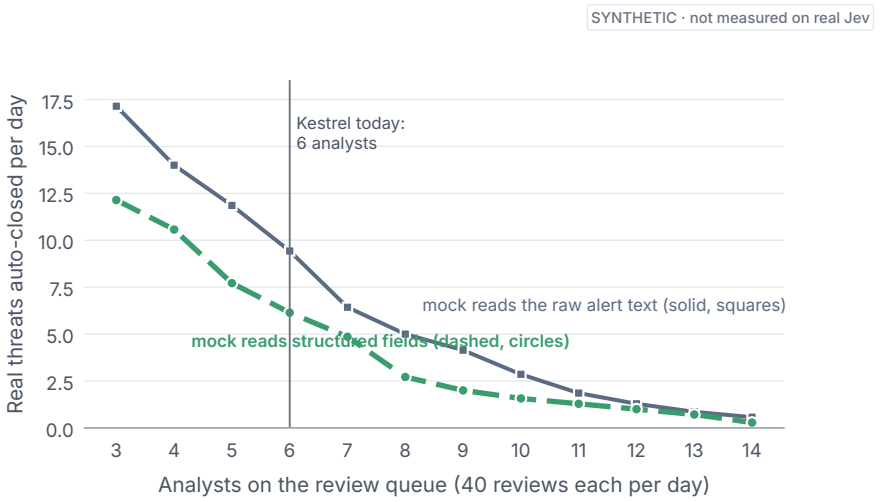


Figure 11.4. Threats auto-closed per day, for every team size from 3 to 14 analysts, with thresholds fitted again for each. The gap between the two curves is what better input is worth: the same model reading structured alert fields instead of raw alert text.

Two more rules

What is at risk matters, not only how likely. A model can be 99% sure that an alert on the payroll server is harmless, but the remaining 1% is still on the payroll server. So nothing on a critical system should close automatically, however confident the model is. And any case that does not clearly belong to another door should go to a person. When in doubt, review.

Randomly check what the machine closed. Cases sent to people come back with an answer. Cases closed by the machine come back with *nothing*, so you never see the machine's own mistakes, exactly where they hurt most. So send a small random share, say 3%, of automatically closed cases to a person anyway.

> **KEY IDEA**
>
> **Cases closed by the machine never tell you if they were wrong, unless you check a random sample.**

Running it every day

Write the policy down with a version number, for example: "Version 3: act below 0.031, escalate above 0.34, set on 1–21 September, six people". Then it can be reviewed, and changed back if needed. Record every decision:

what the model saw, what it said, and which policy turned that into an action. If the model fails or does not answer in time, send the case to review, never to automatic closing.

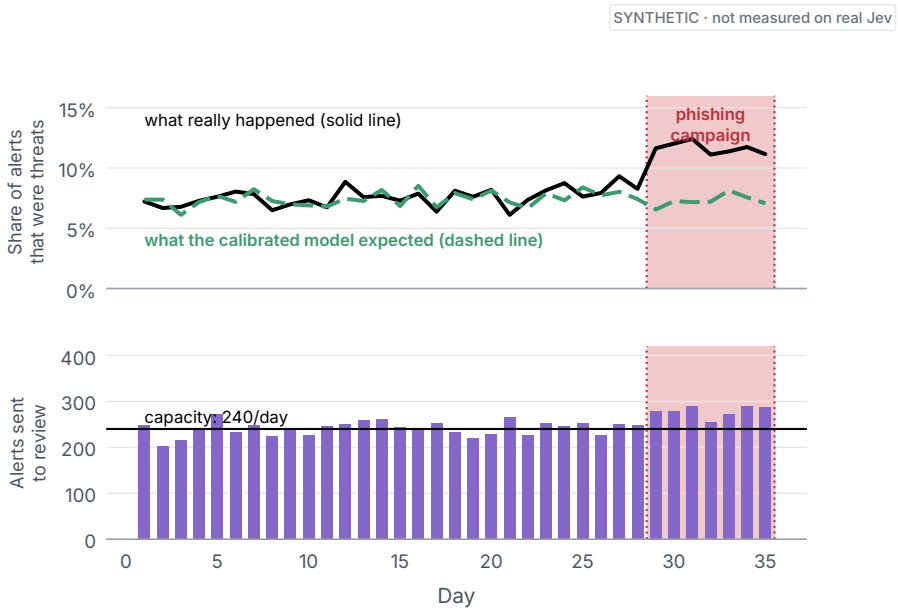


Figure 11.5. Top: the share of alerts that were real threats each day (black, solid line) against what the calibrated model expected (green, dashed line). Bottom: alerts sent to review each day against the team's capacity. In week five (shaded, between dotted lines), reality moves away from the model, and the queue overflows.

Then watch three numbers every day: the share of cases at each door, the review queue compared with the team's capacity, and the predicted number of attacks compared with the confirmed number. When an attack campaign starts in week five, no single alert looks unusual. But the model expected 7% of alerts to be attacks while the real figure was 12%, and the review queue overflowed (Figure 11.5). Both problems are clear on a daily chart. Then a person decides what to do: re-calibrate with the latest week, tighten the act threshold for that type of alert for some time, or add more reviewers.

⚠ **WHERE THIS BREAKS**

Thresholds are set using past data, but attackers act in the future. An attacker who learns what makes the model relax will try to make their activity look exactly like that. Random checks, keeping the model's scores private, and never auto-closing critical systems make this harder, but not impossible. The cost estimates also need to be reviewed whenever the business changes, not only when the model changes.

KEEP THESE

- Act, review, escalate: the machine handles the clear cases at both ends, and people handle the unclear middle.
- Never set a threshold on a probability you have not checked against recent real outcomes.
- A policy that ignores how many people are available is only a wish.
- Cases closed by the machine never tell you if they were wrong, unless you check a random sample.

EXERCISES

1. Design an act/review/escalate policy for something outside security, such as refund requests or insurance claims. What does each door mean, which mistake is expensive, and what limits the review door?

Next chapter: common designs for using decision models in real systems, and an agent built from them.

CHAPTER 12

Decision patterns, and a hybrid agent

An experienced cook does not need a recipe for every dish, because most dishes use the same few steps: fry, add water, cook till thick. Software also has such common steps. We call them **patterns**: well-known solutions with a name, so that a team can simply say "put a gate before that action" and everyone understands. In this chapter we collect the patterns for using decision models, and then build an agent with them.

BY THE END OF THIS CHAPTER

- Recognise six common patterns for using decision models in real systems.
- Find a decision step that is wrong even when nothing crashes.
- Divide work correctly between a decision model and an LLM.

All six patterns in Figure 12.1 follow one rule.

● KEY IDEA

A decision model never writes and never acts by itself. It only answers a clear question, and normal program code takes the action.

Six patterns

1. Extract, then decide. The input is messy, like an email or a PDF, but the decision is small. LLMs are very good at reading messy text, but weaker at the decision itself. So we let the LLM convert the text into clean fields, and let a decision model decide using those fields.

In our example, extracting first improved the ranking (AUC) from 0.708, when the LLM decided alone, to 0.865. But the company's own clean data fields did even better (0.892), because the alert text never mentions things

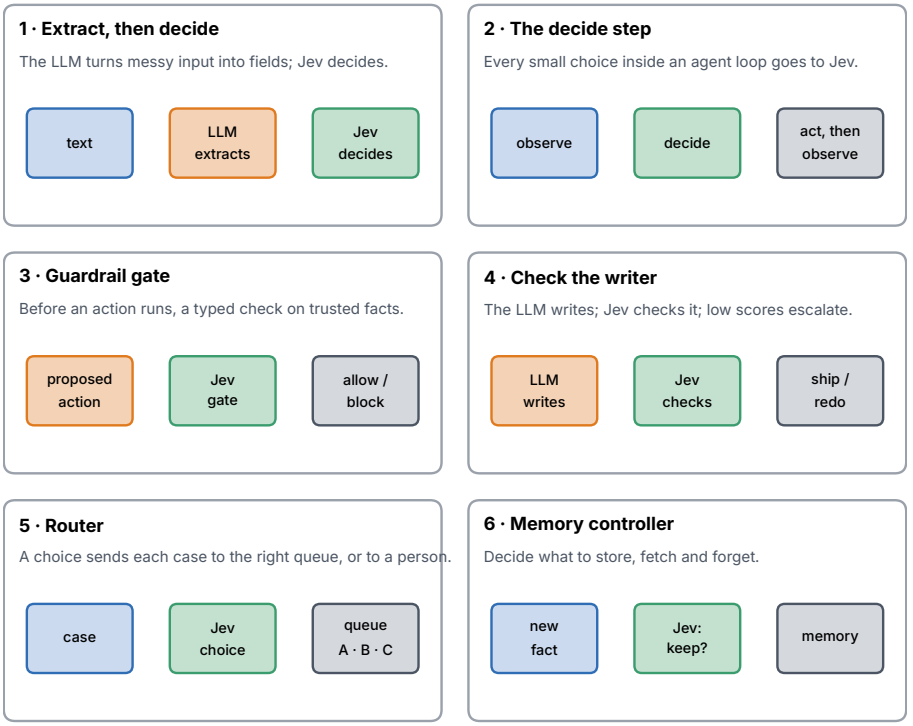


Figure 12.1. The six patterns in this chapter, as small diagrams. In every one, something else writes or acts, and a typed question decides.

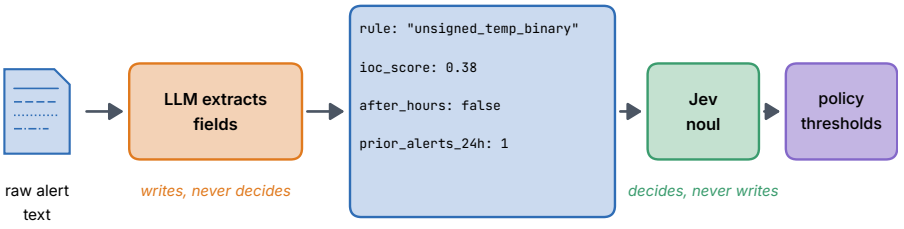


Figure 12.2. Extract, then decide. The LLM reads the raw text and writes structured fields; a decision model answers a typed question about them; policy thresholds turn the answer into an action. The LLM writes but never decides; the decision model decides but never writes.

like how important the computer is. No LLM can extract a fact that is not written. So if clean data already exists, use it.

2. The decide step. Most steps in an agent's loop are small decisions. Give these to a decision model as clear questions, and call the LLM only when something must be written. Since every decision comes with a probability, the agent can notice when it is unsure and stop, instead of continuing confidently.

3. The safety gate. Before any risky action, ask a clear question about it, using only facts the attacker cannot change. A high probability allows the action, a low probability blocks it, and anything in the middle goes to a person. These are the three doors again, now for actions.

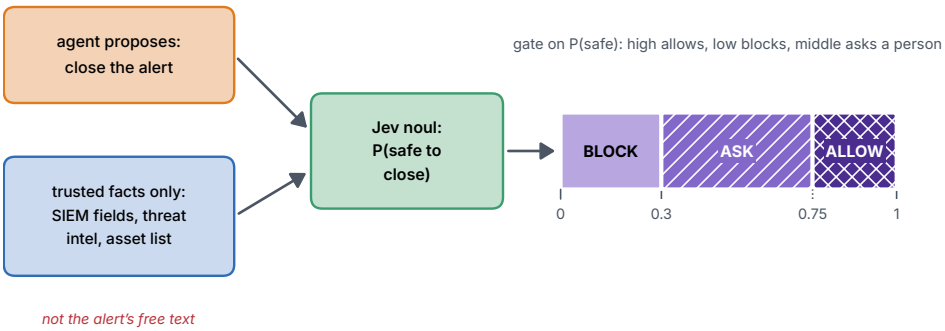


Figure 12.3. A guardrail gate. The agent proposes an action. A decision model estimates whether it's safe, from trusted facts only, never from text an attacker could write. High probability allows the action, low probability blocks it, and the middle asks a person.

4. Check the writer. Some outputs must be written, such as a report or a reply to a customer. After the LLM writes a draft, ask clear questions *about* the draft: is every statement supported by the evidence? Is the tone suitable? Then send it, return it once for correction, or give it to a person. Limit the number of corrections, or the loop may never end.

5. The router. To send each case to the right team, use a pick-one question, but send automatically only when the top option is clearly ahead. In our example, sending every alert by its top answer sent 88% to the right team. Sending automatically only when the top probability was at least 0.8 covered 80% of alerts, and 94% of those went to the right team. The rest went to a general queue for a person to sort.

6. The memory manager. Agents save notes and look things up. Each save and each lookup is a small decision: is this worth remembering? Where should I search? A published system called Jev-Mem uses a System One

model for these small decisions and calls a slower reasoning model only for the final answer; its authors report better and faster answers than the systems they compared it with (Jiang et al. 2026). Memory mistakes are silent: a fact that was wrongly thrown away is simply missing when needed. So keep a record of what gets thrown away.

A hybrid agent

Now let us design the security agent from Chapter 6 as a **hybrid**: a decision model makes every decision, and the LLM writes only what a person will read. Every run leaves a **trace**, a step-by-step record of what the agent saw and what each decision said, with its probability. A trace works as a debugging tool, an audit record and an explanation, all at once.

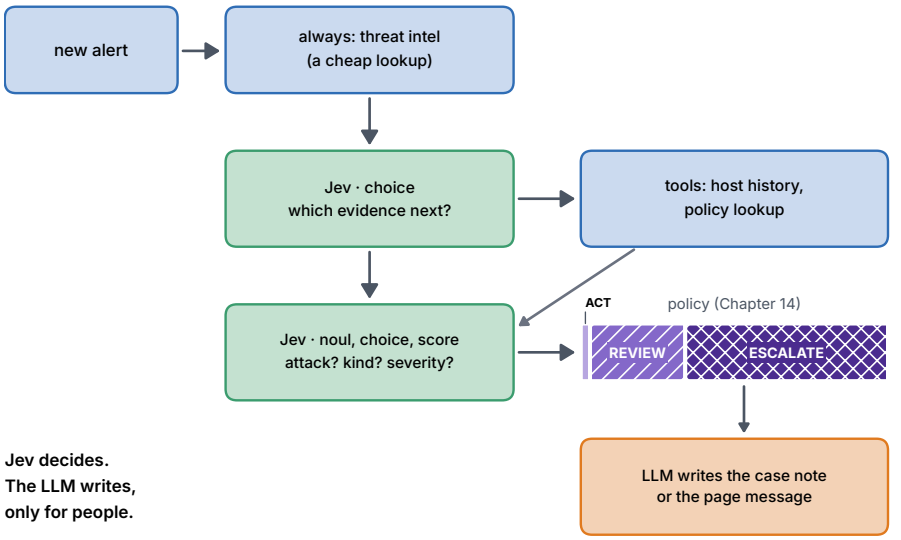


Figure 12.4. The hybrid agent. A cheap lookup always runs first. Jev decides which other evidence to collect, then answers the three verdict questions together. The policy turns the verdict into an action, and the LLM writes a case note or urgent message only when a person will read it.

A mistake with no error message

Suppose we run the agent for a week. Everything looks fine: no errors, every answer is readable, and the actions are spread across all three doors. Then we do one simple check that should be done for every decision step of every agent: *count how often each answer is chosen.*

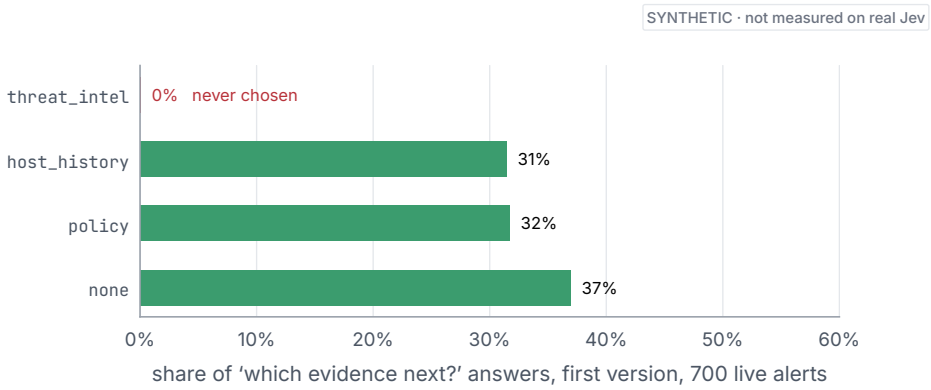


Figure 12.5. How often the first version of the agent chose each option at its "which evidence next?" step, across 700 live alerts. It never once chose threat intel.

In our example, across 700 alerts, the step "which evidence should I collect next?" *never* chose the threat-intelligence lookup, the most useful information about an alert (Figure 12.5). So all verdicts were made without it. The ranking dropped from 0.887 to 0.723, and about 22.6 real attacks per day were closed automatically.

A clean answer format does not guarantee a correct answer, and nothing later in the system will complain.

> ● KEY IDEA
>
> **Count how often each answer is chosen at every decision step. An option that is never chosen, or always chosen, is a warning sign, even without any error.**

The fix was not a better model. The threat lookup takes a fraction of a second and costs almost nothing. Deciding *whether* to do the lookup costs about the same as simply doing it.

> ● KEY IDEA
>
> **When checking something is cheaper than deciding whether to check it, just check it.**

Hybrid agent against an LLM-only agent

Now compare it with an agent where the LLM makes every decision, using the same steps, tools and thresholds.

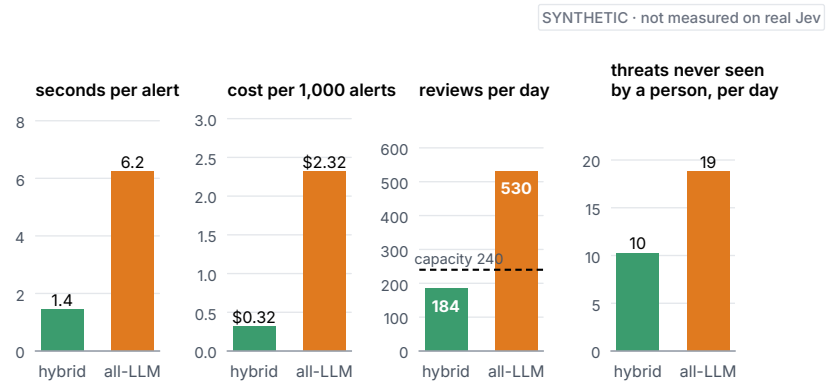


Figure 12.6. Four measures for the two agents on the same 700 live alerts, scaled to a day. The LLM's stated confidences sent far more alerts to review than the queue can take, so more real threats were never seen by a person.

In our example, the hybrid agent took 1.4 seconds per alert, while the LLM-only agent took 6.2. The bigger difference is the review queue (Figure 12.6). The LLM's written confidences are mostly a few round numbers, so the thresholds cannot separate cases well, and it sent about 530 alerts per day to a team that can review 240. The hybrid sent 184. When the queue is twice too long, the extra alerts are simply not seen. Counting both automatically closed attacks and attacks never reached in the queue, the hybrid missed about 10 per day, and the LLM-only agent about 19.

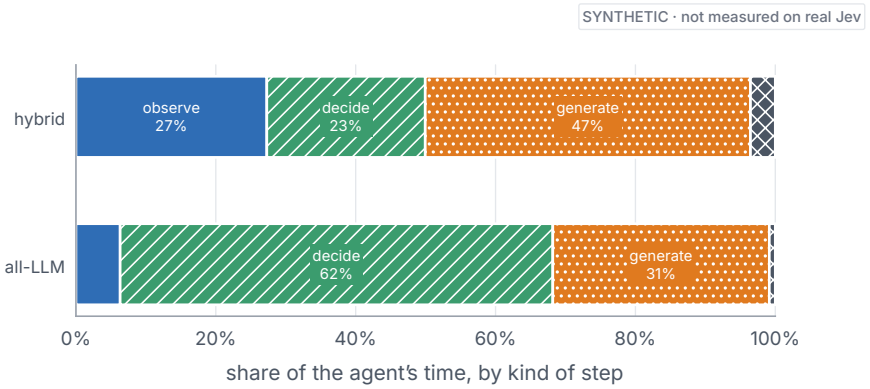


Figure 12.7. Share of each agent's total time spent on each kind of step. In the hybrid, the largest share is the LLM writing for people. In the all-LLM agent, most of the time goes on deciding.

Figure 12.7 shows where the time goes. The LLM-only agent spends most of its time deciding. The hybrid agent spends its biggest share on *writing* notes that people will read, which is exactly the work an LLM is best at.

● KEY IDEA

Use the LLM for the work that only an LLM can do: reading messy text and writing for people.

⚠ WHERE THIS BREAKS

All numbers in this chapter are illustrative, including the timings, and part of the difference between the two agents comes from how the example was set up. Treat the speed, the queue and the design lessons as the main findings. Also, combining several patterns adds more places to fail: an extractor, then a gate, then a router means three places where a mistake can happen.

● KEEP THESE

- A decision model never writes and never acts by itself. It only answers a clear question, and normal program code takes the action.
- Count how often each answer is chosen at every decision step. An option that is never chosen, or always chosen, is a warning sign, even without any error.
- When checking something is cheaper than deciding whether to check it, just check it.
- Use the LLM for the work that only an LLM can do: reading messy text and writing for people.

EXERCISES

1. Design a safety gate for an email agent that can send messages to customers. What is the clear question, which facts can be trusted, and where are the thresholds?

Next chapter: one team's complete journey from start to finish, and the same design in fields far from security.

CHAPTER 13

One team, start to finish, and beyond

So far, each chapter looked at one part of the problem. In this chapter we follow the whole journey, the way a real team would: how the team worked before, how the new system was tested quietly before it was allowed to act, how it was switched on, and what happened when an attack campaign hit. Then we leave security and see the same design in other fields.

BY THE END OF THIS CHAPTER

- Test a new decision system in shadow mode before trusting it.
- See why the order of a queue matters as much as its size.
- Recognise decision layers in banking, healthcare, customer support and more.

Before: one long queue

In our example, the security team gets about 718 alerts per day and can review 240. Before the new system, alerts flagged by the senior expert's rules went to the front of one queue, and all other alerts waited in the order they arrived. Any alert not reached within a day was closed without anyone looking at it.

This is a very common way of working, and it hides its own failures. Nobody sees the alerts that were closed unseen, so nobody knows how many of them were real. In our example we do know: about 27.1 real attacks per day were never seen by anyone.

Testing quietly: shadow mode

The new system uses only what we have already seen: clear questions (Chapter 8), calibration on past data (Chapters 3 and 9), three doors sized

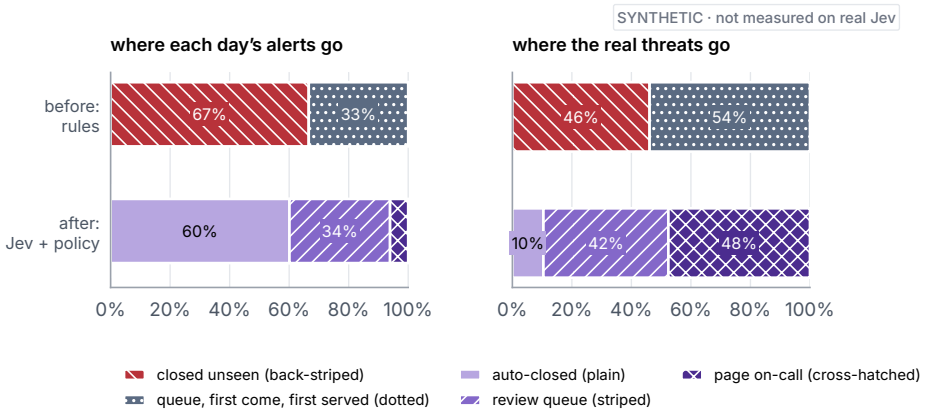


Figure 13.1. Where each day's alerts go, and where the real threats go, before and after. Before, most alerts aged out of the queue unseen, and so did a large share of real threats. After, most alerts are closed by the policy on purpose, and most real threats are queued or paged.

to the team (Chapter 11), and an LLM that writes only for people (Chapter 12).

But we do not switch on such a system and simply hope. First we run it in **shadow mode**: it sees every alert and records what it *would* have done, while people continue working as before. Nothing it decides has any effect yet.

> ● KEY IDEA
>
> **Before a new decision system is allowed to act, run it in shadow mode and compare its choices with reality.**

After the shadow week, the team discusses one main number: how many real attacks the system *would* have closed automatically. In our example it was about 6.1 per day. That sounds bad, until we compare it with the 27.1 per day that the old queue missed. It was also a little higher than predicted from past data. The honest thing is to say this clearly in the meeting, and watch the daily numbers to see whether it is chance or a trend. The team switched the system on, with random checks of closed alerts and a rule to send any failed case to review.

The order of the queue matters

The new review queue is sorted by probability, with the most likely attacks first. If we kept the old way, in order of arrival, the typical real at-

tack would wait about 158 minutes. Sorted by probability, it waits about 4 minutes. Same people, same alerts, same queue; only the order changed. The share of real attacks seen by a person within one hour went up from 37% to 83%.



Figure 13.2. How long it takes for a real threat to reach a person. Before, rule-flagged threats were seen quickly and the rest waited up to a day or were never seen. After, urgent calls reach someone in about fifteen minutes and the queue takes the likeliest threats first.

● KEY IDEA

A probability decides not only *whether* a person looks at a case, but also *which* case they look at first.

When an attack campaign hits

Then a phishing campaign started. Alerts about suspicious links came three times more often than usual, and many more of them were real. On the first day, the review queue was already over capacity. Also, the probabilities for link alerts were now too *low*, because they had been calibrated on weeks when such alerts were mostly harmless. The daily monitoring caught both problems on day two. On day three, the team re-estimated the base rate for link alerts, corrected those probabilities (Chapter 9), and approved overtime.

It helped. Over the week, real attacks never seen by anyone fell from about 144 to about 121 (Figure 13.3). But with correct probabilities, the queue became even longer than the overtime capacity, because many more alerts really needed checking.

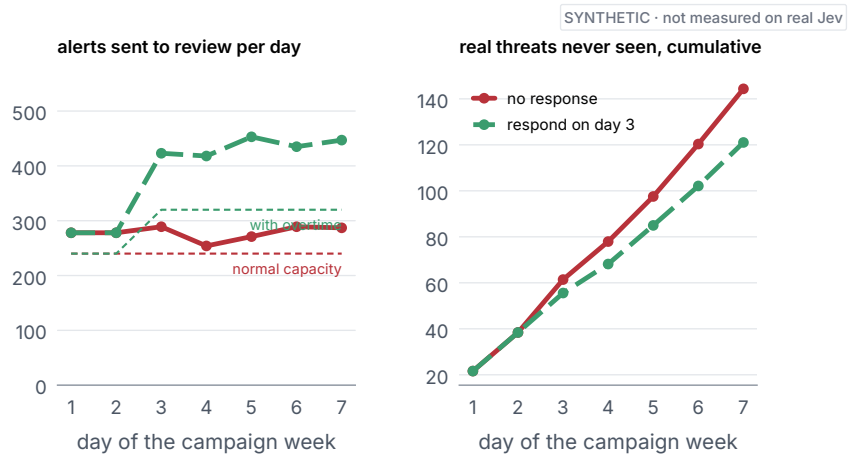


Figure 13.3. The campaign week, day by day, with and without a response on day 3. Left: alerts sent to review, against the team's capacity. Right: real threats never seen by a person, added up over the week. Recalibrating link alerts and adding overtime helped, but it didn't make the problem go away.

● KEY IDEA

Good calibration tells you honestly how much work there is. It cannot reduce the work.

The team was never the problem. The old queue decided what they saw by accident; the new system decides it on purpose.

Beyond security

Security was only our teaching example. Once you understand decision layers, you can see them everywhere.

Figure 13.4 describes six examples in the same way: what the model reads, the clear questions it answers, the "none of these" option, and who makes the final decision. What changes from card to card is *who decides*.

Where decisions are many and each mistake is cheap, like checking social media posts or card payments for fraud, no team can look at everything. So the model decides, and people check a sample (Figure 13.5). Where decisions are few and each mistake is costly, like a hospital deciding which patient needs care first, the model should never close a case. Its job is to put the right patient at the top of the list, and a doctor or nurse decides.

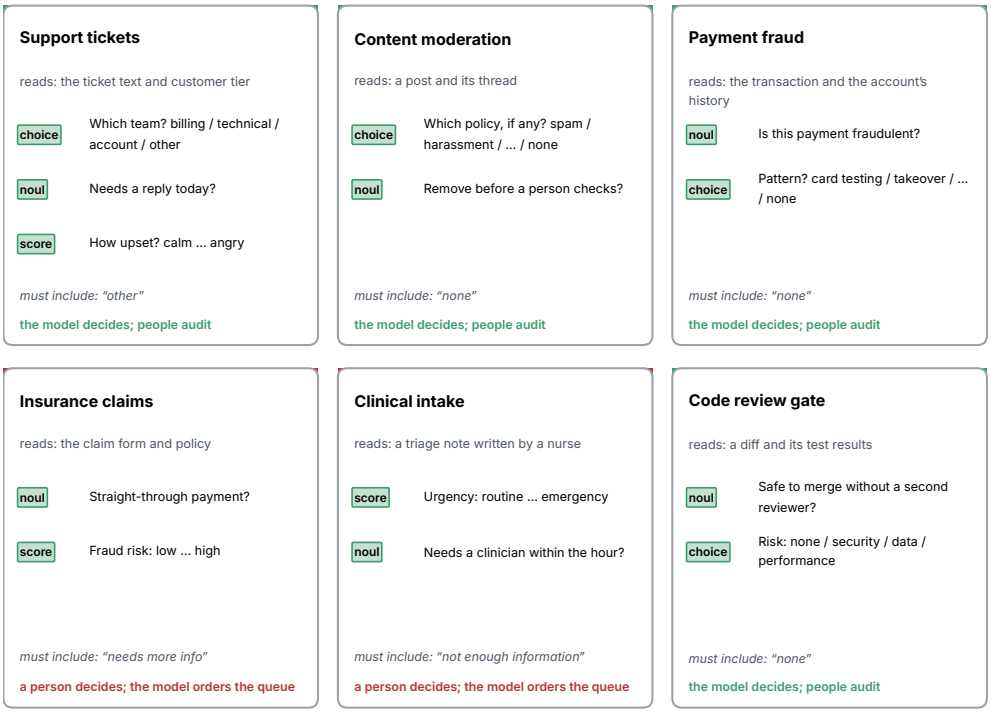


Figure 13.4. Six decision layers, each described the same way: what the model reads, the typed questions it answers, the option every choice must include, and who makes the final call. Green cards, which end "the model decides; people audit", let the model decide. Red cards, which end "a person decides", keep a person as the decider.

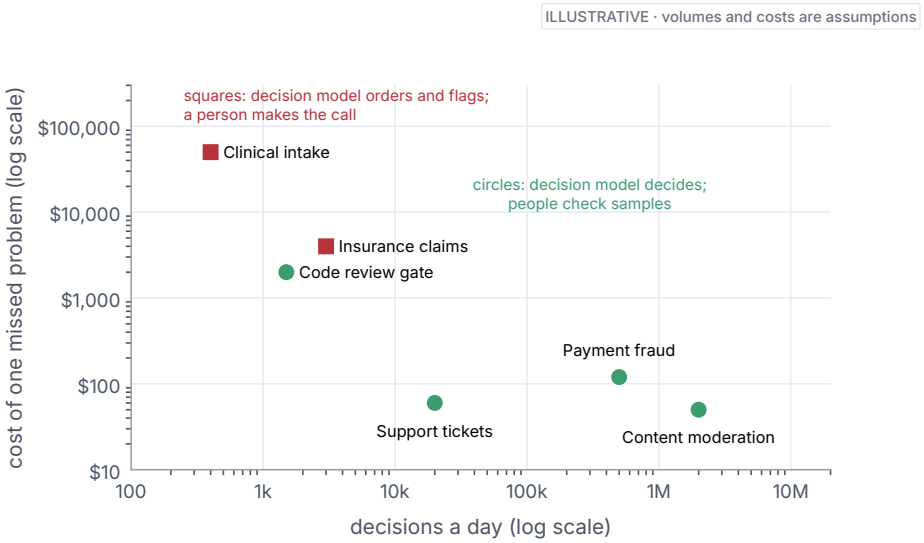


Figure 13.5. The six domains placed by how many decisions a day they involve and how much one missed problem costs (illustrative). Where volume is high and each mistake is cheap, the model decides and people check samples. Where each mistake is expensive, the model orders the queue and a person makes the call.

KEY IDEA

When decisions are many and cheap, the model decides. When they are few and costly, a person decides.

The threshold formula stays the same; only the costs change. A late reply to a support ticket costs a little customer goodwill, so the threshold may be near 5%. A missed emergency at a hospital costs far more than seeing someone early, so the threshold falls to about 0.3%. Such a low threshold needs probabilities that are correct even below one percent.

Always test before you trust. In our example with 400 support tickets, the same model that worked well on security alerts sent only 74% of tickets to the right team. Account problems went to the right team only 54% of the time, because “I can’t log in” uses similar words to “the app is not working”. Sending automatically only when the top answer was above 0.6 covered 55% of tickets, with 85% correct.

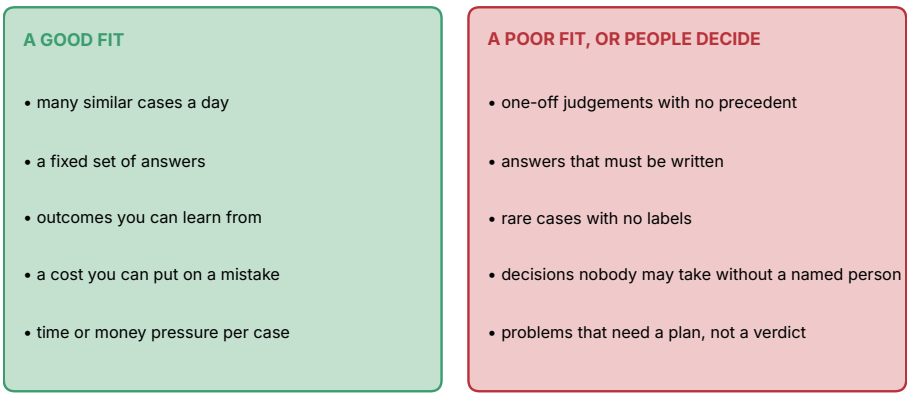


Figure 13.6. Signs that a decision suits a System One model, and signs that it doesn’t, or that a person must make it.

Not every problem needs a decision layer. It fits well when there are many similar cases, a fixed set of answers, outcomes you can learn from, and a cost you can put on each mistake (Figure 13.6). It does not fit one-time judgements, answers that must be written, or problems that need a plan rather than a verdict.

⚠ **WHERE THIS BREAKS**

All numbers here are examples, and real teams differ in many ways. Shadow mode is only as good as the outcomes you compare it with. And "a person decides" protects people only if that person has time to think. A doctor who must approve four hundred suggestions an hour is not really deciding. Watch the workload of the people in the loop as carefully as you watch the model.

● **KEEP THESE**

- Before a new decision system is allowed to act, run it in shadow mode and compare its choices with reality.
- A probability decides not only *whether* a person looks at a case, but also *which* case they look at first.
- Good calibration tells you honestly how much work there is. It cannot reduce the work.
- When decisions are many and cheap, the model decides. When they are few and costly, a person decides.

EXERCISES

1. Choose a decision from your own work and write its card: what the model reads, the clear questions, the "none of these" option, and who decides.

Next chapter: what it takes to run a decision model as a reliable service, and what is inside a small System One model.

CHAPTER 14

From model to service

In every project there comes a day when the prototype works and someone asks: “Can we use this from Monday?” A model that makes good decisions on a laptop is not yet a service. A real service must answer at 3 a.m. even when the network is slow, explain its decisions to an auditor a year later, warn someone when the situation changes, and change safely when it is improved. This chapter explains what that needs, and then looks inside a small System One model.

BY THE END OF THIS CHAPTER

- Use a model you do not own safely: fixed versions, retries and failures.
- Make every decision explainable with a record.
- Understand what is inside a small System One model.

Using a model you do not own

Test with a practice version first. Most AI services give an official software library. The library prepares the request, retries when there is a problem, and reads the answer. Only its last layer actually talks to the internet. If we replace that last layer with a practice version that answers locally, we can test all our logic for free, and then switch to the real service by changing one setting.

Fix the model version. A name like “latest” can point to a new model at any time, without warning.

● KEY IDEA

Fix the model version, so that any change in behaviour is your decision, not a surprise.

Retry only the right errors. Your own mistakes, such as a wrong key or a wrong model name, will fail the same way every time, so fix them instead

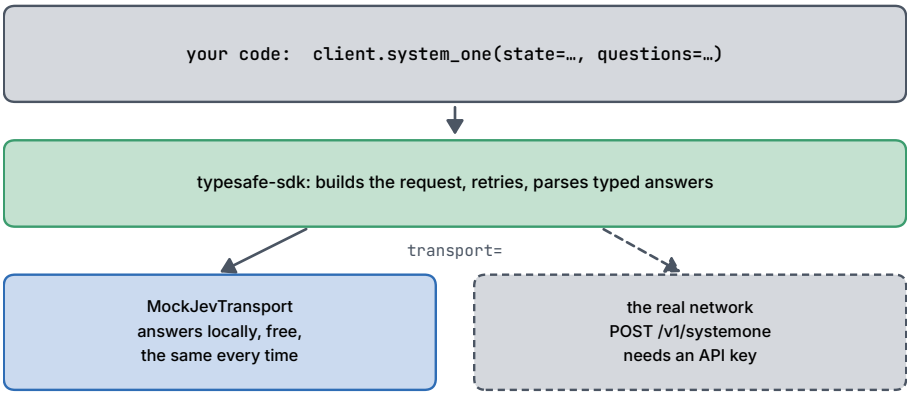


Figure 14.1. The official library does all the work of building requests, retrying and parsing answers. The transport decides where requests go. A practice version answers locally; the default sends them over the network to the real service.

of retrying. Problems outside your control, such as a busy server or a broken connection, are worth retrying. In our example, if one request in five fails, only 79% of calls succeed with no retry, but 99.2% succeed with two retries. Retries take time, though, and a login check may have only 300 milliseconds.

Decide in advance what a failure means. Treat a failed call as one more kind of case. For a login, the right rule may be "allow it, but mark it for review". For closing an alert, it should be "do not close; send to a person". Write this rule in the policy, where everyone can see it.

> ● KEY IDEA
>
> **When time is short, decide in advance what a failed call means.**

Ask questions together. In our example, four questions in one request used about 179 input tokens, and the same four questions in four separate requests used 343. One request is also faster.

The service

You will recognise every part of Figure 14.2 from earlier chapters. Four habits turn these parts into a reliable service.

One settings file, with a version. The model version, the calibration numbers, the thresholds and every business rule are kept in one place, and every decision is stamped with that version.

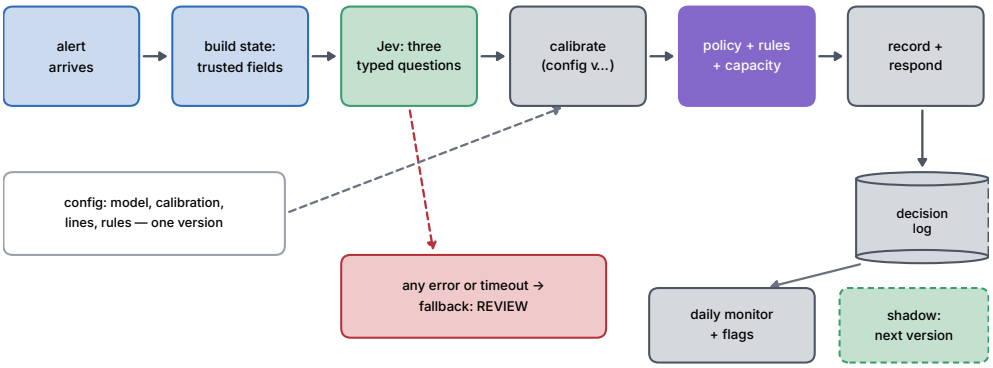


Figure 14.2. The decision service. Each alert's trusted fields go to Jev as three typed questions; the answer is calibrated, and the policy, business rules and capacity turn it into a zone. Every decision is recorded before it is returned. Any failure falls back to review. A daily monitor reads the log, and a candidate version can run in shadow beside the current one.

● KEY IDEA

Everything that can change a decision is kept in one place, with a version number.

A record for every decision. A decision model gives probabilities, not reasons. But if we record the raw and calibrated probabilities, the thresholds, the rule that was applied and the version, then every *action* can be fully explained, even though the model itself explains nothing.

Thresholds and rules set together. In our example, a test run over one week showed that the review queue was over capacity from the first day. The thresholds had been set for 240 reviews per day, but *after* that, someone added a sensible rule: never close an alert on a critical system automatically. Every alert this rule moved to review was extra work that the thresholds had not planned for. When the thresholds were set again *with* the rule included, the average queue went from about 275 to 245 per day.

● KEY IDEA

Thresholds and business rules are one decision, and must be set together.

Daily monitoring, a responsible person, and shadow testing. Every day, check four numbers: the share of cases at each door, the review queue

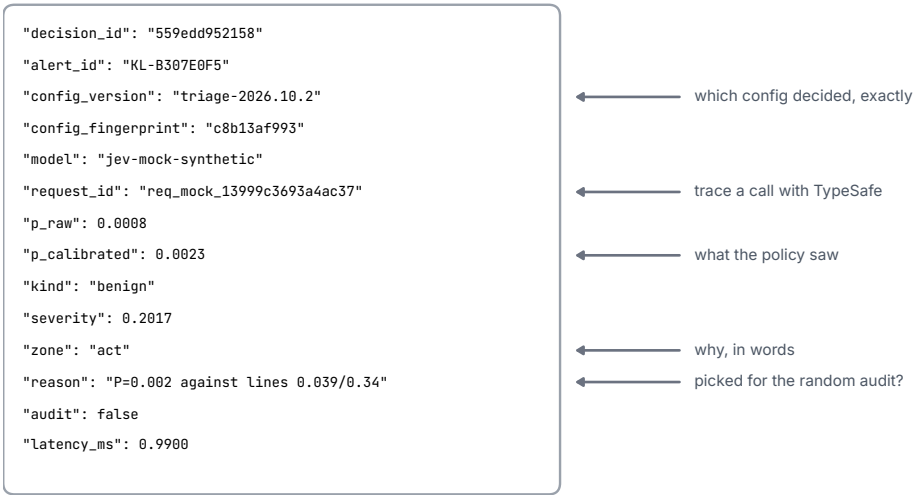


Figure 14.3. One decision record, with notes on why each field is there. Together they say what the service saw, what the model said, what the policy did, which version decided, and why.

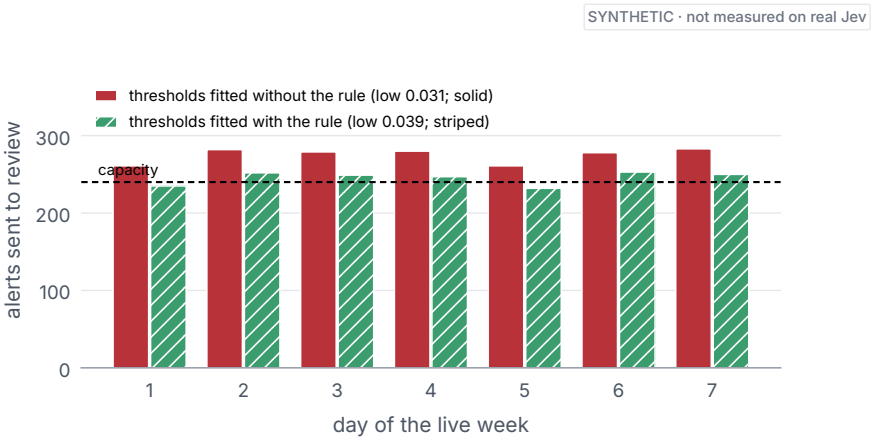


Figure 14.4. Alerts sent to review each day of the live week, before and after the fix. Fitted without the critical-asset rule, the thresholds let the queue run over capacity every day. Fitted with it, the act threshold rises slightly and the queue fits.

compared with capacity, how often calls fail, and predicted attacks compared with confirmed attacks. When a campaign starts, the last number shows the problem on day one, in plain numbers: the probabilities are no longer correct. A named person should have the authority to switch the service to "review everything" while the cause is found. And every improvement should first run in shadow mode beside the current version. In our example, the new version agreed with the old one on 96% of alerts, and the shadow test showed, before anyone paid the price, how many attacks its changes would let through.

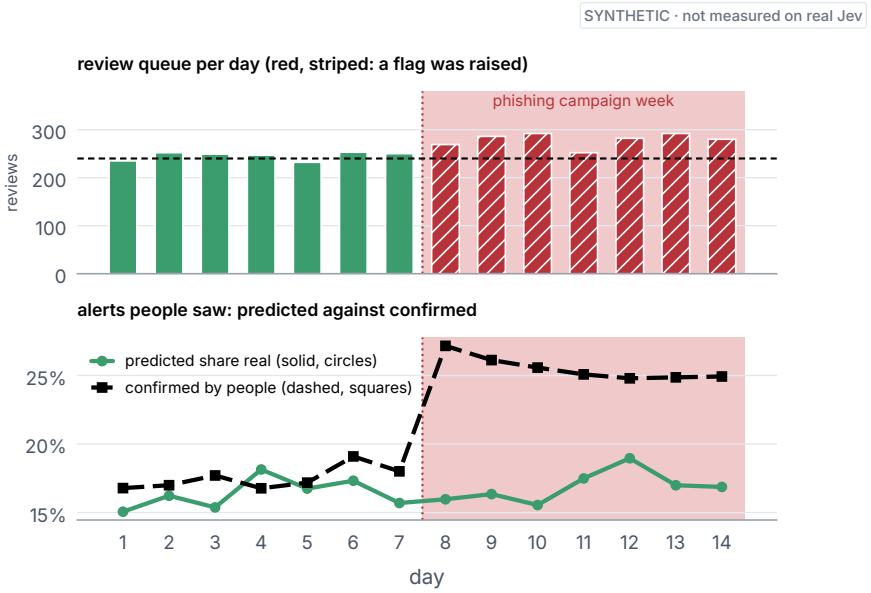


Figure 14.5. Two weeks of the daily monitor: the live week, then the phishing campaign. Top: the review queue, red and striped on days a flag was raised. Bottom: for alerts people saw, the share the model predicted to be real against the share they confirmed. The campaign opens a gap on its first day.

What is inside a small System One model

You do not need to build a System One model to use one. But it helps to understand what one looks like inside.

Imagine a small model built for alerts. One shared part, the **encoder**, reads an alert once and turns it into a short list of numbers. Three small parts, called **heads**, read those numbers, one for each question type (Figure 14.6). The yes/no head uses the S-curve from Chapter 2. The pick-one head makes probabilities that add up to 1. The scale head uses one "dial" with ordered cut points, so that "high" always stays above "medium". All

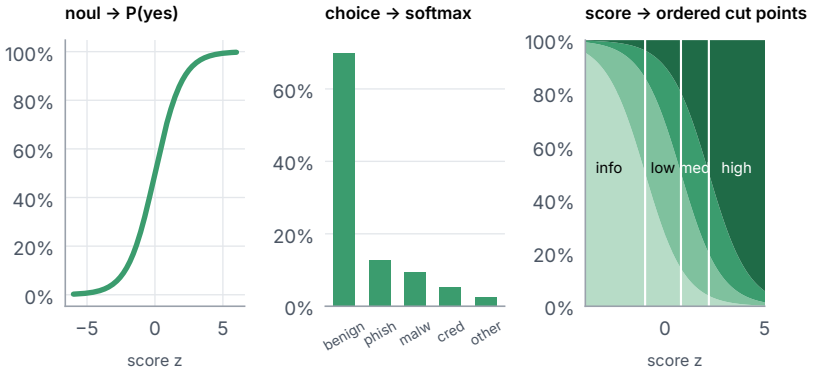


Figure 14.6. The three heads. A noul turns one number into P(yes) with the S-curve. A choice turns one number per label into probabilities that add up to 1. A score turns one number into probabilities for levels in order, by cutting it at points that must stay in order.

three heads are trained together with log loss, which rewards honest probabilities, and then each head gets a temperature correction on unseen data (Chapter 3). Such a model can be very small, about 2,091 numbers.

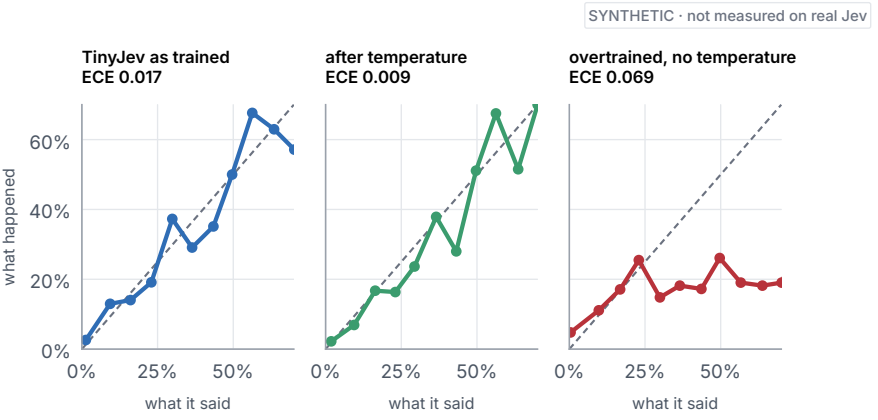


Figure 14.7. Reliability of P(attack) on the live week. Left: the small model as trained, slightly overconfident. Middle: after dividing by the fitted temperature. Right: the overtrained network before any temperature; its problems are ones no single temperature can fix.

In our example, the temperature correction reduced the calibration error from 0.017 to 0.009, without changing the ranking. A much bigger version trained for too long was a warning: its error was 0.069, and after a temperature correction it was no better, 0.077 (Figure 14.7). A correction can soften answers, but it cannot bring back what was lost by memorising.

● KEY IDEA

Train with an honest scoring rule, stop before the model memorises, and then calibrate on data it has never seen.

⚠ WHERE THIS BREAKS

A real service also needs login security for its users, limits on usage, safe storage for its records, control over who can change the settings, and protection for any personal data. Shadow tests show only where two versions *disagree*; if both are wrong in the same way, only random checks will catch it. And a small model trained on one company's alerts will drift as soon as the situation changes.

● KEEP THESE

- Fix the model version, so that any change in behaviour is your decision, not a surprise.
- When time is short, decide in advance what a failed call means.
- Everything that can change a decision is kept in one place, with a version number.
- Thresholds and business rules are one decision, and must be set together.
- Train with an honest scoring rule, stop before the model memorises, and then calibrate on data it has never seen.

EXERCISES

1. For an online service you depend on, list three errors you would retry and three you would fix. For each decision it makes, what should a failed call mean?

Next chapter: what changes, for the industry and for you, when decisions become cheap, fast and calibrated.

CHAPTER 15

What changes now

Predictions about AI often turn out wrong, so this last chapter stays with three things: what this book has shown, what it suggests, and what nobody knows yet.

BY THE END OF THIS CHAPTER

- See the decision layer as a separate part of an AI system.
- Understand why cheaper decisions can mean more work for people.
- Leave with a clear plan for your first decision.

A separate layer for decisions

For the last few years, the usual AI system has been one big model doing everything: reading, deciding, writing and using tools, all through text.

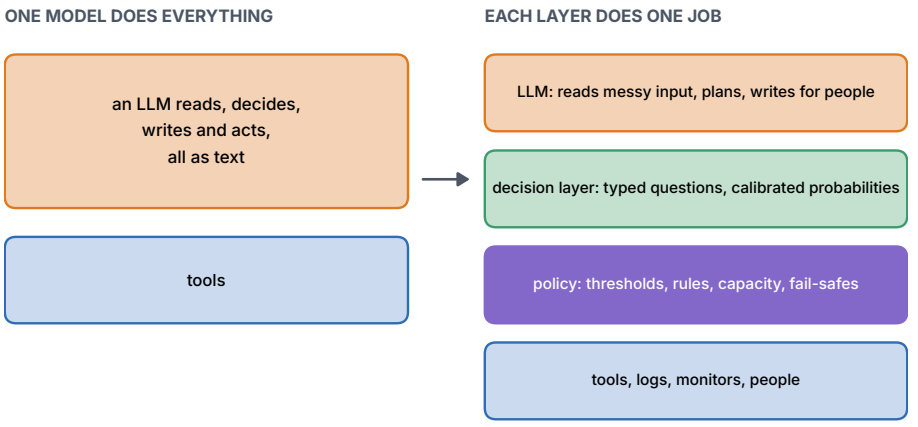


Figure 15.1. Two ways to build an AI system. Left: one model does everything as text. Right: each layer does one job. An LLM reads and writes, a decision layer answers typed questions with probabilities, a policy turns those into actions, and tools, logs and people sit underneath.

The right side of Figure 15.1 shows the design this book keeps reaching. The LLM does what only an LLM can do: read anything and write for people. The decision layer answers small, clear questions in one step, with probabilities we can test. The policy sets thresholds based on costs and team capacity. And at the bottom, tools take actions and people review. Each part is smaller than the whole system, can be tested on its own, and can be replaced without rebuilding everything else. This is the main reason this design is likely to last.

● KEY IDEA

Divide the system: the LLM reads and writes, the decision layer decides, the policy acts, and people own the costs.

Cheaper decisions, more decisions

In 1865, the economist William Stanley Jevons noticed something surprising about coal (Jevons 1865). James Watt's steam engine produced much more work from each tonne of coal. People expected Britain to use less coal. Instead, it used much more, because cheaper coal power was now used for many more things. We have seen the same with mobile data in India: when data became very cheap, people did not just save money; they started watching videos all day. This pattern is called the **Jevons paradox**. The name Jev comes from Jevons, and this time the "coal" is decisions.

A decision is worth making only when the loss it prevents is bigger than its cost, and when it can be made in time. A company makes a few thousand careful decisions per day, but it *could* make millions: one for every email, every login and every line in its logs. Each of these is worth only a tiny amount to check. At the example LLM price of $0.00066 per decision, almost none are worth checking. At Jev's reported price of $0.000021, many are.

The interesting part is where the change shows up. For work a company already does, cheaper decisions mean many more decisions (in our example, 8 times more), but the total bill goes *down*. The bill goes up only when cheap decisions make *new* work worth doing, such as checking every file shared outside the company.

● KEY IDEA

Cheaper decisions rarely increase the bill for existing work. They increase it by making new work worth doing.

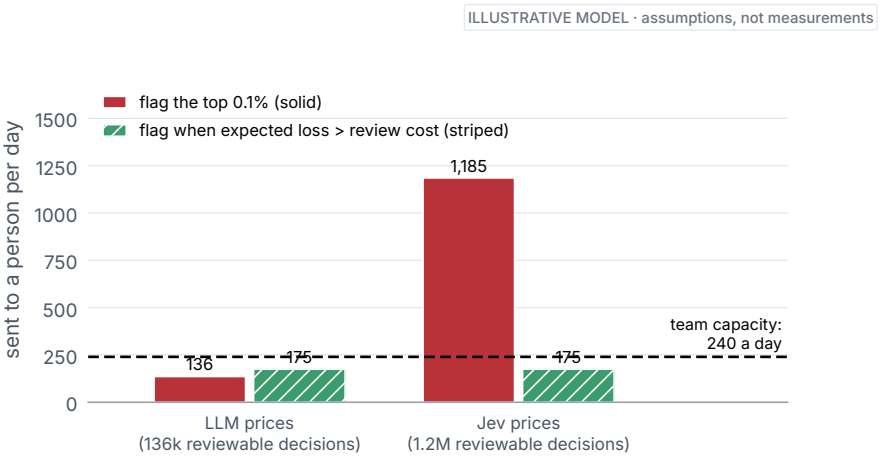


Figure 15.2. Alerts sent to a person per day, under two rules (illustrative). A rule that flags a fixed share of what the model looks at grows with the number of decisions, and at Jev volumes it sends far more than the team can handle. A cost-based threshold, which flags only when the expected loss is larger than the cost of a review, barely moves.

There is a hidden cost: more decisions also means more cases sent to people. A rule that seemed fine at low volume, such as "send the top 0.1% to a person", sends 136 cases per day, within the team's capacity of 240. At the new, higher volume, the same rule sends 1,185 per day, even though nobody changed it (Figure 15.2). A cost-based rule, which sends a case only when its expected loss is higher than the cost of a review, sends 175 per day in both cases.

● KEY IDEA

When decisions become cheap, set thresholds by cost, not by a fixed percentage.

What becomes valuable

If deciding becomes cheap, what becomes rare and valuable? Good judgement about the decision itself. Someone still has to decide what each mistake costs, where the thresholds should be, how many cases the team can review, and whether the probabilities are still correct this week. No model can do this for you. This is why the first chapters of this book, about probability, calibration and costs, are the foundation. They are the difference between a system that makes a million cheap decisions well and one that makes a million cheap mistakes.

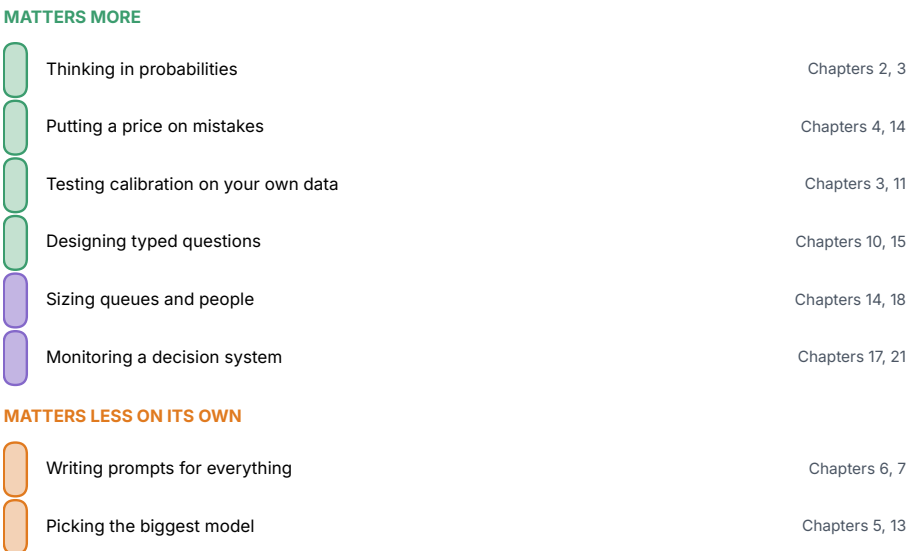


Figure 15.3. Skills that matter more when decisions get cheap, and skills that matter less on their own, with the chapters that teach each.

● KEY IDEA

When deciding becomes cheap, good judgement about the decision becomes valuable.

Many things are still unknown: how well any System One model is calibrated on your data, how it changes between versions, how it handles attackers, and what prices will finally be. None of these is a reason to wait. Calibration can be tested this week. Version changes can be handled with a fixed version and daily monitoring. Prices can be checked in a spreadsheet. And this design does not depend on any one company: a clear question in, a probability out, and a policy in your own code can be moved to a different model in a day.

● KEY IDEA

What you need is decisions you can measure, thresholds you can change, and a model you can replace.

What to do on Monday

If you do only one thing after reading this book, make a list of every decision your system makes, and count how often each one happens. Most

? **Does Jev's calibration hold on data like yours?**

Chapter 11 tells you how to find out.

? **How fast do System One models improve, and do they drift between versions?**

Pin versions; re-test.

? **Will prices stay this low?**

Early-access prices change. Chapter 12's model shows what depends on it.

? **Who else builds decision models?**

The interface pattern outlives any one vendor.

? **How will rules on automated decisions evolve?**

Keep people deciding where stakes are high.

Figure 15.4. Open questions, and what you can do about each today.

people are surprised by how many there are, and how many are small, clear and frequent.

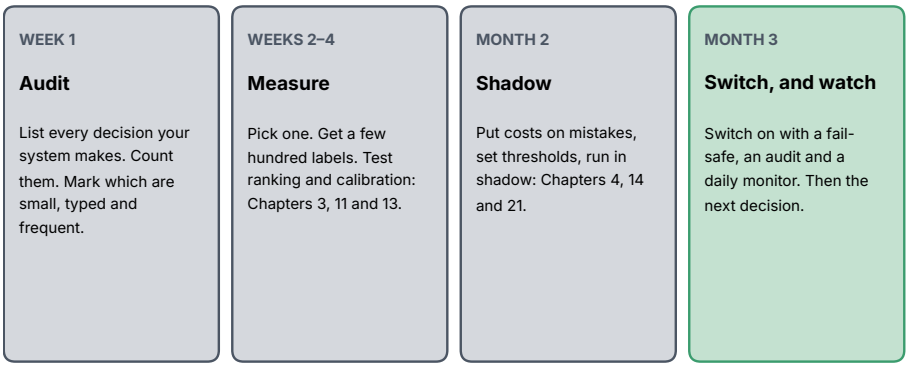


Figure 15.5. A first three months, one decision at a time.

Then choose *one* decision. Collect a few hundred cases with known outcomes and test the calibration. Put a cost on each type of mistake and set the thresholds. Run it in shadow mode. Switch it on with a rule for failures, random checks and daily monitoring. Then move to the next decision. This is not exciting work, but it is how every system in this book became better.

⚠ **WHERE THIS BREAKS**

This chapter is the least certain part of the book. The cost and speed numbers are examples, and Jev's numbers are as reported by the company during early access. The idea that a separate decision layer will become standard is a reading of the evidence, not a proven fact. It could be wrong if general models become fast and well calibrated enough. Even then, everything about costs, thresholds and calibration will still apply.

SET THE THRESHOLD

Take the most frequent decision in your own system. What does a wrong "yes" cost? What does a wrong "no" cost? Where should the threshold be? And how will you know, one month from now, whether the probabilities behind it are still correct?
(There is no answer at the back of the book for this one. If you can answer all four questions, you have learned what this book set out to teach.)

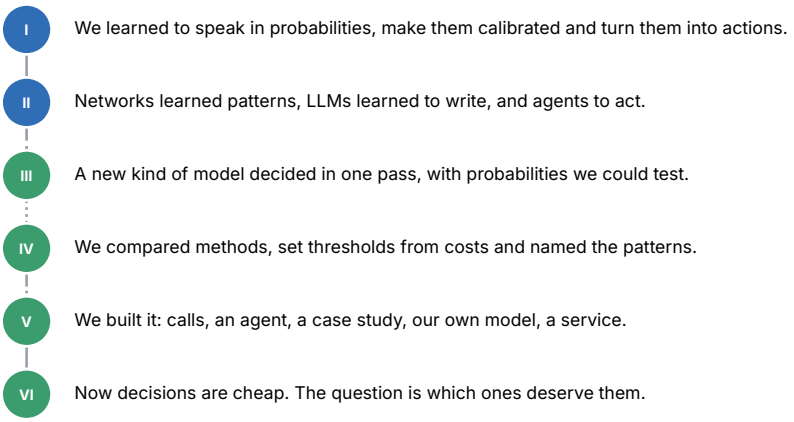


Figure 15.6. The whole book as one story.

Alan Turing ended his most famous paper with a line that suits this moment: "We can only see a short distance ahead, but we can see plenty there that needs to be done" (Turing 1950).

KEEP THESE

- Divide the system: the LLM reads and writes, the decision layer decides, the policy acts, and people own the costs.
- Cheaper decisions rarely increase the bill for existing work. They increase it by making new work worth doing.
- When decisions become cheap, set thresholds by cost, not by a fixed percentage.
- When deciding becomes cheap, good judgement about the decision becomes valuable.
- What you need is decisions you can measure, thresholds you can change, and a model you can replace.

EXERCISES

1. Make a decision list for one system you know: every decision, how often it happens, what a mistake costs, and who is responsible for it.

Thank you for reading.

Glossary

One plain sentence for each term, and the chapter where it's explained fully.

Act, review, escalate The three zones of a decision policy: handle it automatically, send it to a person, or raise it urgently (Chapter 11).

Agent A program that uses a model to observe, decide and act in a loop, calling tools along the way (Chapter 6).

Attention A way for each word in a sequence to look at the others and decide which ones matter for its meaning (Chapter 5).

Audit (random) Sending a random share of automatically handled cases to a person, so that mistakes the system is sure about still get seen (Chapter 11).

AUC The chance that a model gives a randomly chosen real case a higher score than a randomly chosen harmless one. It measures ranking, not calibration (Chapter 3).

Base rate How common something is before you look at any evidence, like the share of alerts that are real threats (Chapter 2).

Brier score The average squared gap between a predicted probability and what happened. Lower is better (Chapters 2 and 3).

Calibration A model is calibrated when the things it calls 70% likely happen about 70% of the time: its probabilities are honest (Chapter 3).

Capacity How many cases the people in the loop can handle in a day. A policy that ignores it isn't a policy (Chapter 11).

Cassette A file of recorded requests and responses that a test can replay without calling the real API (Chapter 14).

CISO Chief information security officer: the person in charge of an organisation's security (Chapter 13).

Choice A typed question with a set of named options. The answer is a probability for each option, adding up to 1 (Chapter 8).

Confidence In a Jev `choice` or `score` answer, a number from 0 to 1 computed from the shape of the probabilities. It's high when most of the probability is on one option, and low when it's spread out. It isn't simply the top probability (TypeSafe's quick-start shows 0.78 beside 0.85). The book's mock uses the top probability as a stand-in (Chapter 7).

Context window The amount of text an LLM can read at once (Chapter 6).

Cost-based threshold The probability above which acting is cheaper than not acting, set by what each kind of mistake costs (Chapter 4).

Decision layer The part of a system that turns probabilities into actions: typed questions, calibration, thresholds, rules and fail-safes (Part IV).

Decision record The logged account of one decision: what was seen, what the model said, what the policy did, which version decided and why (Chapter 14).

Drift A change in the world, or in the data, that makes a model's probabilities less accurate than when they were checked (Chapter 11).

DLP (data loss prevention) Security tools that warn when company data may be leaving the organisation, such as a large file sent to a personal email address (Chapter 3).

ECE (expected calibration error) The average gap between what a model said and what happened, across groups of similar predictions (Chapter 3).

Elasticity How much more of something people use when its price falls (Chapter 15).

Embedding A list of numbers that places a word, sentence or item in a space where similar things sit close together (Chapter 5).

Epoch One full pass of training over all the training examples (Chapter 5).

Escape option The "other", "none" or "not enough information" option that gives cases fitting nowhere somewhere to go (Chapter 8).

Extract, then decide A pattern in which an LLM turns messy input into fields and a decision model decides on those fields (Chapter 12).

F1 One score that balances how many flagged cases are real and how many real cases get flagged. It measures ranking and labelling, not calibration (Chapter 3).

Fail-safe The decided meaning of a failure, such as "if the call fails, send the case to review" (Chapters 11 and 14).

Fine-tuning Further training of an existing model on examples from one task (Chapter 5).

Gradient descent Improving a model step by step by nudging its numbers in the direction that reduces the loss (Chapter 2).

Guardrail gate A typed check on trusted facts that runs before an agent's action is allowed to happen (Chapter 12).

Hallucination An LLM stating something fluent and false (Chapter 5).

Hybrid agent An agent in which a decision model answers the small decisions and an LLM only writes (Chapter 12).

Isotonic regression A calibration method that learns any increasing mapping from scores to probabilities. It's flexible, but needs more data than Platt scaling (Chapter 3).

Jevons paradox When making something more efficient to use makes people use more of it in total, not less (Chapter 15).

JSON mode An LLM setting that makes its output valid JSON, though not necessarily correct JSON (Chapter 5).

Label The known right answer for an example, used to train or test a model (Chapter 1).

LLM (large language model) A model trained to predict the next piece of text, which lets it read and write almost anything (Chapter 5).

LLM-as-judge Using an LLM to rate or check an output, often on a numbered scale (Chapter 10).

Log loss A score that punishes confident wrong answers heavily. It's a proper scoring rule (Chapter 2).

Logistic regression A model that adds up weighted evidence and passes the total through an S-curve to get a probability (Chapter 2).

Loss A number that measures how wrong a model is on some examples. Training tries to make it smaller (Chapter 2).

Noul TypeSafe's name for a yes-or-no question. The answer is the probability of yes (Chapter 8).

Observe, Decide, Act The loop an agent runs: read something, choose what to do, do it, repeat (Chapter 6).

Ordinal Having a meaningful order, like severity levels. An ordinal model keeps that order in its probabilities (Chapter 14).

Overconfidence Probabilities that are more extreme than the evidence justifies. It's common in large models and in overtrained small ones (Chapter 5).

Overfitting Learning the training examples so closely that performance on new examples gets worse (Chapter 2).

Platt scaling A calibration method that fits an S-curve mapping a model's scores to calibrated probabilities (Chapter 3).

Policy The rules that turn a probability into an action, usually as thresholds between zones (Chapter 11).

Prior shift A change in the base rate between where a model was calibrated and where it's used. You can correct it if you know the new base rate (Chapter 9).

Prompt injection Text planted in an input that tries to make a model follow the attacker's instructions (Chapter 6).

Proper scoring rule A way of scoring predictions under which reporting your honest probability gives the best average score (Chapter 2).

RAG (retrieval-augmented generation) Giving an LLM relevant documents, found by search, before it answers (Chapter 6).

Reliability diagram A chart of what a model said against what happened. A calibrated model sits on the diagonal (Chapter 3).

RLCD TypeSafe's name for reinforcement learning for calibrated decisions, which it says trains Jev. The details aren't published (Chapter 7).

RLHF Reinforcement learning from human feedback, which trains an LLM towards answers people prefer (Chapter 5).

Router A pattern that sends each case to a destination with a choice question, falling back to a person when unsure (Chapter 12).

Score A typed question with ordered levels. The answer is a probability for each level and an expected level (Chapter 8).

Shadow mode Running a new system or version beside the current one, recording what it would do without letting it act (Chapters 13 and 14).

SIEM (security information and event management) The tool that collects a company's logs and raises security alerts. Kestrel's alerts come from one (Chapter 10).

Softmax The function that turns a list of scores into probabilities that add up to 1 (Chapters 5 and 8).

State What a decision model reads: the situation, as text or structured fields (Chapter 7).

Structured output An LLM answer constrained to a schema, such as JSON with named fields (Chapter 5).

System 1, System 2 Kahneman's names for fast, automatic thinking and slow, deliberate thinking (Chapter 7).

System One model TypeSafe's term for a model that answers typed questions with probabilities in one fast pass, without writing text (Chapter 7).

Temperature (calibration) One number that softens or sharpens every probability a model gives, fitted on held-out data (Chapter 3).

Temperature (sampling) A setting that makes an LLM's choice of next token more or less random (Chapter 5).

Three-zone policy A policy with two thresholds, splitting cases into act, review and escalate (Chapter 11).

Threshold A number that splits cases into two groups, like a line on the probability scale: above it you do one thing, below it another (Chapters 1 and 4).

Token A piece of text, often part of a word, that an LLM reads and writes one at a time (Chapter 5).

Trace The step-by-step record of what an agent observed, decided and did (Chapter 12).

Transformer The network design built from attention layers that underlies modern LLMs (Chapter 5).

Transport The last layer of a software library, which actually sends a request over the network. Replacing it with a practice version lets you test without the real service (Chapter 14).

Triage Sorting cases by how likely they are to be real and how urgent they are, so the most important ones are handled first (Preface, Chapter 13).

Typed question A question whose answer has a fixed shape: a noul, a choice or a score (Chapter 8).

Vendor-reported A claim made by the company selling the product, not independently checked (Chapter 7).

Cheat sheets

Everything on these pages is explained fully in the chapter given in brackets. They're here for the day you need the formula and not the explanation.

Setting the threshold (Chapter 4)

For a yes-or-no decision, act when the probability of a problem is above

$$p^* = \frac{C_{\text{false alarm}}}{C_{\text{false alarm}} + C_{\text{miss}}}$$

A false alarm costing $400 and a miss costing $10,000 give $p^* = 400/10{,}400 \approx 0.04$.

Three zones (Chapter 11)

Say a review costs C_r and misses a share m of real problems. A miss costs C_{miss} and a false urgent call costs C_{page}. Then the cheapest thresholds for calibrated probabilities are

$$\text{low} = \frac{C_r}{C_{\text{miss}}\,(1-m)} \qquad \text{high} = \frac{C_{\text{page}} - C_r}{C_{\text{page}} + m\,C_{\text{miss}}}$$

Then check the queue. If the review zone holds more cases than people can clear, raise the low (act) threshold until it fits, and write down what that costs.

Measuring probabilities (Chapters 2 and 3)

For predictions p_i and outcomes $y_i \in \{0, 1\}$:

$$\text{Brier} = \frac{1}{n}\sum_i (p_i - y_i)^2 \qquad \text{log loss} = -\frac{1}{n}\sum_i \left[y_i \log p_i + (1-y_i)\log(1-p_i)\right]$$

$$\text{ECE} = \sum_b \frac{n_b}{n}\left|\,\bar{p}_b - \bar{y}_b\,\right|$$

where each bin b holds n_b predictions with average $\bar{p}_b$ and observed rate $\bar{y}_b$. AUC measures ranking only: a model can have a perfect AUC and useless probabilities.

Fixing probabilities (Chapters 3 and 9)

Method	What it fits	Use when
Temperature	$p' = \sigma(z/T)$, one number	the model is too confident (or not confident enough) everywhere
Platt scaling	$p' = \sigma(az + b)$, two numbers	a few hundred labels; the usual first fix
Isotonic	any increasing curve	thousands of labels
Prior shift	odds × $\dfrac{b'/(1-b')}{b/(1-b)}$	only the base rate changed, from b to b'

Here $z = \log \frac{p}{1-p}$ and σ is the S-curve.

Testing a model's calibration on your data (Chapter 9)

1. Collect labels for a random sample of real cases, not just the interesting ones.
2. Measure ranking (AUC) and calibration (ECE, with a bootstrap range) separately.
3. Compare the average predicted probability with the observed base rate. A big gap usually means the base rate has changed (prior shift).
4. Fix the base rate first; then use Platt scaling on a few hundred labels if needed.
5. Test again on labels you didn't fit on. Repeat on a schedule and whenever the model version changes.

Designing typed questions (Chapter 8)

- **Noul**: one fact per question; say what counts as yes and as no; no double negatives.
- **Choice**: options that don't overlap; always an escape option ("other", "none", "not enough information").
- **Score**: put thresholds on the probability of the levels you care about, not on the average level.
- Ask related questions in one call: the state is sent once.
- Decide with all the probabilities and the costs of each action, not just the top label.

Calling a decision model (Chapters 7, 8 and 14)

- Ask several typed questions about the same situation in **one call**: yes/no (noul), pick one (choice), rate on a scale (score).
- Every choice needs an escape option: "none", "other", "not enough information".

- **Pin the model version.** "latest" can change under you.
- Retry rate limits and server errors; never retry your own mistakes, fix them.
- Decide in advance what a **failed call means**; the safe default is "send to a person".
- Log the probabilities, the policy version and the request ID with every decision.

Choosing a method (Chapter 10)

If you have…	Start with
thousands of labels, clean fields, one stable decision	logistic regression
a rule an expert can write and an auditor must read	rules
no labels, raw text or many different decisions	a decision model, checked on a few hundred labels
messy input and a clear decision	an LLM to extract, a decision model to decide
output a person must read	an LLM, checked by typed questions

Before a decision service runs (Chapter 14)

A fixed model version · a config with a version number and a fingerprint · a meaning for every failure · a record for every decision · random audits of automatic actions · a daily monitor · shadow runs for new versions · an owner who can switch everything to review.

Bibliography

Casscells, Ward, Arno Schoenberger, and Thomas B. Graboys. 1978. 'Interpretation by Physicians of Clinical Laboratory Results'. *New England Journal of Medicine* 299 (18): 999–1001.

Chow, C. K. 1970. 'On Optimum Recognition Error and Reject Tradeoff'. *IEEE Transactions on Information Theory* 16 (1): 41–46.

Elkan, Charles. 2001. 'The Foundations of Cost-Sensitive Learning'. *Proceedings of the Seventeenth International Joint Conference on Artificial Intelligence (IJCAI)*, 973–78.

Frederick, Shane. 2005. 'Cognitive Reflection and Decision Making'. *Journal of Economic Perspectives* 19 (4): 25–42.

Gigerenzer, Gerd, and Ulrich Hoffrage. 1995. 'How to Improve Bayesian Reasoning Without Instruction: Frequency Formats'. *Psychological Review* 102 (4): 684–704.

Greshake, Kai, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. 2023. 'Not What You've Signed up for: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection'. *Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security*.

Guo, Chuan, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. 'On Calibration of Modern Neural Networks'. *Proceedings of the 34th International Conference on Machine Learning (ICML)*.

Jevons, William Stanley. 1865. *The Coal Question: An Inquiry Concerning the Progress of the Nation, and the Probable Exhaustion of Our Coal-Mines*. Macmillan.

Jiang, Dongming, Yi Li, and Bingzhe Li. 2026. *Jev-Mem: System-One-Controlled Agentic Memory for Efficient AI Agents*. arXiv:2609.23986. Accessed 29 September 2026. https://arxiv.org/abs/2609.23986.

Kahneman, Daniel. 2011. *Thinking, Fast and Slow*. Farrar, Straus and Giroux.

Kahneman, Daniel, and Gary Klein. 2009. 'Conditions for Intuitive Expertise: A Failure to Disagree'. *American Psychologist* 64 (6): 515–26.

Kaplan, Jared, Sam McCandlish, Tom Henighan, et al. 2020. *Scaling Laws for Neural Language Models*. https://arxiv.org/abs/2001.08361.

LeCun, Yann, Yoshua Bengio, and Geoffrey Hinton. 2015. 'Deep Learning'. *Nature* 521: 436–44.

Lewis, Patrick, Ethan Perez, Aleksandra Piktus, et al. 2020. 'Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks'. *Advances in Neural Information Processing Systems 33 (NeurIPS)*.

Molas, Alex. 2026. *Jev Can't Be Calibrated*. Blog post, 23 September 2026. Accessed 29 September 2026. https://www.alexmolas.com/2026/09/23/jev-cant-be-calibrated.html.

Murphy, Allan H., and Robert L. Winkler. 1977. 'Reliability of Subjective Probability Forecasts of Precipitation and Temperature'. *Journal of the Royal Statistical Society, Series C (Applied Statistics)* 26 (1): 41–47.

Naeini, Mahdi Pakdaman, Gregory F. Cooper, and Milos Hauskrecht. 2015. 'Obtaining Well Calibrated Probabilities Using Bayesian Binning'. *Proceedings of the AAAI Conference on Artificial Intelligence*.

OpenAI. 2023. *GPT-4 Technical Report*. https://arxiv.org/abs/2303.08774.

Ouyang, Long, Jeff Wu, Xu Jiang, et al. 2022. 'Training Language Models to Follow Instructions with Human Feedback'. *Advances in Neural Information Processing Systems 35 (NeurIPS)*.

Platt, John C. 1999. 'Probabilistic Outputs for Support Vector Machines and Comparisons to Regularized Likelihood Methods'. In *Advances in Large Margin Classifiers*. MIT Press.

regolo.ai. 2026. *Jev and System One Models: Benchmarks, Open-Source Alternatives, and When to Use Them*. Blog post. Accessed 29 September 2026. https://regolo.ai/jev-and-system-one-models-benchmarks-open-source-alternatives-and-when-to-use-them/.

Rumelhart, David E., Geoffrey E. Hinton, and Ronald J. Williams. 1986. 'Learning Representations by Back-Propagating Errors'. *Nature* 323: 533–36.

Saerens, Marco, Patrice Latinne, and Christine Decaestecker. 2002. 'Adjusting the Outputs of a Classifier to New a Priori Probabilities: A Simple Procedure'. *Neural Computation* 14 (1): 21–41.

Turing, Alan M. 1950. 'Computing Machinery and Intelligence'. *Mind* 59 (236): 433–60.

TypeSafe AI. 2026. *Introducing System One Models & Jev*. TypeSafe AI blog. Accessed 29 September 2026. https://typesafe.ai/blog/introducing-system-one-models-and-jev.

Vaswani, Ashish, Noam Shazeer, Niki Parmar, et al. 2017. 'Attention Is All You Need'. *Advances in Neural Information Processing Systems 30 (NeurIPS)*.

Yao, Shunyu, Jeffrey Zhao, Dian Yu, et al. 2023. 'ReAct: Synergizing Reasoning and Acting in Language Models'. *International Conference on Learning Representations (ICLR)*.

Zadrozny, Bianca, and Charles Elkan. 2002. 'Transforming Classifier Scores into Accurate Multiclass Probability Estimates'. *Proceedings of the Eighth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 694–99.

Index

A
Act, review, escalate, 63, 64
Agent, 34, 39, 70, 76
Attention, 28, 30
AUC, 16

B
Bake-off, 57, 62
Base rate, 8, 14
Bayes' rule, 10, 52

C
Calibration, 15, 21, 51, 56
Capacity, 66, 68
Choice (question type), 46, 49
Cost-based threshold, 25

D
Decision layer, 38, 77, 91
Decision record, 87
Doom demo, 44
Drift, 68

E
ECE (expected calibration error), 17, 20
Embedding, 28, 29
Extract, then decide, 70, 71

G
Gradient descent, 12
Guardrail gate, 72

H
Harbor Pharma, 52
Hybrid agent, 73, 75

I
Isotonic regression, 18, 19

J
Jev-Mem, 72
Jevons paradox, 92
Jevons, William Stanley, 92
JSON mode, 31, 32

K
Kahneman, Daniel, 5, 40

L
LLM (large language model), 28, 32
LLM-as-judge, 58
Log loss, 11, 12
Logistic regression, 11
Loss, 11, 13

M
Molas, Alex, 51

N
Noul, 46, 47

O
Overconfidence, 28, 32
Overfitting, 13

P
Phishing campaign, 68, 77, 86
Platt scaling, 18, 19
Policy, 63, 69
Prompt injection, 37

R
RAG (retrieval-augmented generation), 34, 35
Reasoning models, 72
Reliability diagram, 15, 18
RLCD, 43
Router, 72, 76

S
Score (question type), 46, 48
Shadow mode, 77, 83, 86, 90
State, 43
System 1, System 2, 40, 45
System One model, 40, 45

T
Temperature (calibration), 18, 20
Threat intel, 74
Three-zone policy, 66
Threshold, 3, 22, 27

Token, 31, 32
Trace, 73
Transformer, 30
Transport, 85
TypeSafe AI, 5, 40, 47, 51

About the Author

Sridhar Mukkandi is an applied AI engineer who builds agents and the decision systems behind them. He writes for engineers who want AI systems they can trust.

sridhar.authorhub@gmail.com